



National Institute for Space Research – INPE



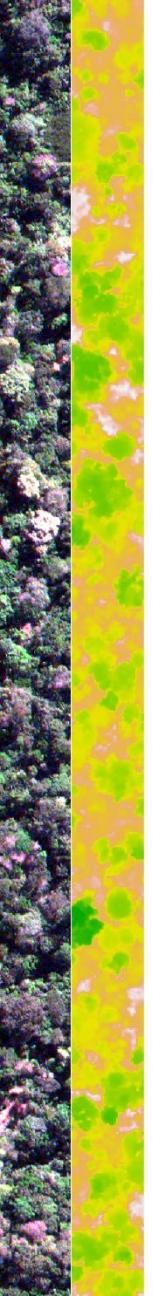
Mini-course @ WORCAP 2021

Deep Learning applied for Remote Sensing images



Ricardo Dal'Agnol da Silva
Post-doc Researcher - INPE

15 Sep 2021



Learning objectives

- Describe what is deep learning
- Explain the innovation brought by convolutional neural networks (CNN)
- Identify the main steps on applying CNN for remote sensing
- Apply an example of CNN with test dataset

Content

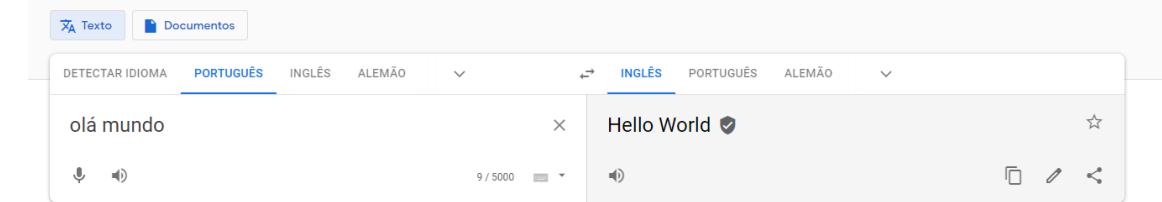
- Lecture - basic concepts
- R and Google Colab

"AI is the new electricity" - Andrew Ng

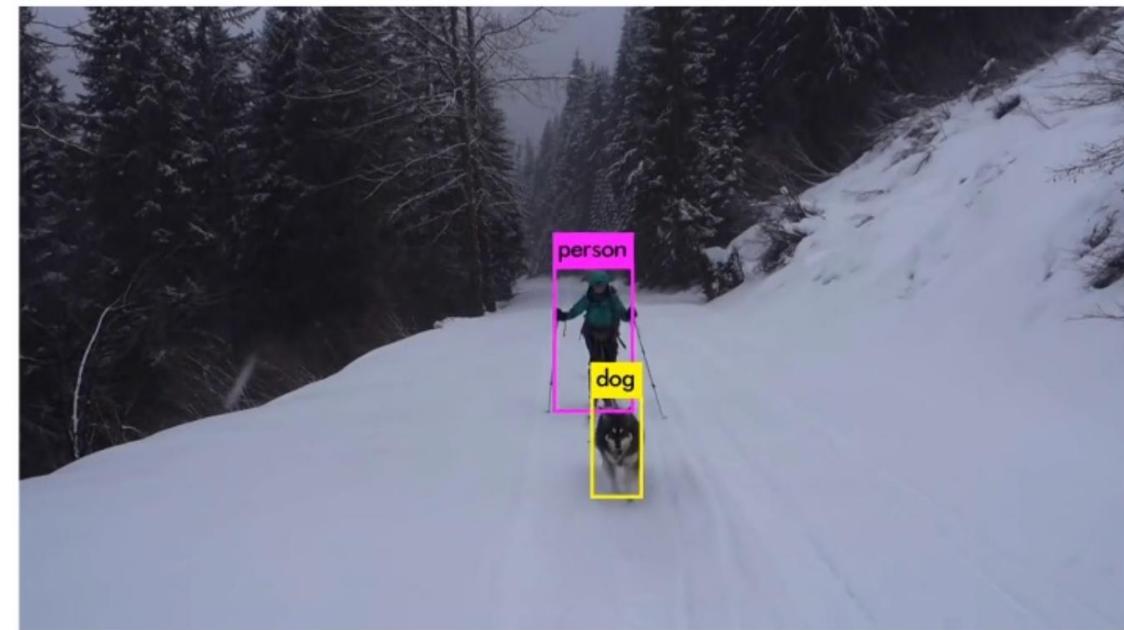
Gary Chavez added a photo you might be in.
about a minute ago · 



OpenAI Five: Dota2
The International 2018



X Texto Documentos DETECTAR IDIOMA PORTUGUÊS INGLÊS ALEMÃO INGLÊS PORTUGUÊS ALEMÃO olá mundo Hello World 9 / 5000

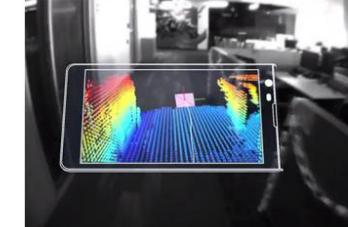


Robot Perception



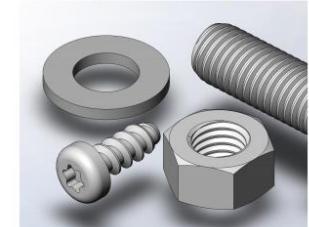
source: Scott J Grunewald

Augmented Reality



source: Google Tango

Shape Design



source: solidsolutions

Classification vs Detection vs Semantic vs Instance segmentation

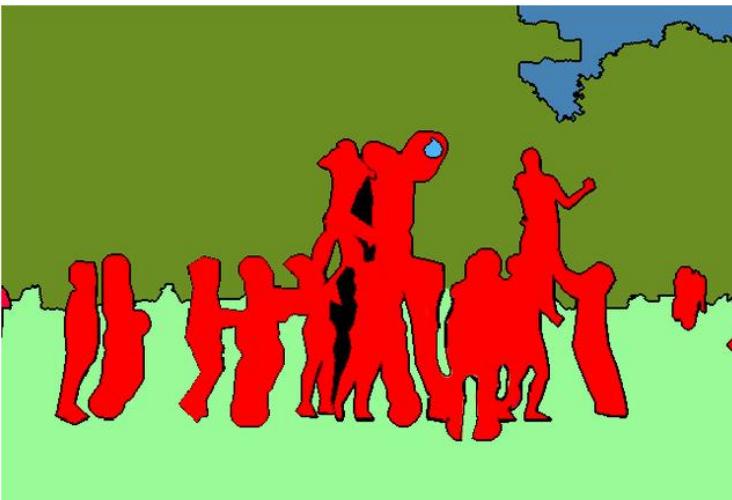
Image Classification



Object Detection

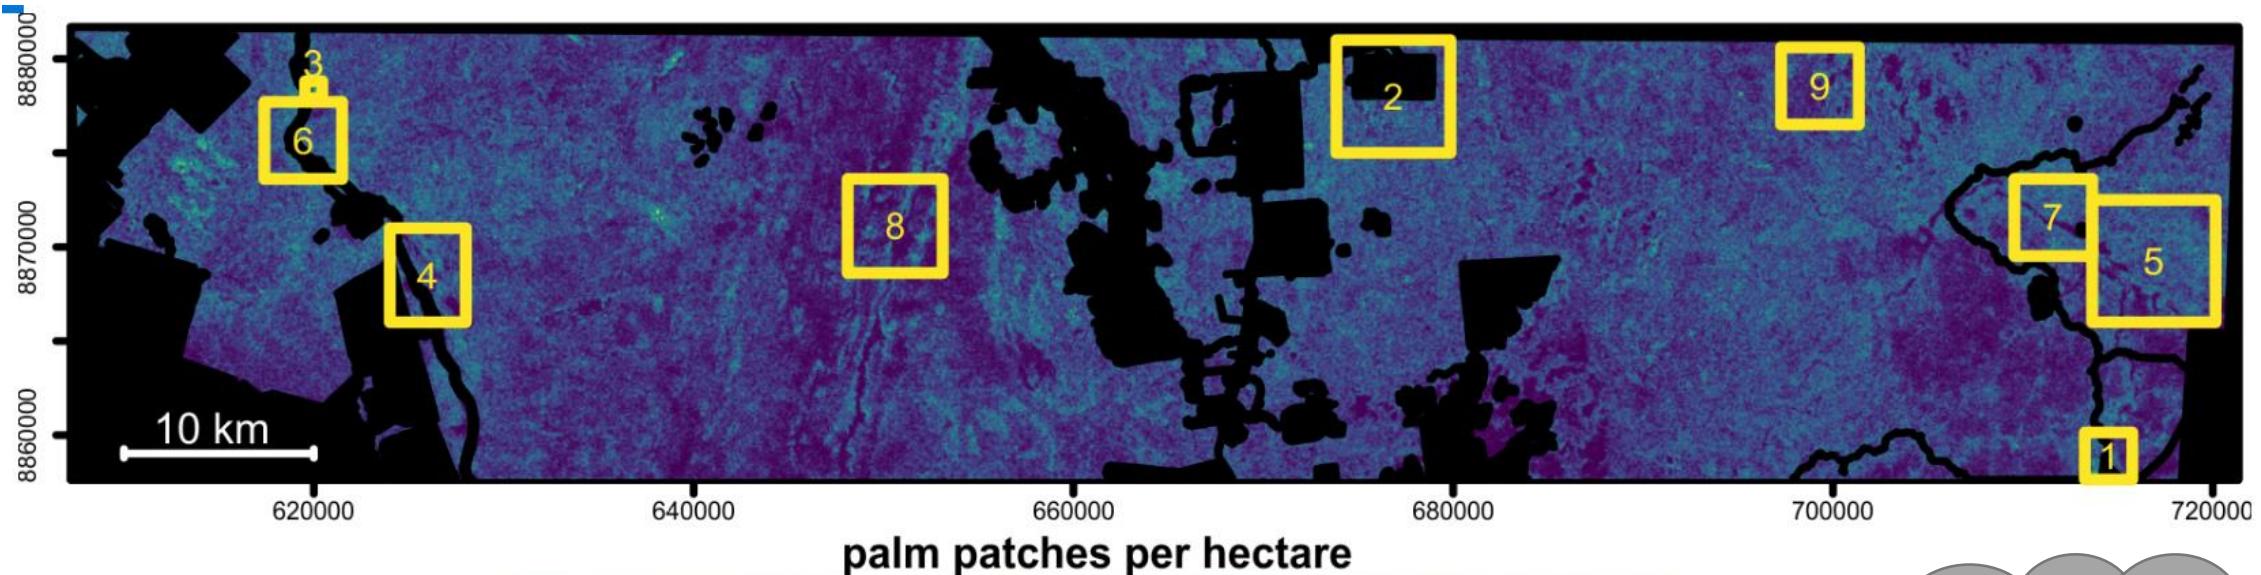


Semantic Segmentation

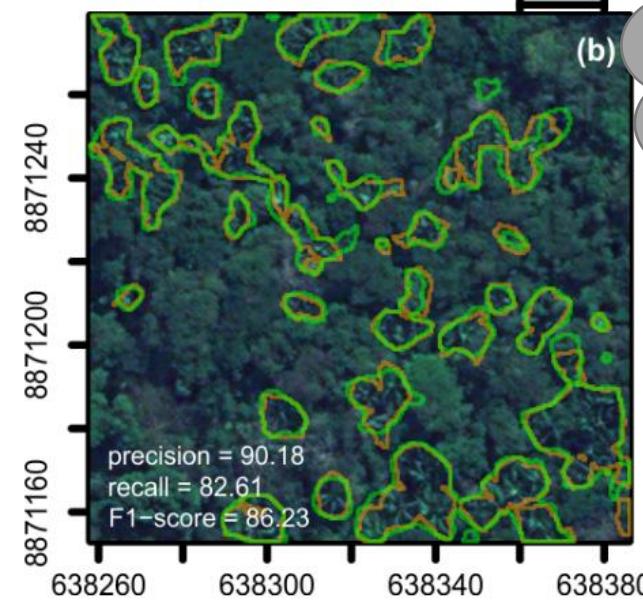
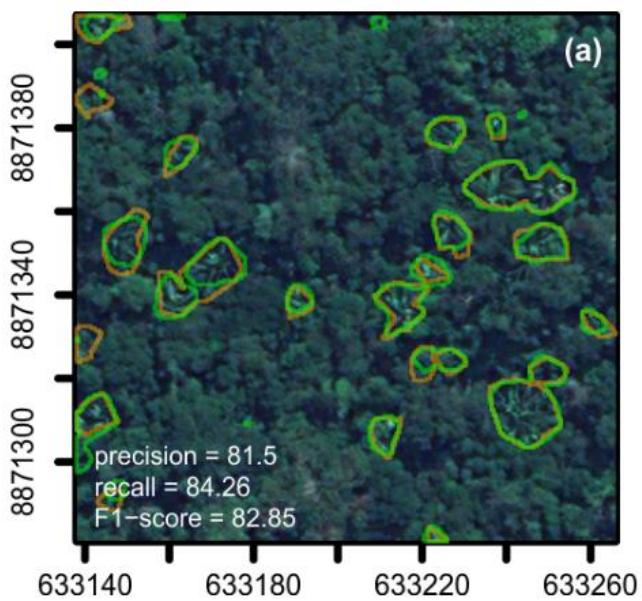


Instance Segmentation





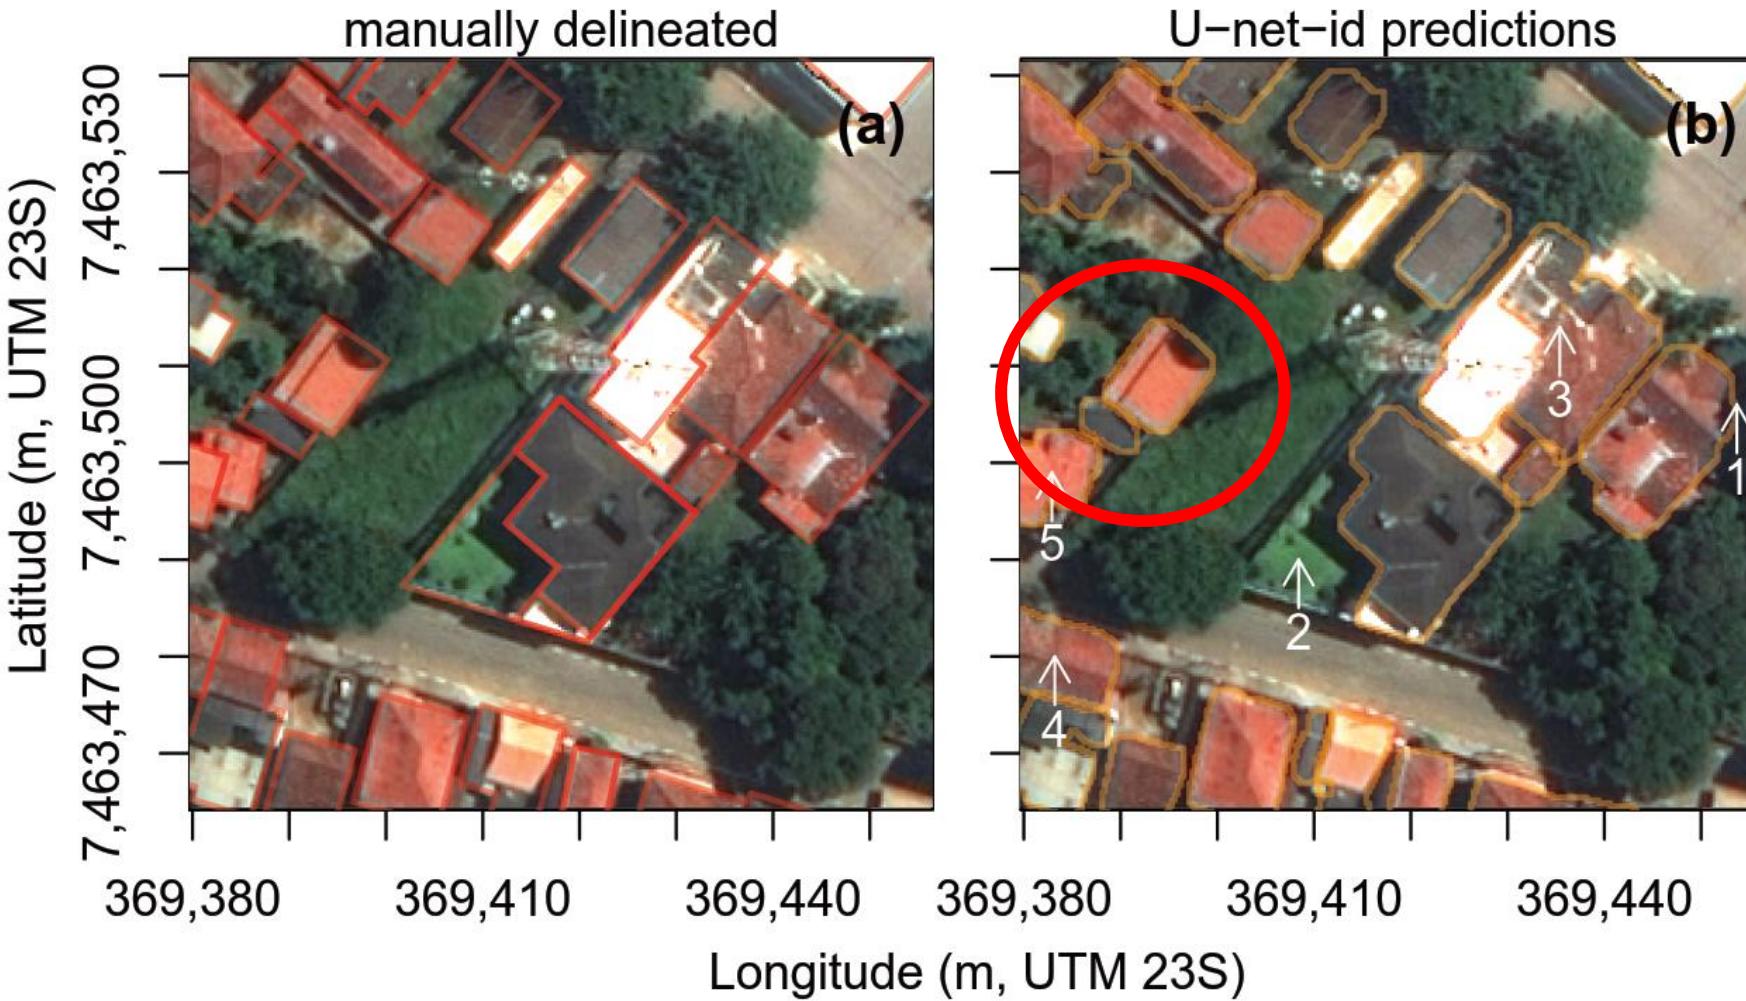
palm patches per hectare



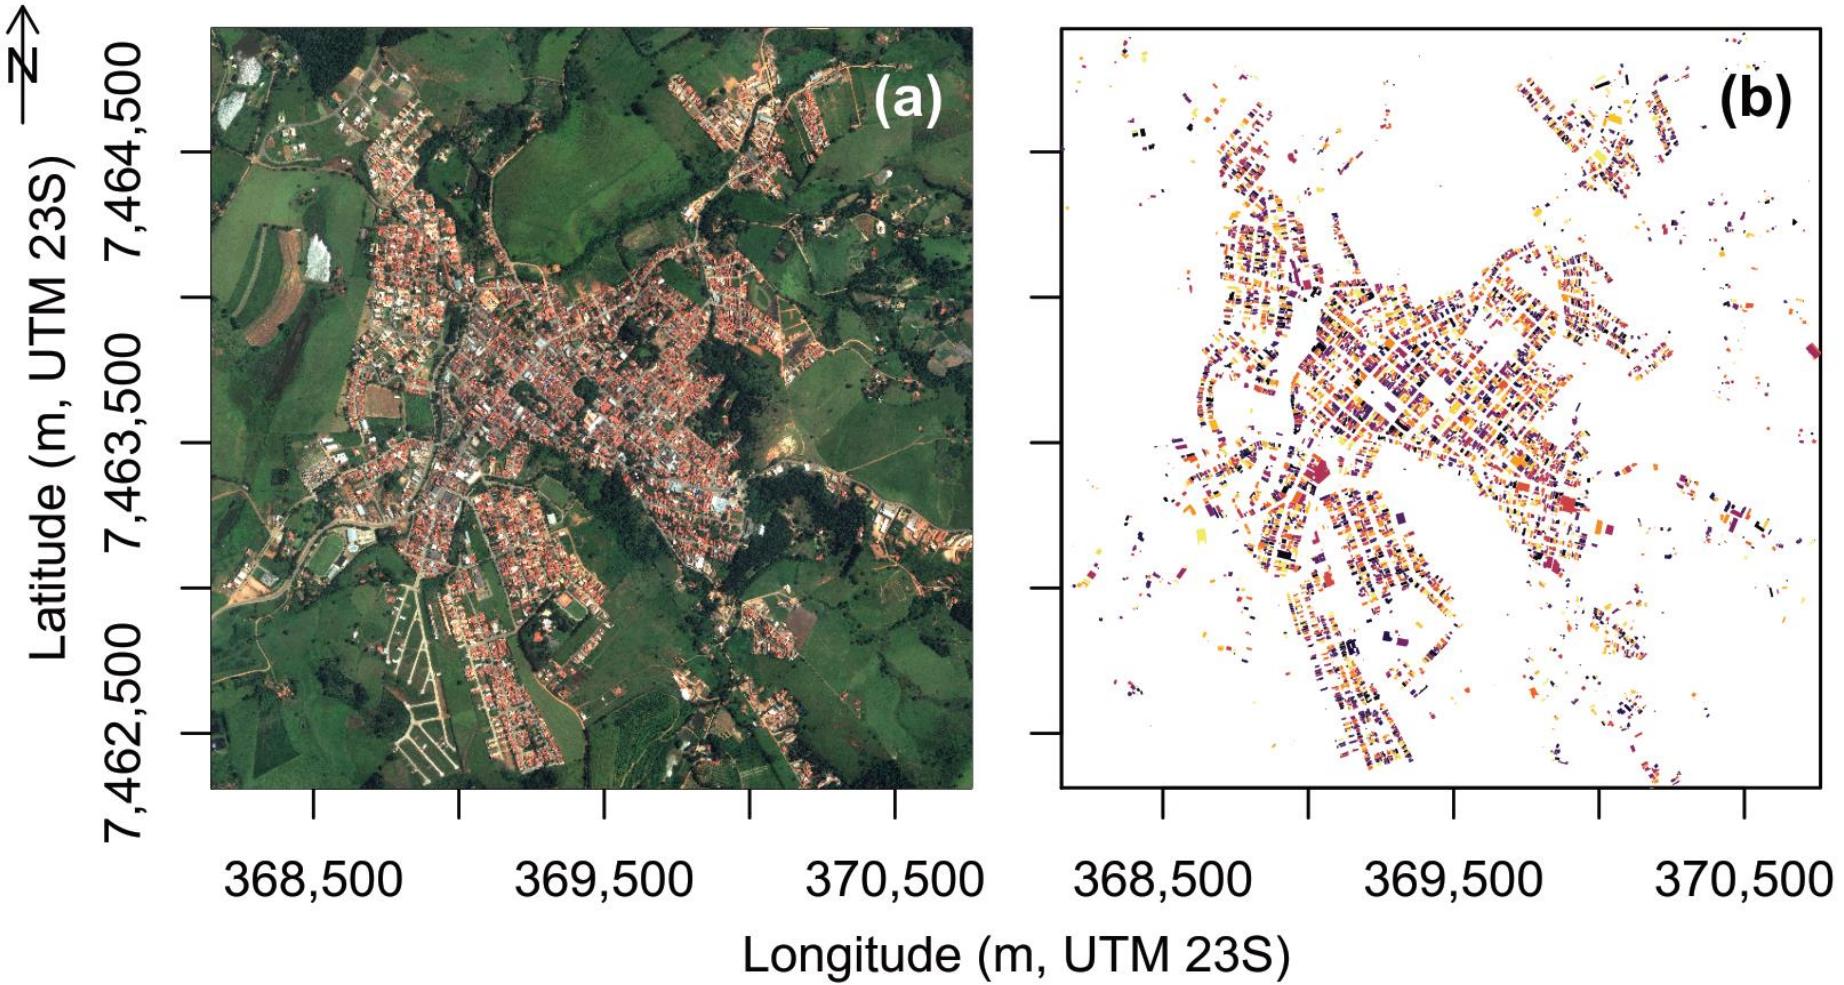
*False Positives are
very often
True Positives...*

- manual delineation
- automatic delineation

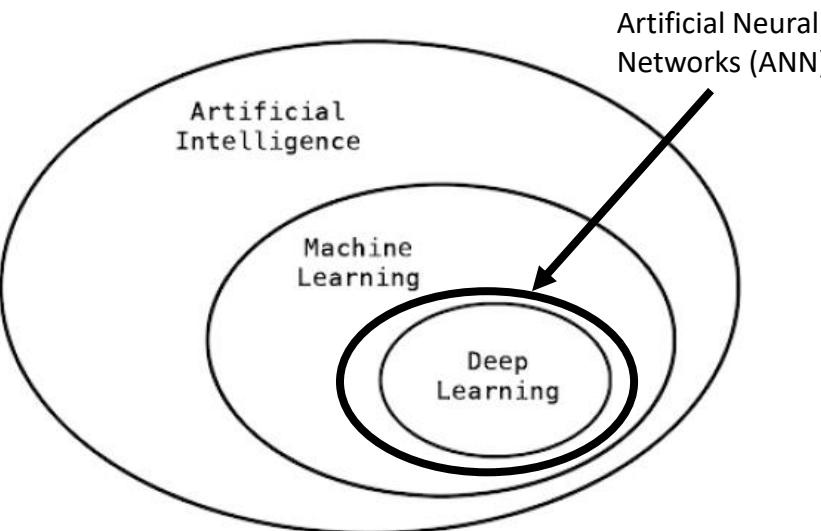
Building segmentation (Wagner et al. 2020)



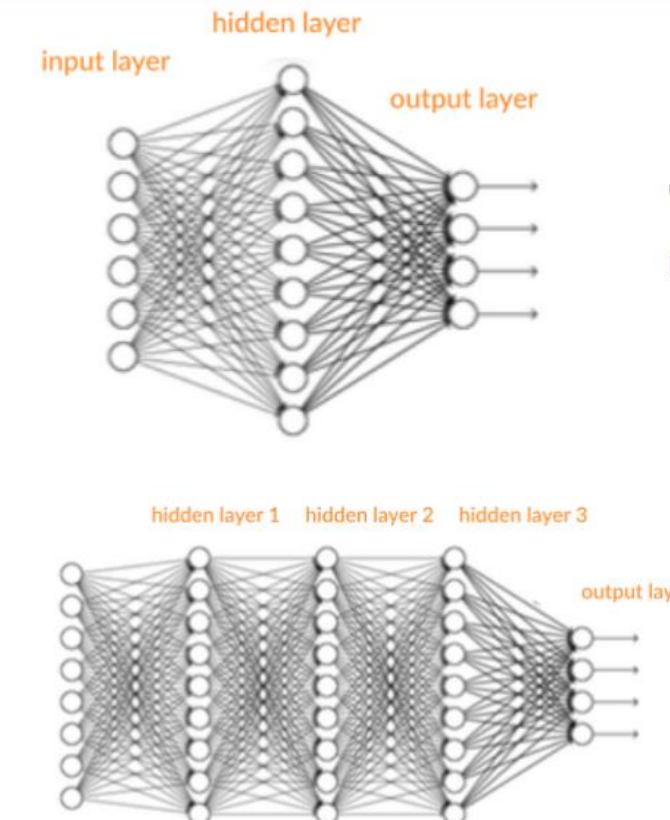
Building segmentation (Wagner et al. 2020)



AI -> ML -> ANN -> DL

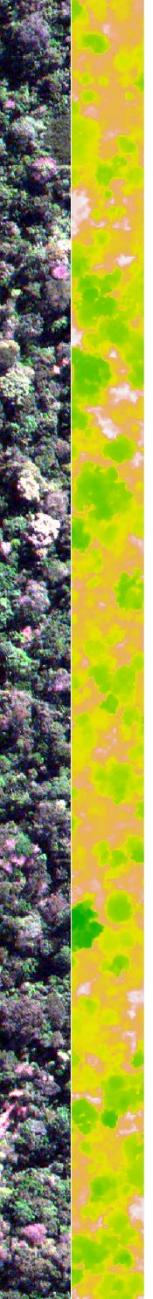


(Adapted from Chollet & Allaire 2017)



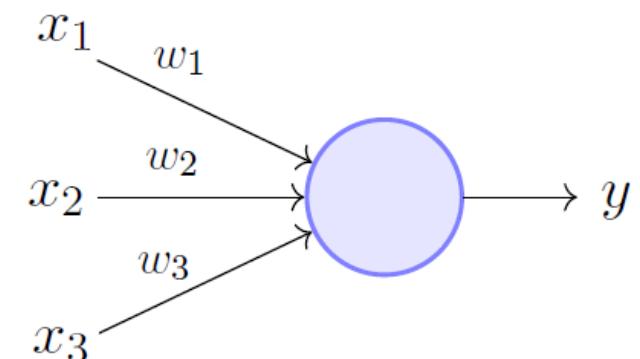
"Non-deep" feedforward neural network

Deep neural network



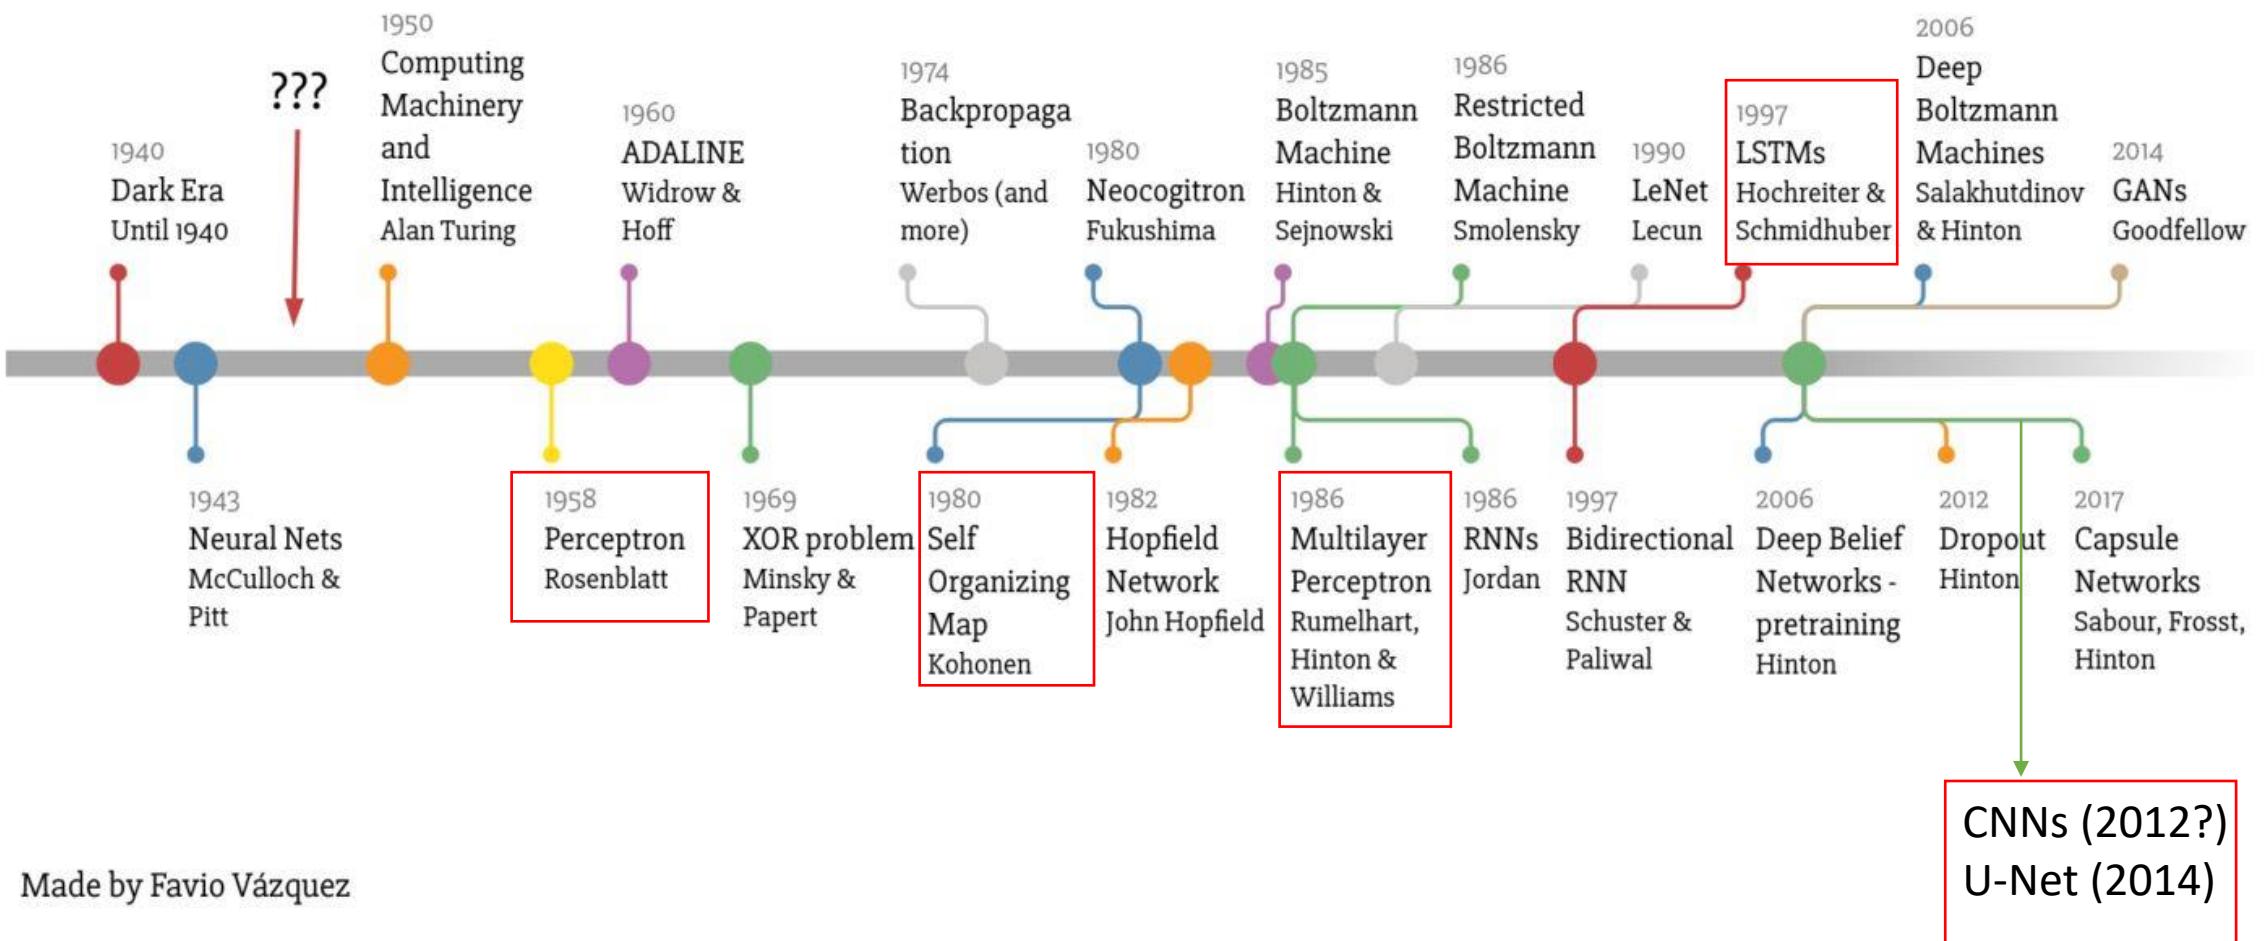
Neural networks

- Goal: approximate some mathematical function
- Relate inputs (x) to outputs (y) using weights (w) in between neurons
- Can have many layers with neurons
- Applications for regression or classification problems
- Usually supervised but can also be unsupervised



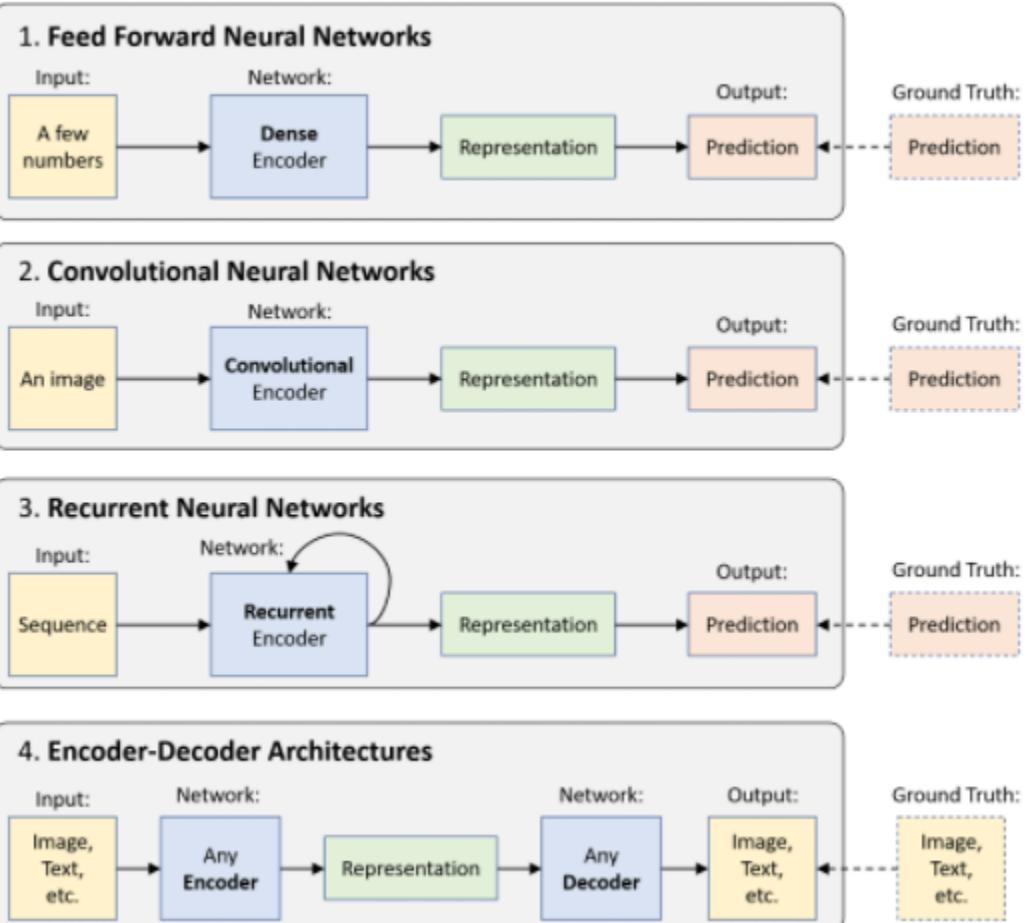
Perceptron Model (Minsky-Papert in 1969)

Deep Learning Timeline

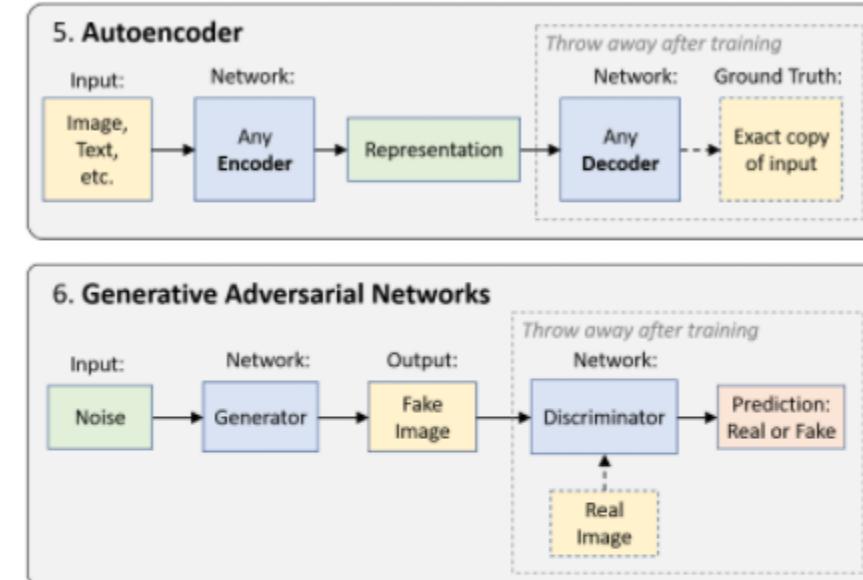


Main DL model types

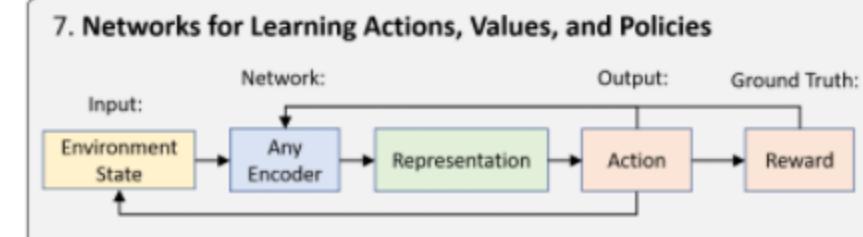
Supervised Learning



Unsupervised Learning

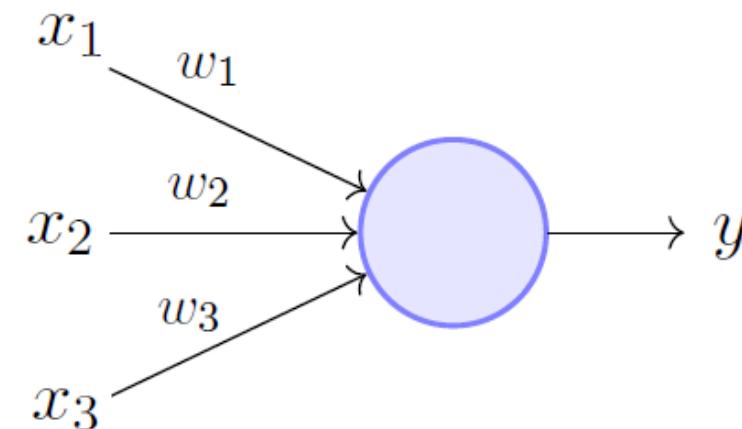


Reinforcement Learning



Quick intro into ANN - Perceptron 1958

- Basically a linear model
- *feed forward* type
- Weights (w) are adjusted and multiplied by inputs (x) to obtain the output (y)

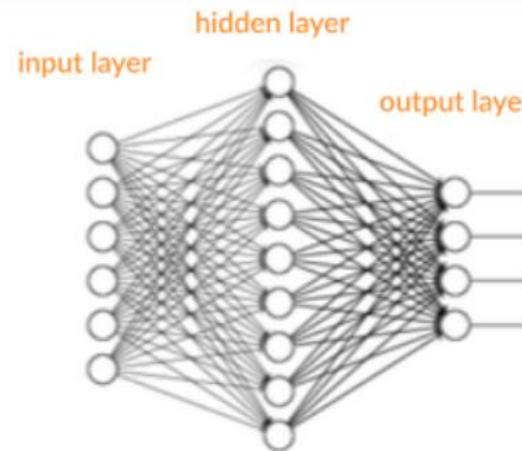


Simplifying to:
 $y = x_1 * w_1 + x_2 * w_2 + x_3 * w_3$

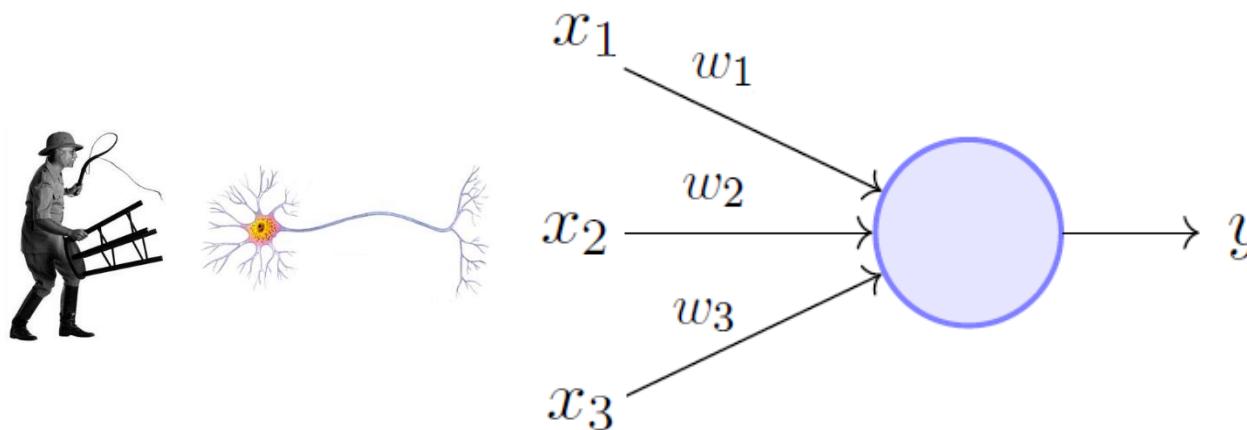
Perceptron Model (Minsky-Papert in 1969)

Perceptron 1958

- Basically a linear model
- *feed forward* type
- Weights (w) are adjusted and multiplied by inputs (x) to obtain the output (y)



Multi-Layer Perceptron (MLP)
1986

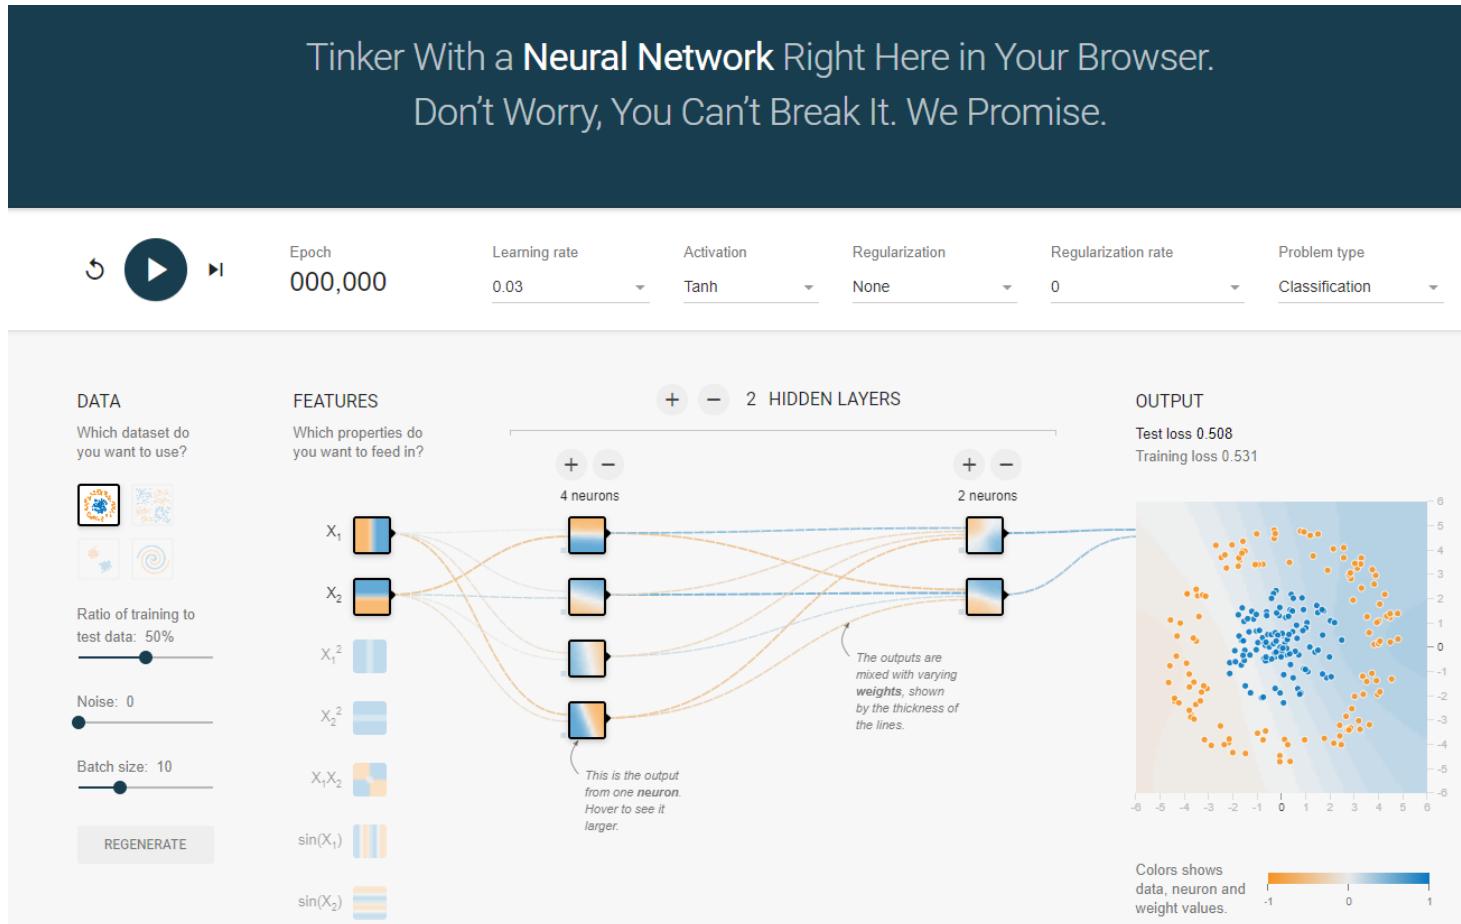


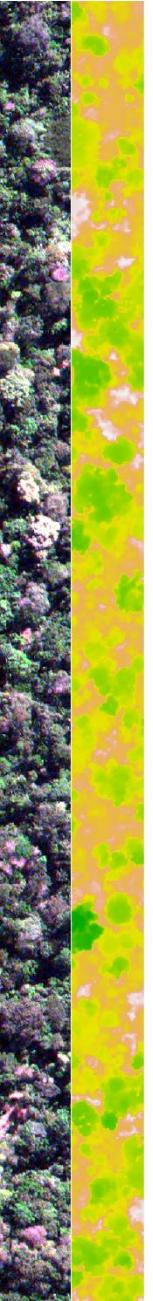
Simplifying to:
 $y = x_1 * w_1 + x_2 * w_2 + x_3 * w_3$

Perceptron Model (Minsky-Papert in 1969)

Playing with a MLP

- <https://playground.tensorflow.org/>





Why deep learning for remote sensing?

- Classification and segmentation, but also regression

>	
Method	
Number of layers	
Accuracy	
Domain	

Past
• Multi-layer Perceptron (MLP)
• A few, e.g. usually < 5
• Similar accuracy to traditional machine learning algorithms
• Pixel-level (or region)

Why deep learning for remote sensing?

- Classification and segmentation, but also regression

>

Method

Number of layers

Accuracy

Domain

Past

- Multi-layer Perceptron (MLP)
- A few, e.g. usually < 5
- Similar accuracy to traditional machine learning algorithms
- Pixel-level (or region)

Present

- Convolutional Neural Networks (CNN)
- Many and growing, e.g. >20
- Above traditional machine learning methods and close to human level
- Patch-level

Convolutional neural networks: state-of-the-art methods for semantic segmentation

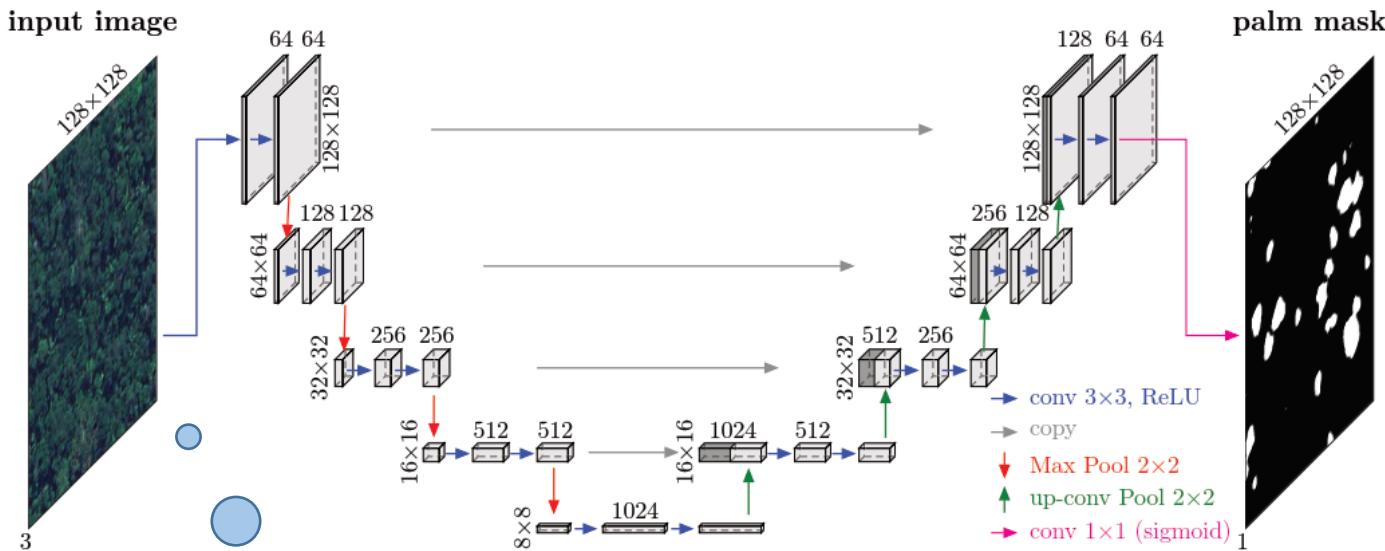


Figure 2. U-net architecture used for canopy palm segmentation, adapted from [19]. The number of channels and the row \times column size in pixels are indicated for each cuboid.

The inputs are
image patches
instead of
pixels



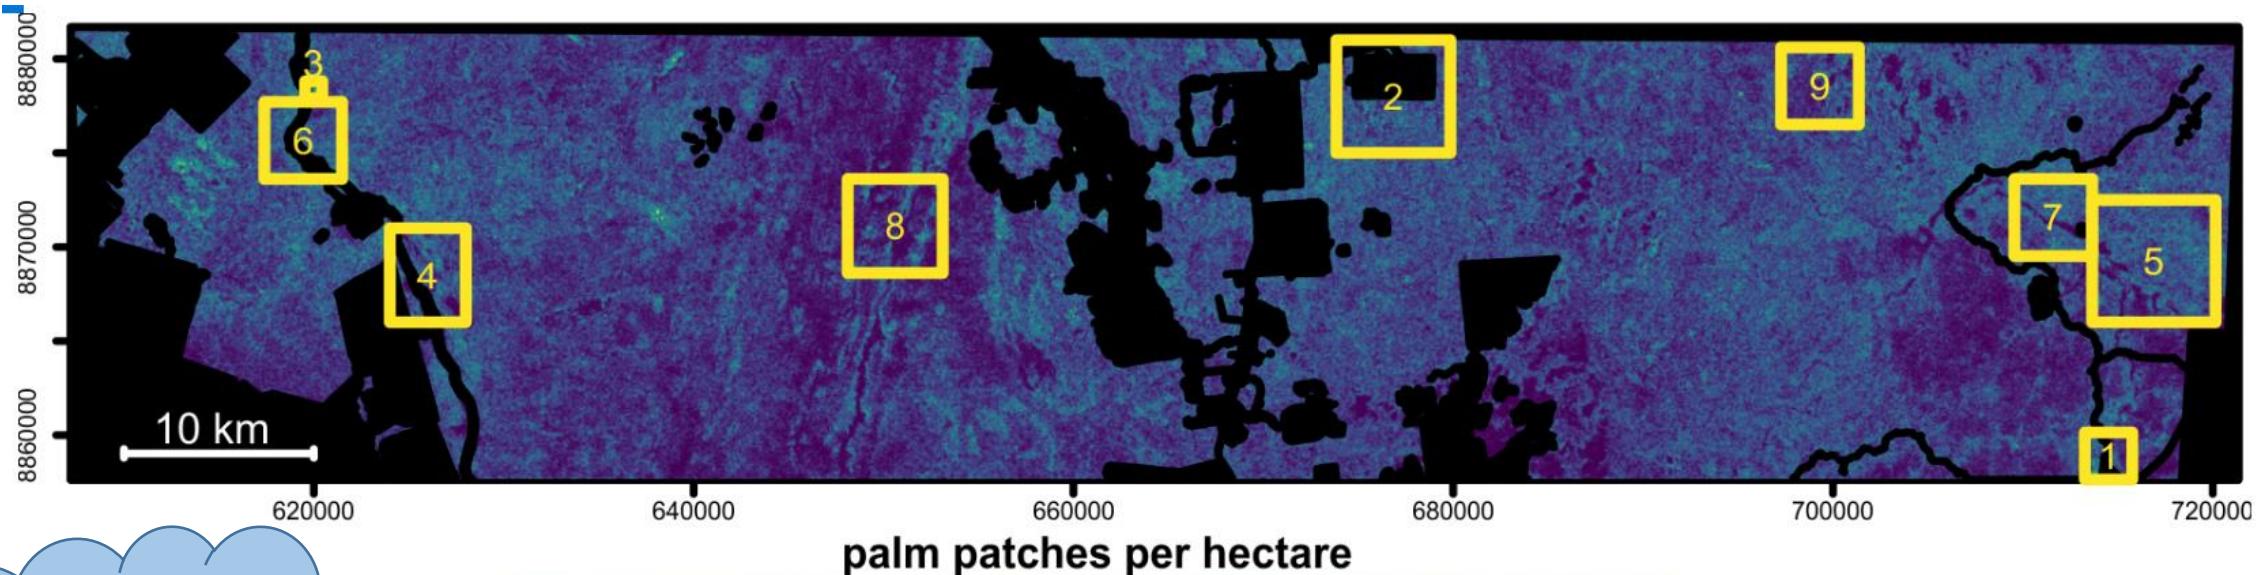
remote sensing

Article

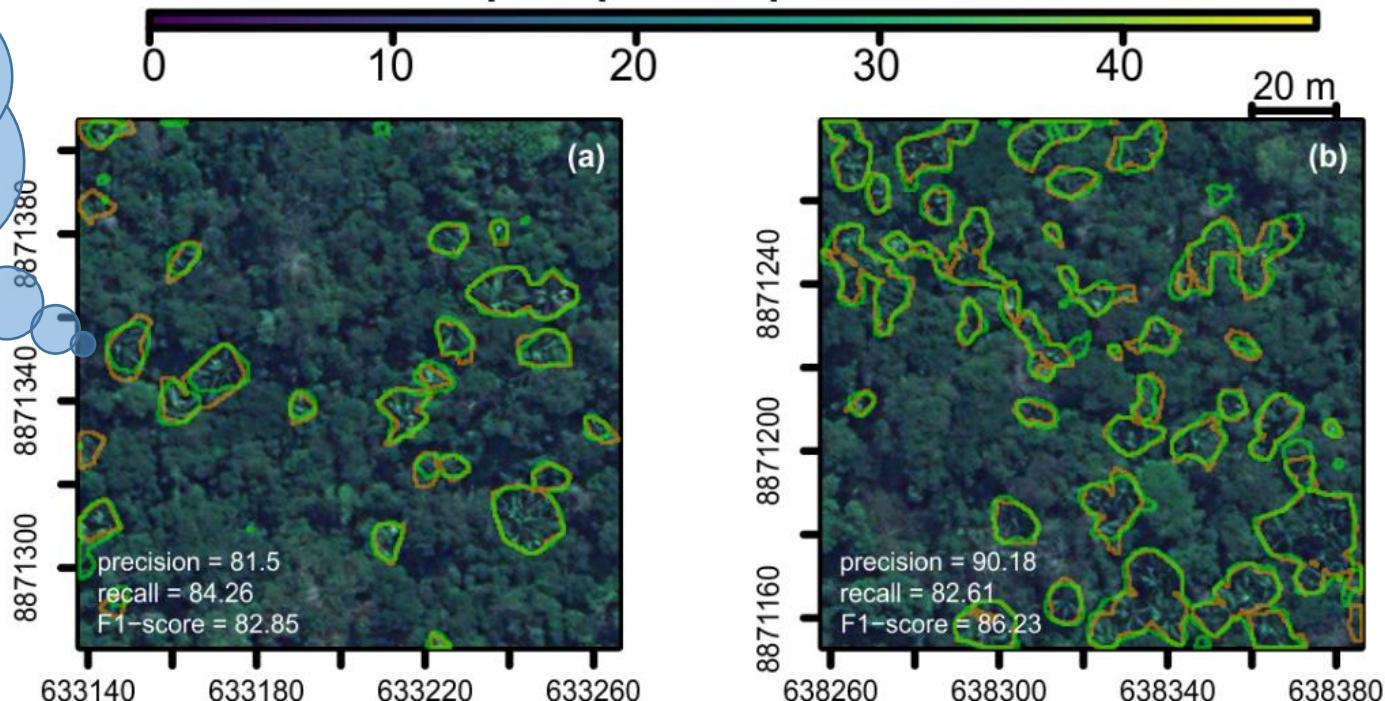
Regional Mapping and Spatial Distribution Analysis of Canopy Palms in an Amazon Forest Using Deep Learning and VHR Images



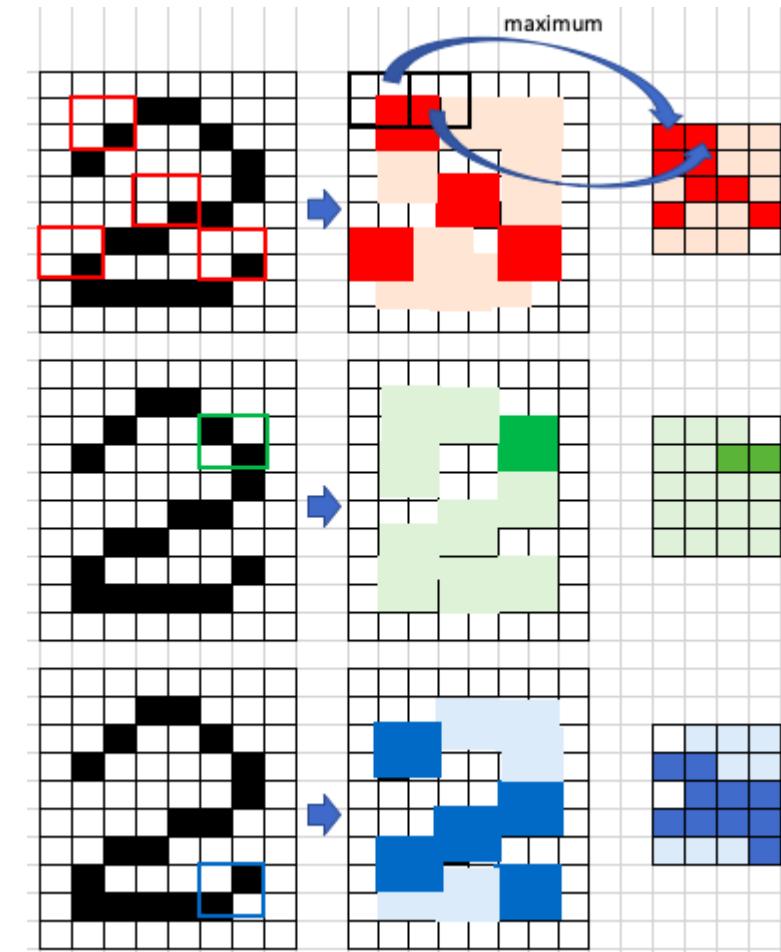
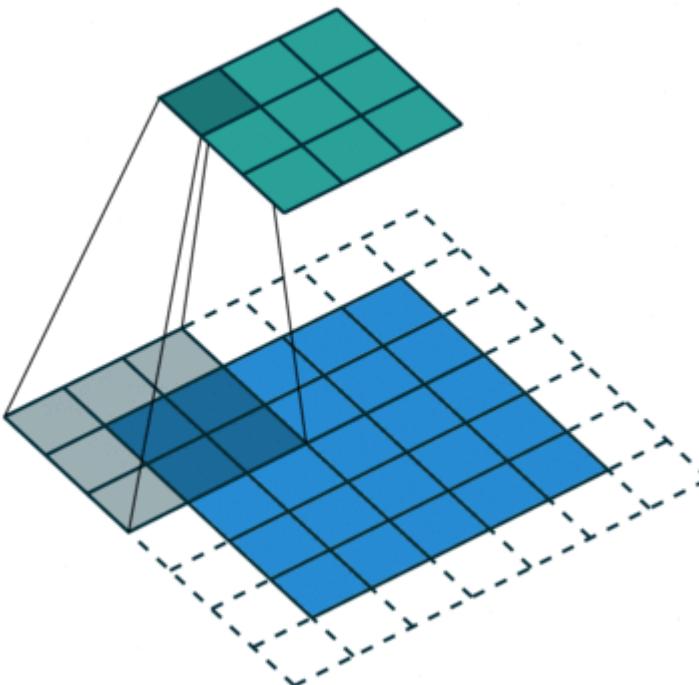
Fabien H. Wagner ^{1,*} , Ricardo Dalagnol ² , Ximena Tagle Casapia ^{3,4} , Annia S. Streher ², Oliver L. Phillips ⁵ , Emanuel Gloor ⁵ and Luiz E. O. C. Aragão ^{2,6}



If you can see
the target in
the image,
deep learning
can map it



Convolution



Convolution

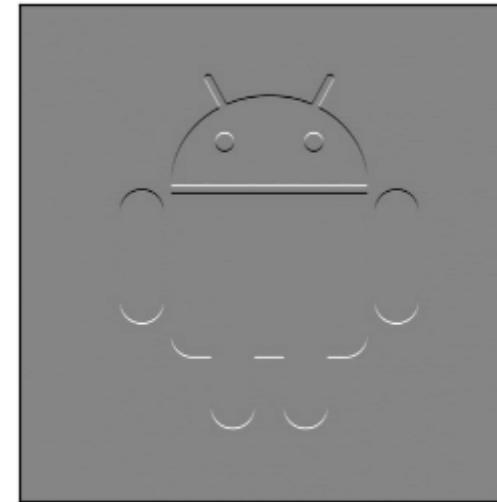


-1	-2	-1
0	0	0
1	2	1

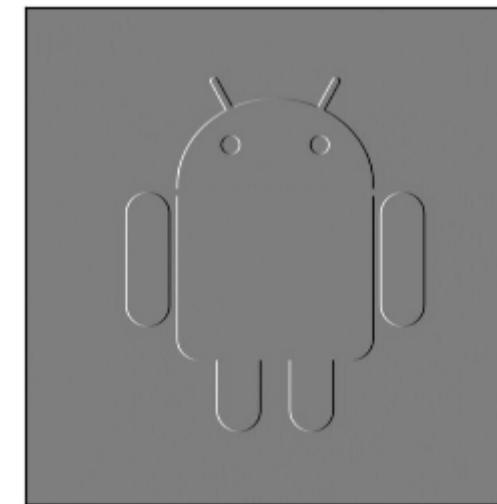
Finds horizontals

-1	0	1
-2	0	2
-1	0	1

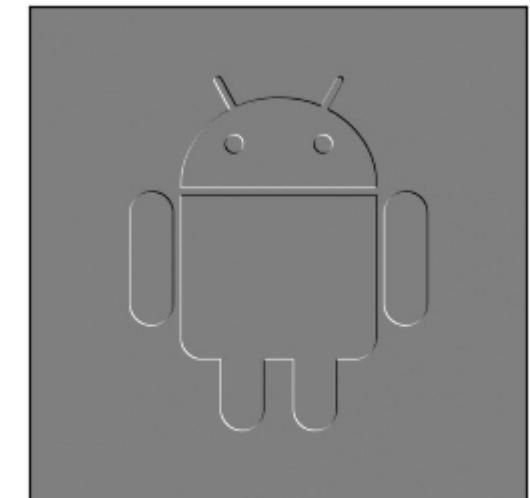
Finds verticals



Horizontal Sobel



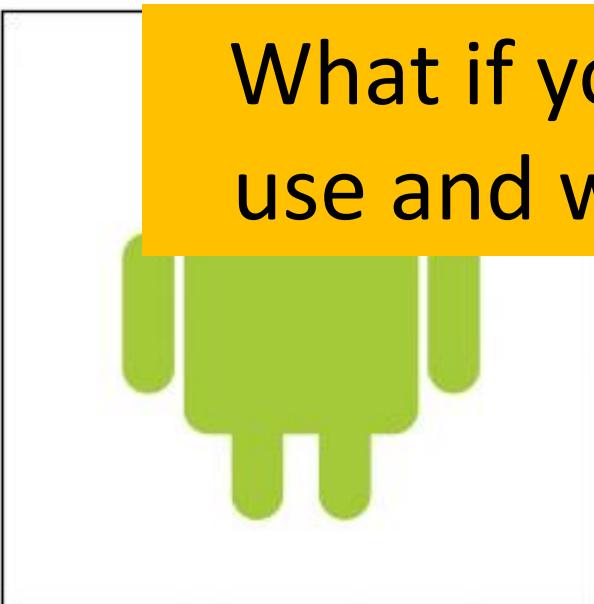
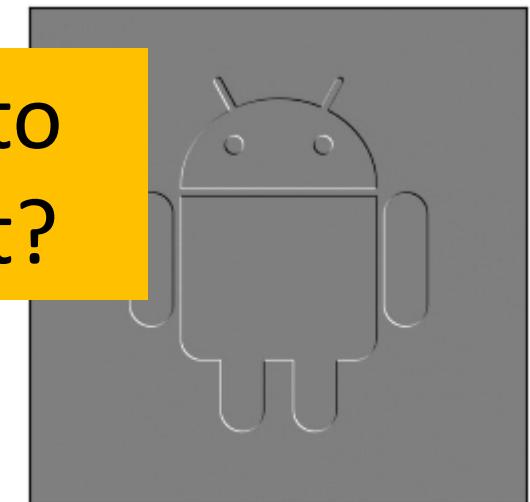
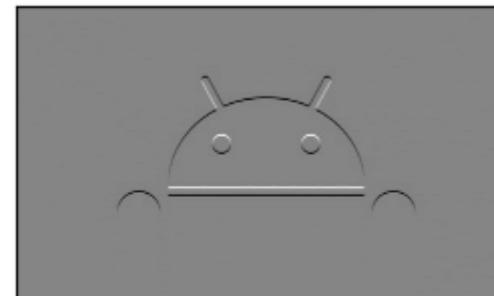
Vertical Sobel



Combined Sobel

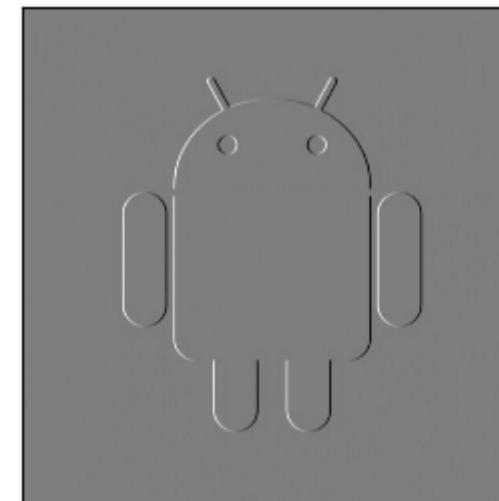
Convolution

-1	-2	-1



-1	0	1
-2	0	2
-1	0	1

Finds verticals



Vertical Sobel

CNN

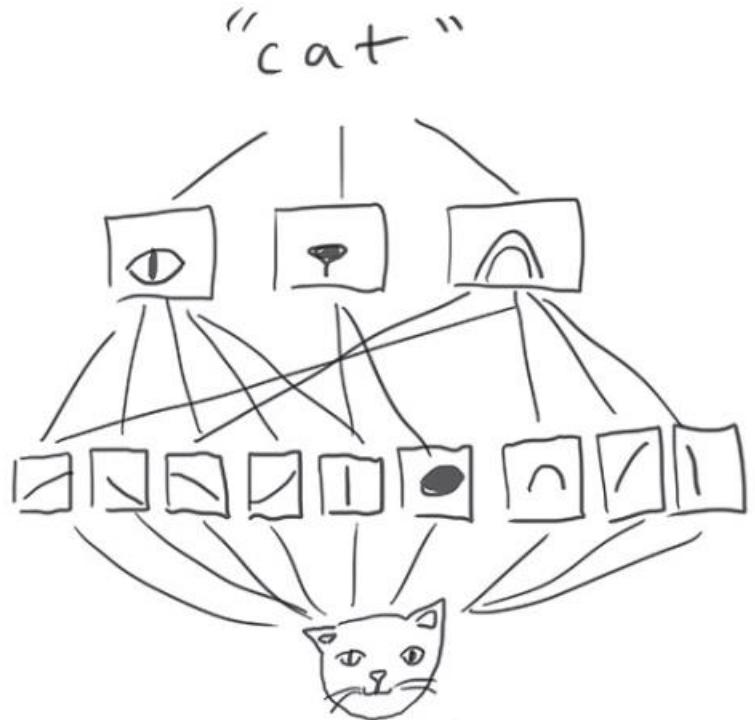
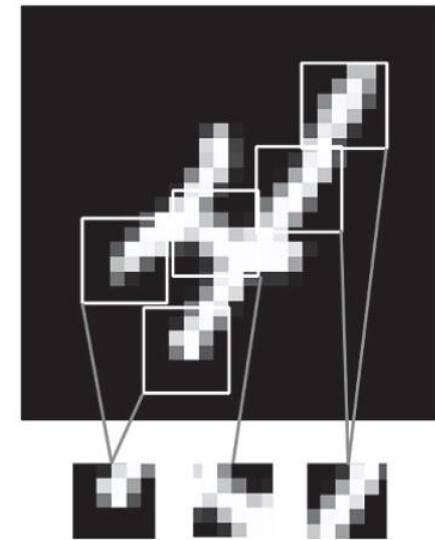
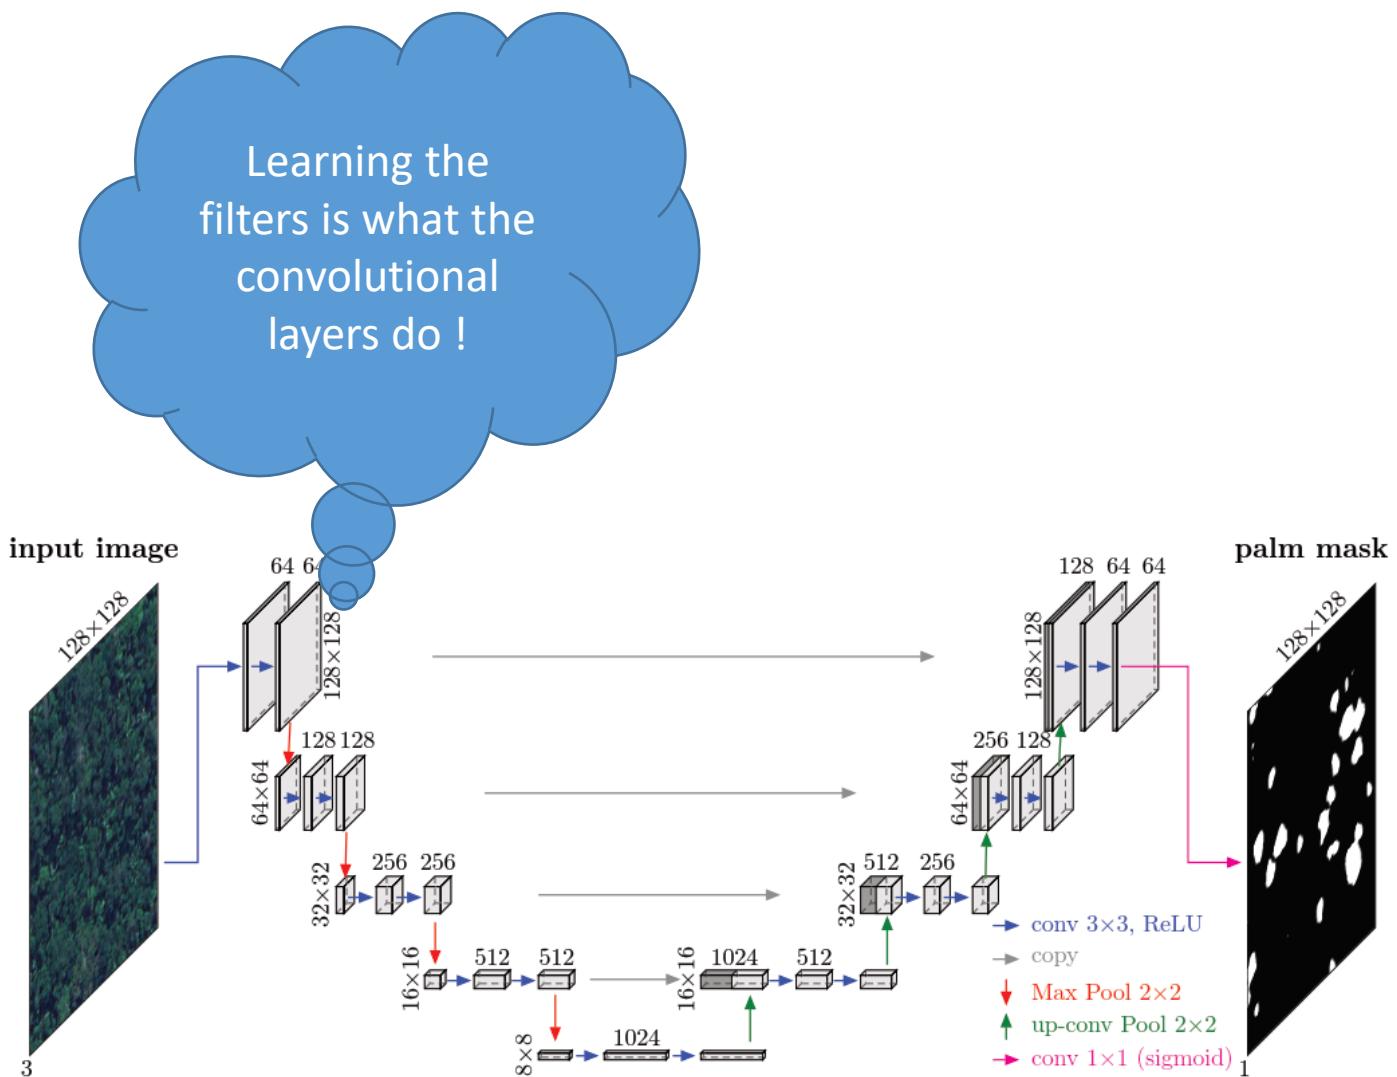
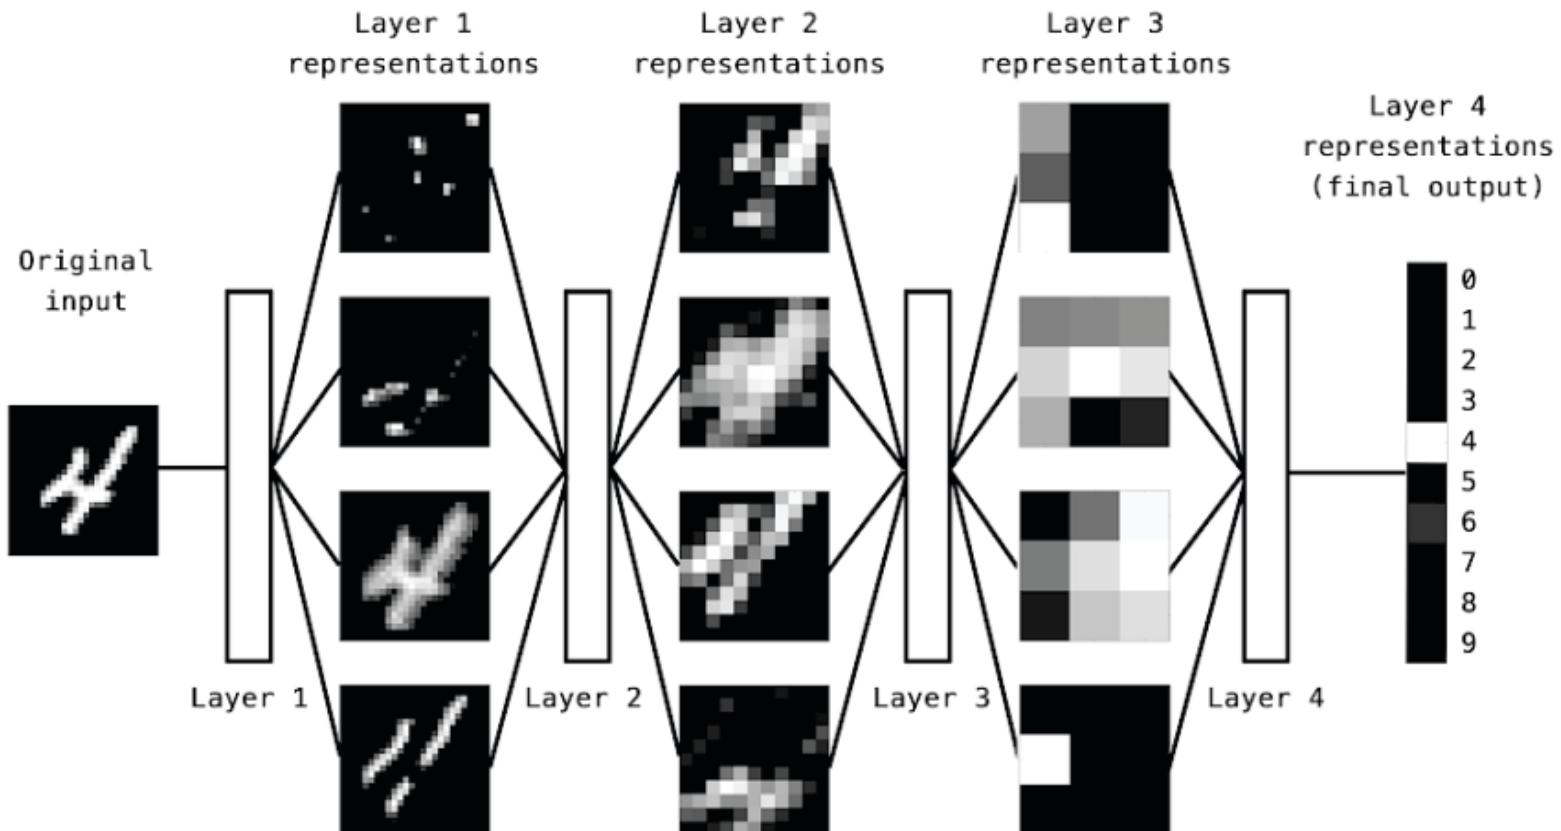


Figure 2. U-net architecture used for canopy palm segmentation, adapted from [19]. The number of channels and the row \times column size in pixels are indicated for each cuboid.

(Chollet & Allaire 2017)

Example with *mnist* dataset



Attention for performance metrics

<https://towardsdatascience.com/metrics-to-evaluate-your-semantic-segmentation-model-6bcb99639aa2>

- Balanced-classification problems
 - Accuracy, Kappa, etc.
- Class-imbalance
 - Precision, Recall, and **F1-score**
- Multi-class segmentation
 - Average precision between classes
- Single-class segmentation
 - **Intersection over Union (IoU)**

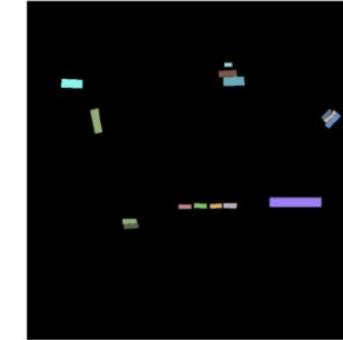
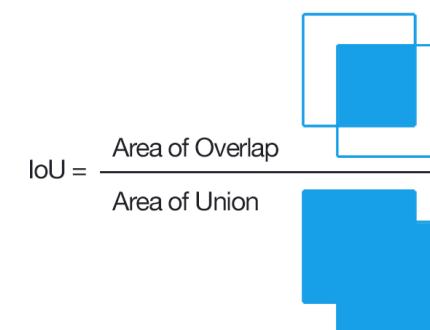
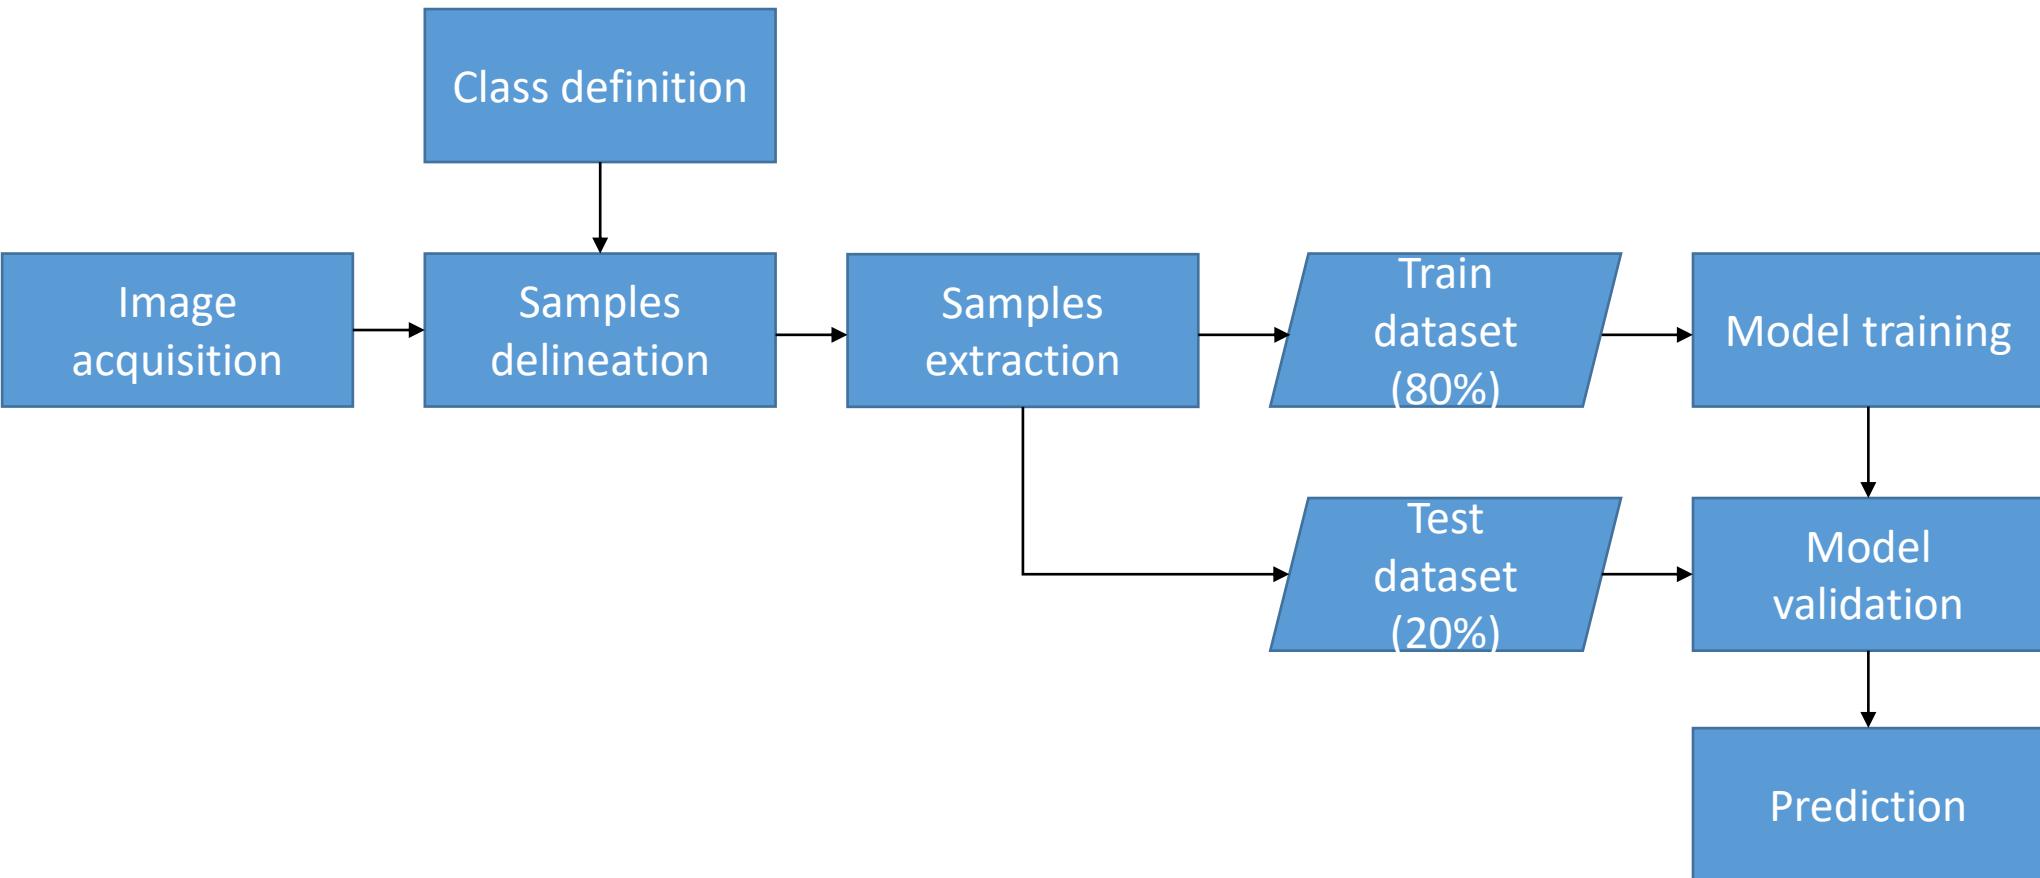


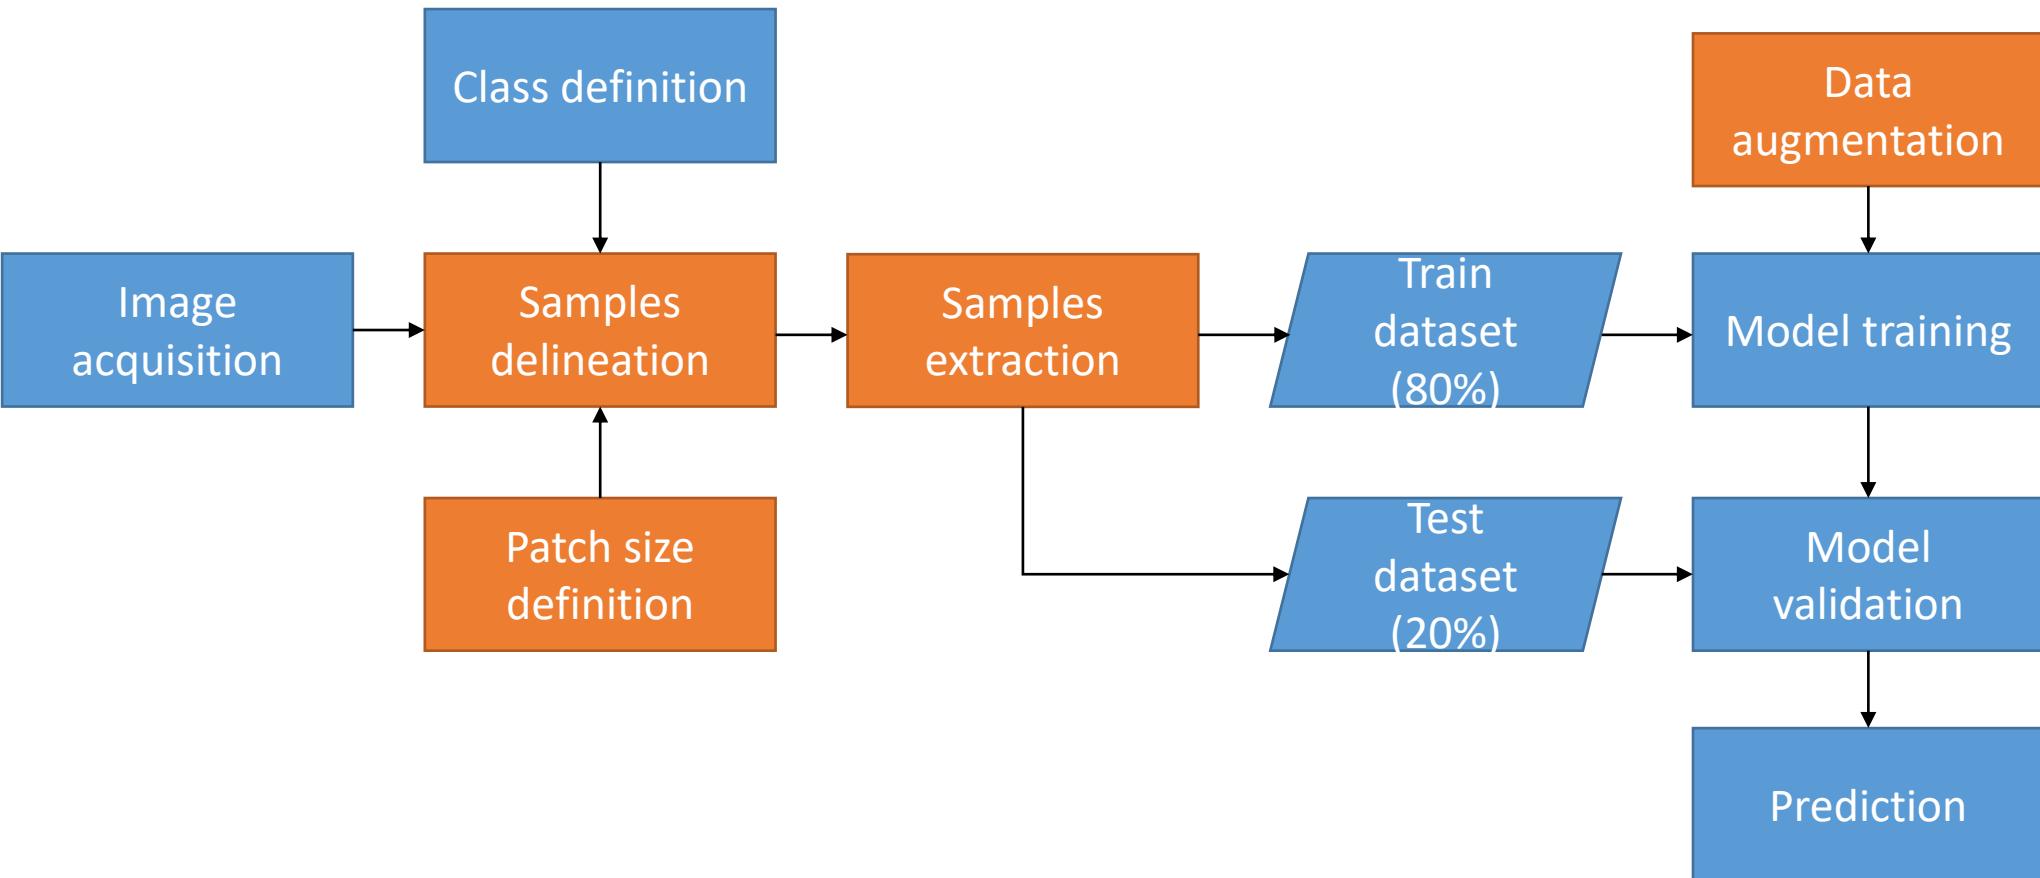
Image from [Vlad Shmyhlo](#) in article: Image Segmentation: Kaggle experience (Part 1 of 2) in TDS



Main steps for RS image segmentation

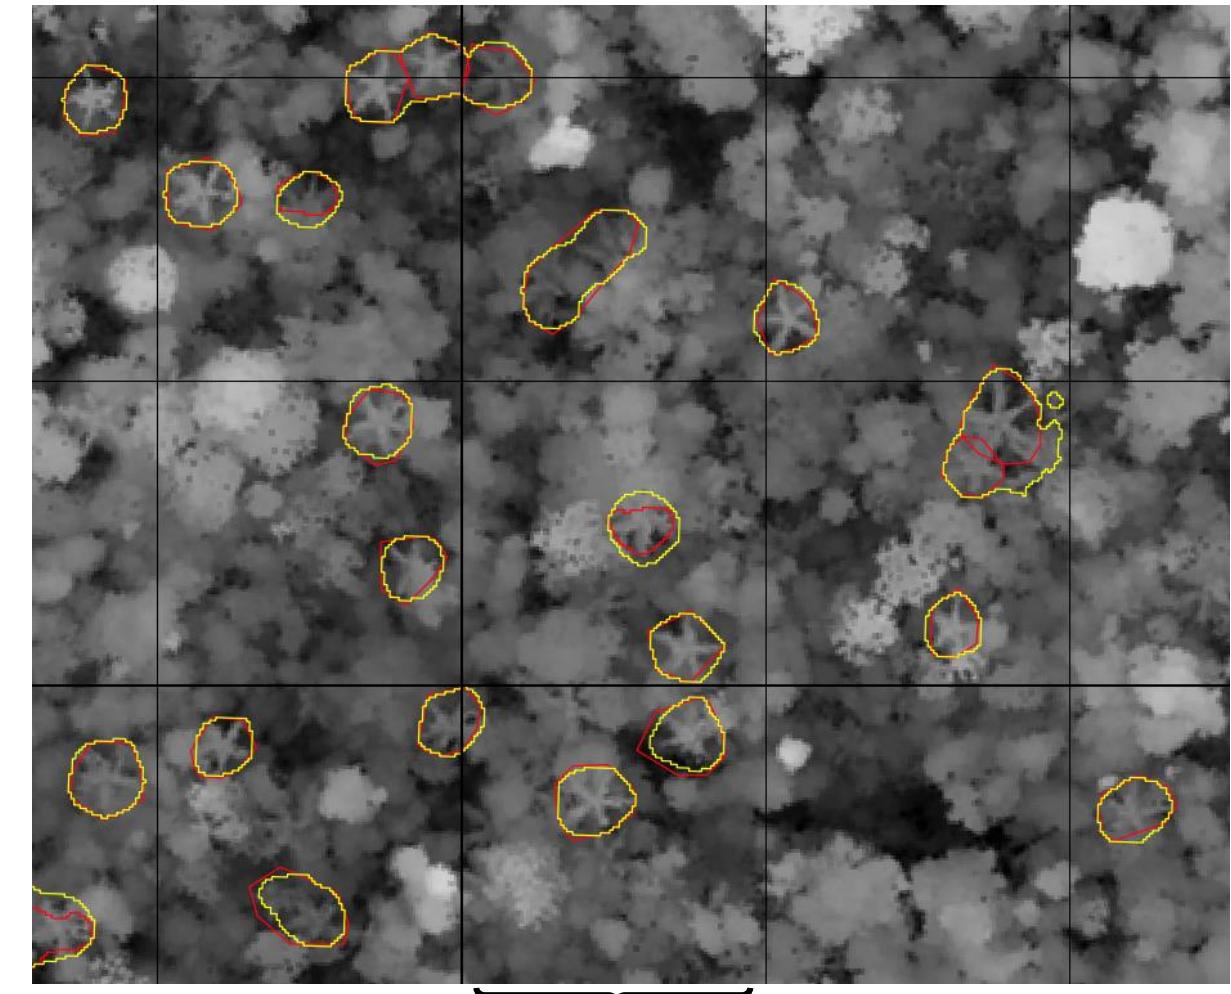


Main steps for RS image segmentation with DL



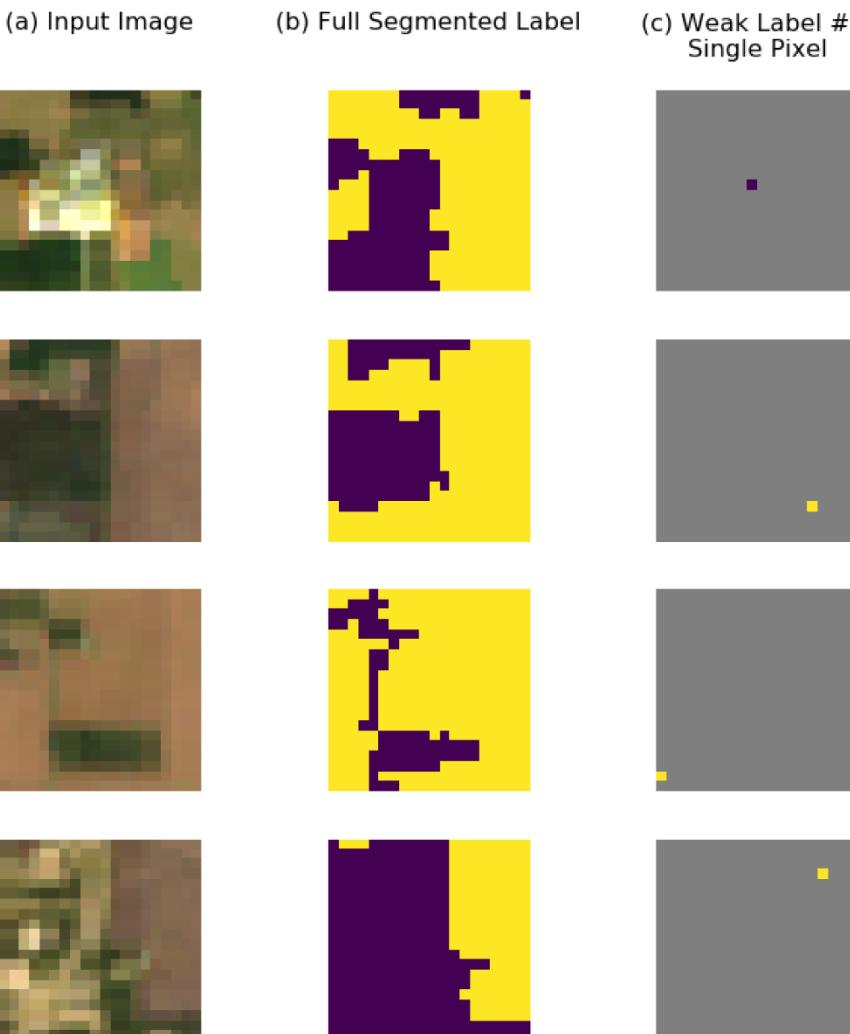
DL vs ML differences: Patch size definition

- Patch of size $m \times m$ pixels
 - $m = \text{multiple of } 32$
 - E.g. 128×128 pixels
- No optimal definition
- The patch should bring contextual information of the target



64 meters
(128x128 pixels of 0.5m)

DL vs ML differences: Samples delineation



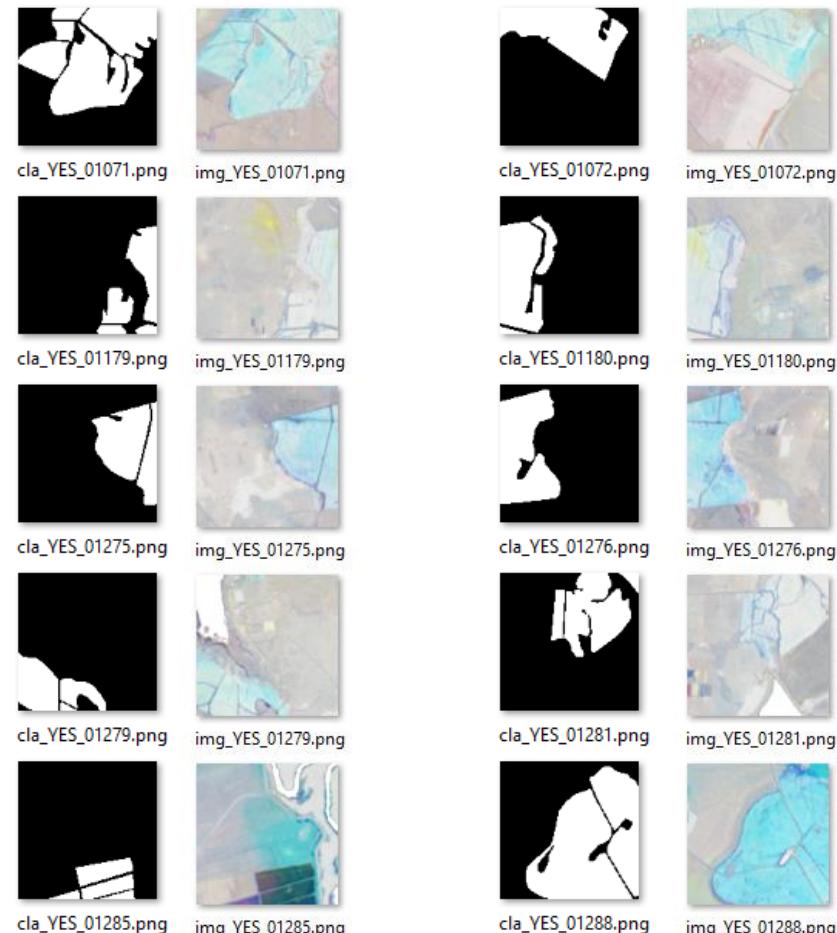
- Samples are usually ‘dense’
- Delineated in GIS or other labelling softwares
- Recent research shows ways to use single pixels called ‘weak supervised learning’ (Wang et al., 2020)

DL vs ML differences: Samples extraction

Traditional ML: pixels to table

	X1	X2	X3	X4	class
2129	0.18409561	0.35243168	0.8963459	0.9284320	Rice
2130	0.18615542	0.48938116	0.9341421	0.9476923	Rice
2131	0.24692002	0.25649321	0.5984405	0.8030888	Rice
2132	0.25300574	0.79422522	0.9033156	0.8696288	Rice
2133	0.30483481	0.69084626	0.9067416	0.9049394	Rice
2134	0.25143707	0.25066489	0.8754706	0.9301759	Rice
2135	0.31929952	0.73389256	0.9191375	0.9151656	Rice
2136	0.44001359	0.47436264	0.4754879	0.4471128	Other
2137	0.62812239	0.60058695	0.4603175	0.4119138	Other
2138	0.66906476	0.59619188	0.5153217	0.4031008	Other
2139	0.51898736	0.56076455	0.2089286	0.2098260	Other
2140	0.47804552	0.37455633	0.3560375	0.3706468	Other
2141	0.57003969	0.56485832	0.5612745	0.6241940	Other

Deep Learning: patches to files



When predicting, we crop, apply model, then mosaic.

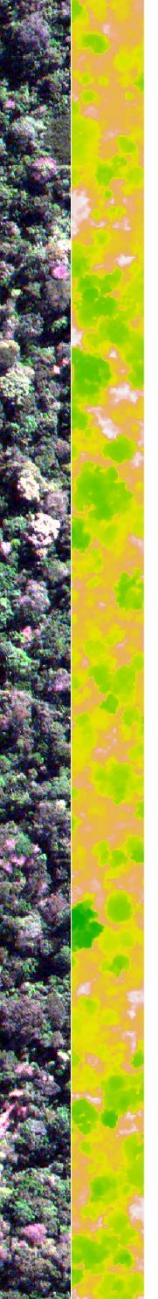
DL vs ML differences: Data augmentation

- Randomized processes that introduces variability in the dataset to teach the model that patterns are not limited to certain angles, values, scales, etc.
 - 0-360° rotations
 - Horizontal and vertical flips
 - Signal change in brightness and saturation (e.g. 90 to 110%)
 - Zoom



(Chollet & Allaire 2017)

What types of RS data can I use? All types!



Optical

- Coarse to high- resolution
 - MODIS
 - Landsat series
 - CBERS series
- Very high resolution (VHR)
 - WorldView-2
 - Planet imagery
 - UAV (Unmanned Aerial Vehicle)
 - Airborne

SAR

- Sentinel-1
- ALOS PALSAR

Lidar

- UAV or Airborne
 - Canopy Height Model (CHM)
 - Digital Terrain Model (DTM)
 - Point cloud

Individually or
combined...
Deep Learning
accepts it all 😊

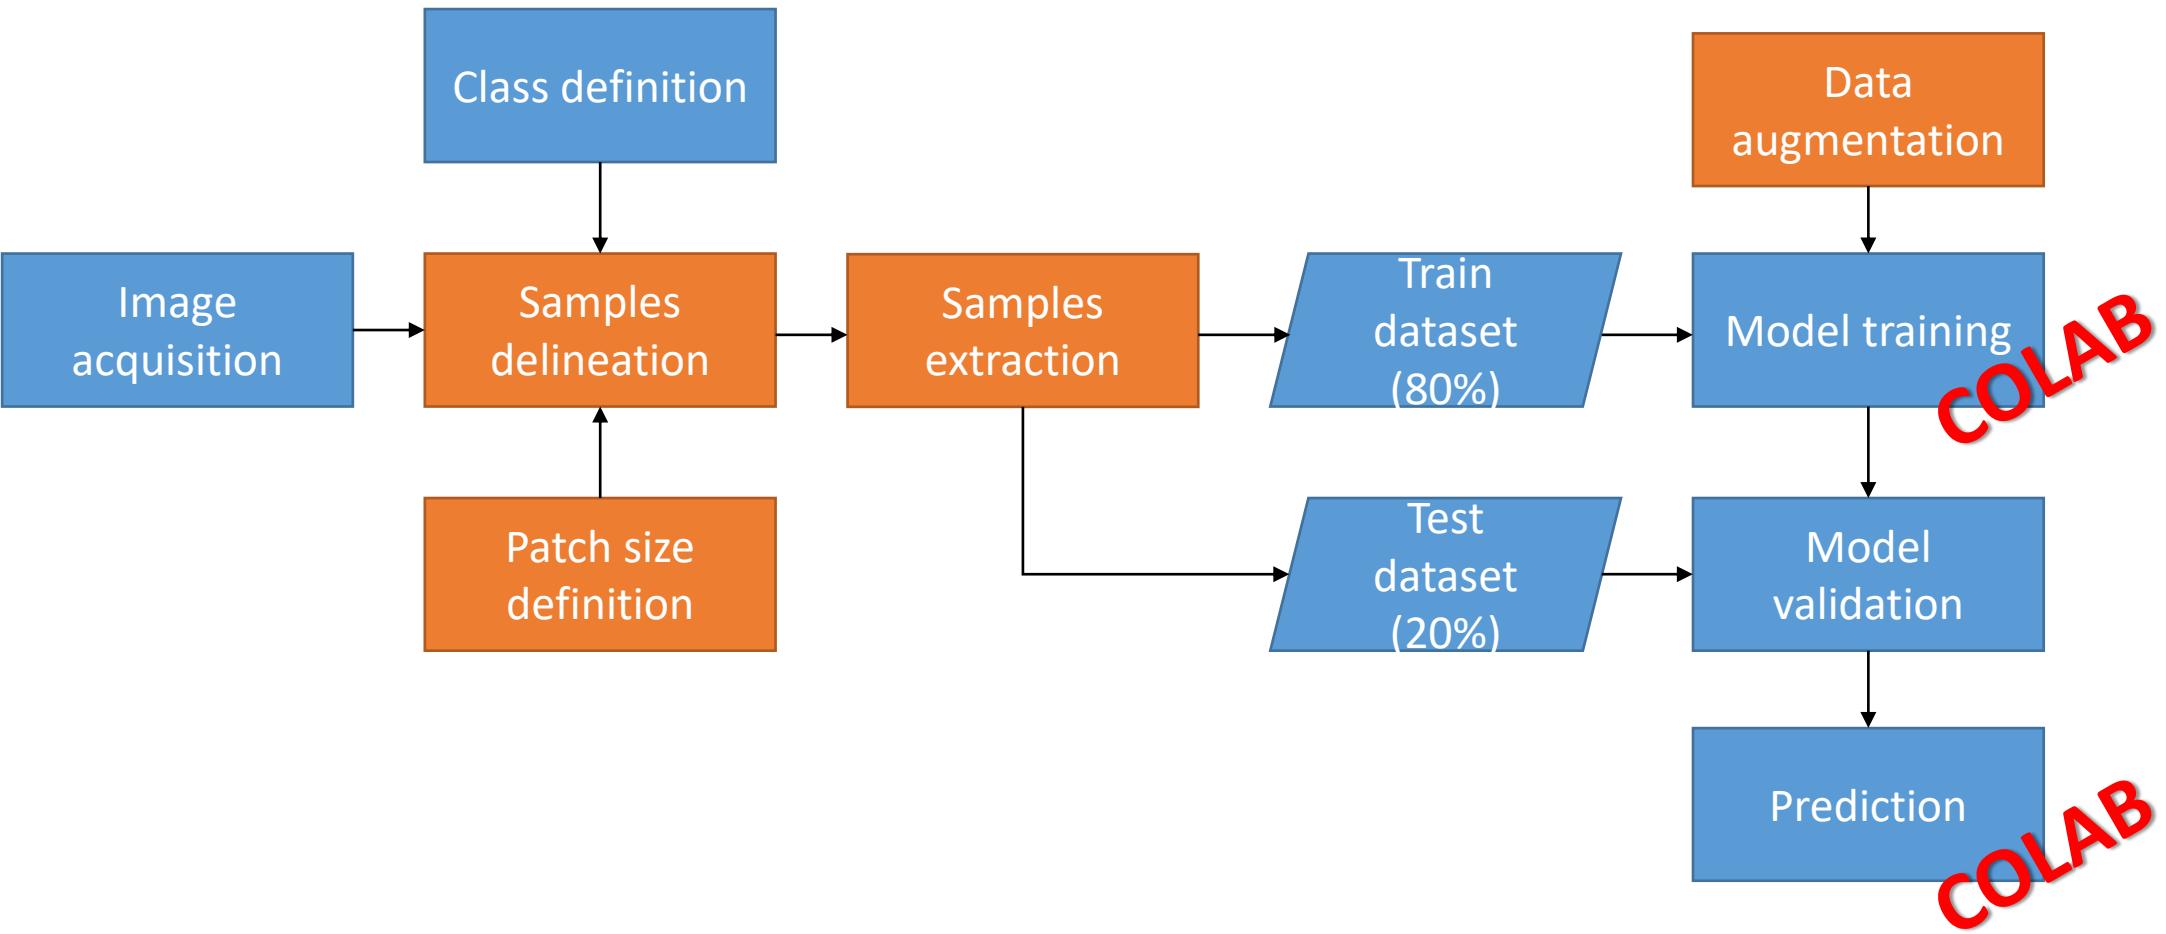
CPU vs GPU



Which language and platform to run DL?

- Python: tensorflow, keras, pytorch
- R: with reticulate / keras <- tensorflow
- Your own computer with GPU
- Google Colab: free, but limited – ideal for beginners
- Paid-virtual machines with GPU (from Amazon, Google, etc.)

Main steps for RS image segmentation with DL





zenodo

Search 🔍

Upload

Communities

 ricds@hotmail.com ▼

July 3, 2020

Lesson Open Access

Edit

New version

1,374

 views

436

 downloads

[See more details...](#)

Indexed in

OpenAIRE

Download the PDF file and follow the instructions. This is for NVIDIA GPU and Windows OS.

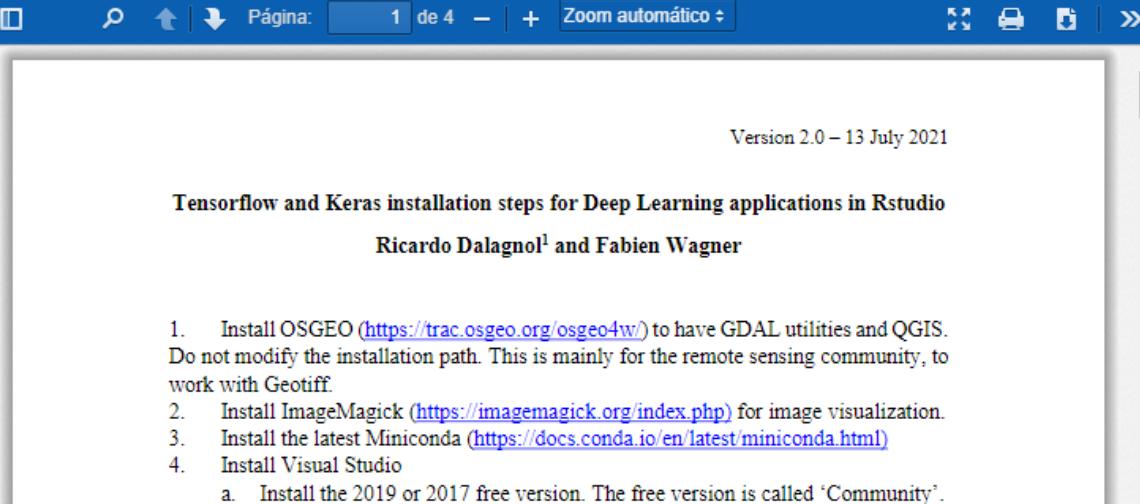
Contacts : Ricardo Dalagnol (ricds@hotmail.com)

Fabien Wagner (wagner.h.fabien@gmail.com)

v2.0 Updated for Python 3.9 and Tensorflow 2.5.

Send an email for collaboration or if you have suggestions for the improvement of this manual.

Preview



Página: 1 de 4

Zoom automático

Version 2.0 – 13 July 2021

Tensorflow and Keras installation steps for Deep Learning applications in Rstudio

Ricardo Dalagnol¹ and Fabien Wagner

1. Install OSGEO (<https://trac.osgeo.org/osgeo4w>) to have GDAL utilities and QGIS. Do not modify the installation path. This is mainly for the remote sensing community, to work with Geotiff.
2. Install ImageMagick (<https://imagemagick.org/index.php>) for image visualization.
3. Install the latest Miniconda (<https://docs.conda.io/en/latest/miniconda.html>)
4. Install Visual Studio
 - a. Install the 2019 or 2017 free version. The free version is called ‘Community’.

Publication date:

July 3, 2020

DOI:

 DOI [10.5281/zenodo.5095433](https://doi.org/10.5281/zenodo.5095433)

Keyword(s):

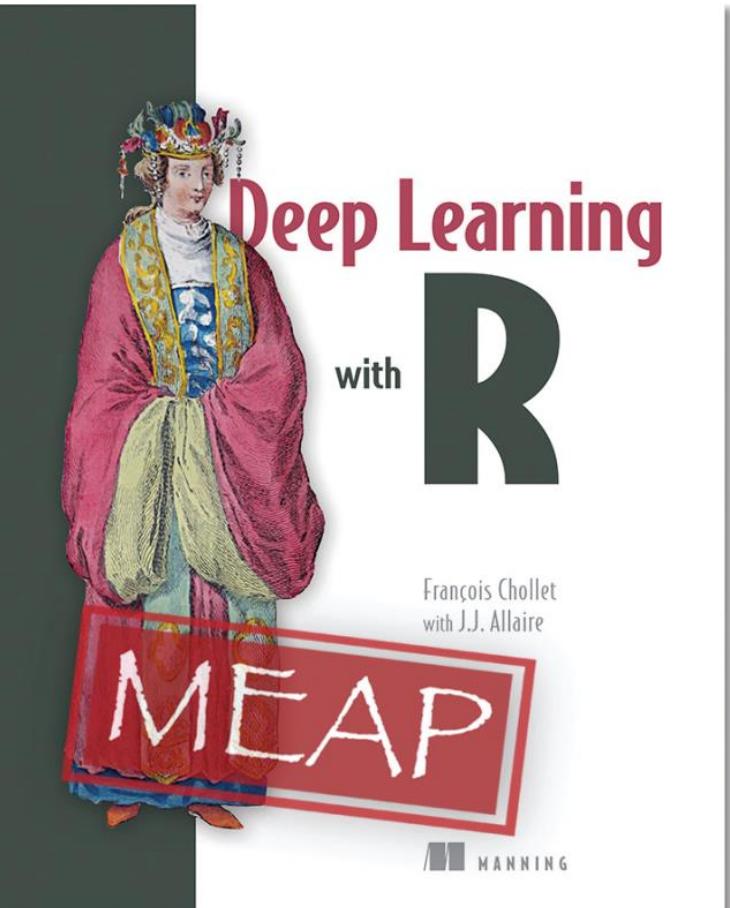
License (for files):

 Creative Commons Attribution 4.0 International

Versions

References

“DL Bible”



(Chollet & Allaire 2017)

Overview

Keras

Getting started

Overview

Keras basics

Sequential API

Functional API

Saving and serializing models

Frequently Asked Questions

Advanced

Getting Started with Keras

Overview

Keras is a high-level neural networks API developed with a focus on enabling fast experimentation. *Being able to go from idea to result with the least possible delay is key to doing good research.* Keras has the following key features:

- Allows the same code to run on CPU or on GPU, seamlessly.
- User-friendly API which makes it easy to quickly prototype deep learning models.
- Built-in support for convolutional networks (for computer vision), recurrent networks (for sequence processing), and any combination of both.
- Supports arbitrary network architectures: multi-input or multi-output models, layer sharing, model sharing, etc. This means that Keras is appropriate for building essentially any deep learning model, from a memory network to a neural Turing machine.

This website provides documentation for the R interface to Keras. See the main Keras website at <https://keras.io> for additional information on the project.

Object Detection and Image Segmentation with Deep Learning on Earth Observation Data: A Review-Part I: Evolution and Recent Trends

by Thorsten Hoeser ^{1,*} and Claudia Kuenzer ^{1,2}

¹ German Remote Sensing Data Center (DFD), German Aerospace Center (DLR), Münchner Straße 20, D-82234 Wessling, Germany

² Department of Remote Sensing, Institute of Geography and Geology, University Würzburg, Am Huband, D-97074 Wuerzburg, Germany

* Author to whom correspondence should be addressed.

Remote Sens. **2020**, *12*(10), 1667; <https://doi.org/10.3390/rs12101667>

Received: 27 April 2020 / Revised: 19 May 2020 / Accepted: 19 May 2020 / Published: 22 May 2020

Hoeser & Kuenzer (2020) *Remote sensing*

<https://www.mdpi.com/2072-4292/12/10/1667/htm>

Summary

- Deep learning are state-of-the-art models for image segmentation. You should start to do DL now 😊
- Hope you learned something new today, thanks !

Contact



ricds@hotmail.com



@RicardoDalagnol