



National Institute for Space Research – INPE



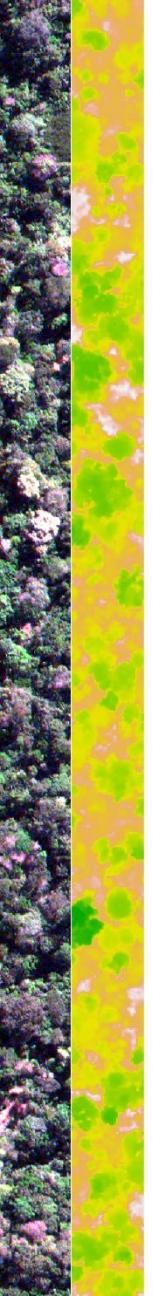
Short course @ IEEE GRSS-ISPRS SC 2021

# Deep Learning applied for Remote Sensing images



Ricardo Dal'Agnol da Silva  
Post-doc Researcher - INPE

04 Nov 2021



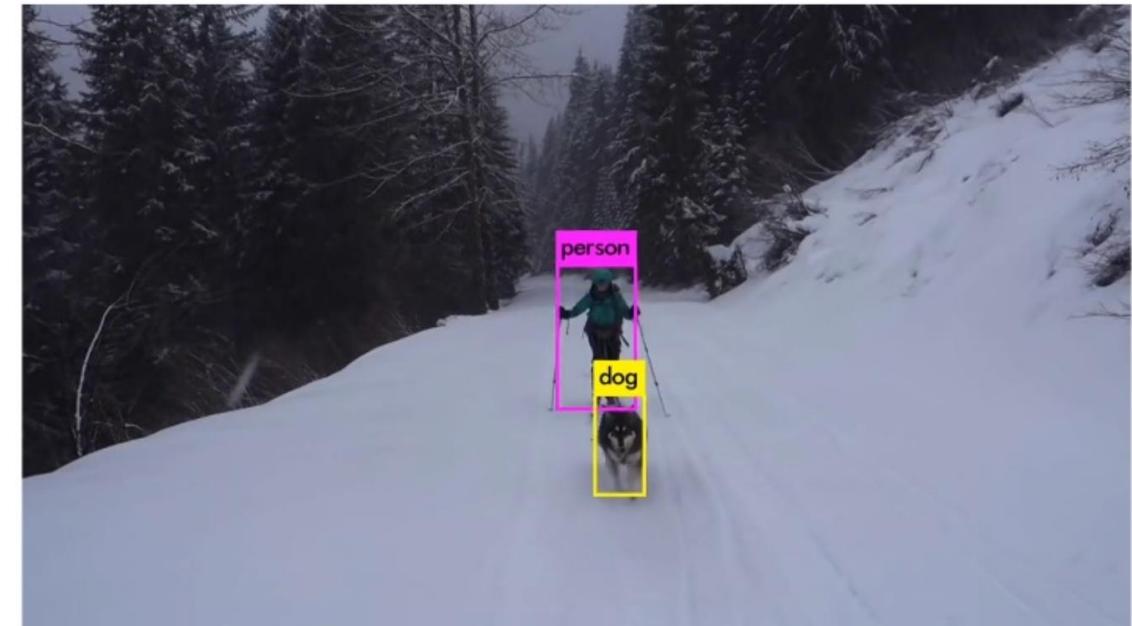
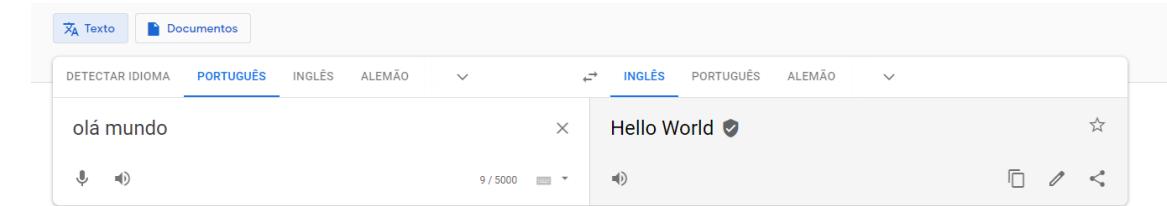
# Learning objectives

- Describe what is deep learning
- Explain the innovation brought by convolutional neural networks (CNN)
- Identify the main steps on applying CNN for remote sensing
- Apply an example of CNN with test dataset

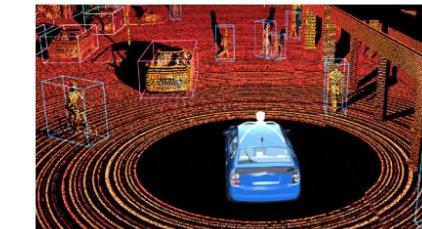
# "AI is the new electricity" - Andrew Ng



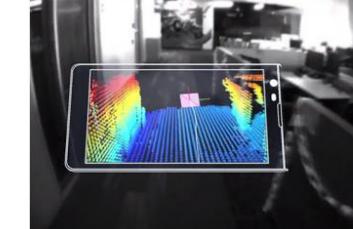
OpenAI Five: Dota 2  
The International 2018



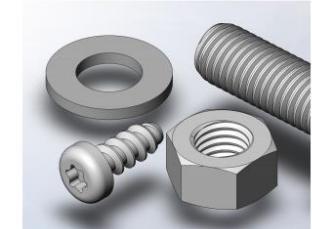
Robot Perception



Augmented Reality



Shape Design



# Classification vs Detection vs Semantic vs Instance segmentation

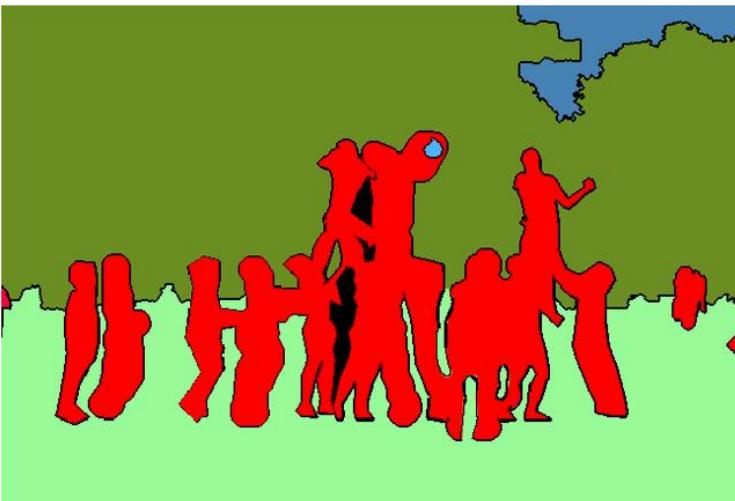
Image Classification



Object Detection

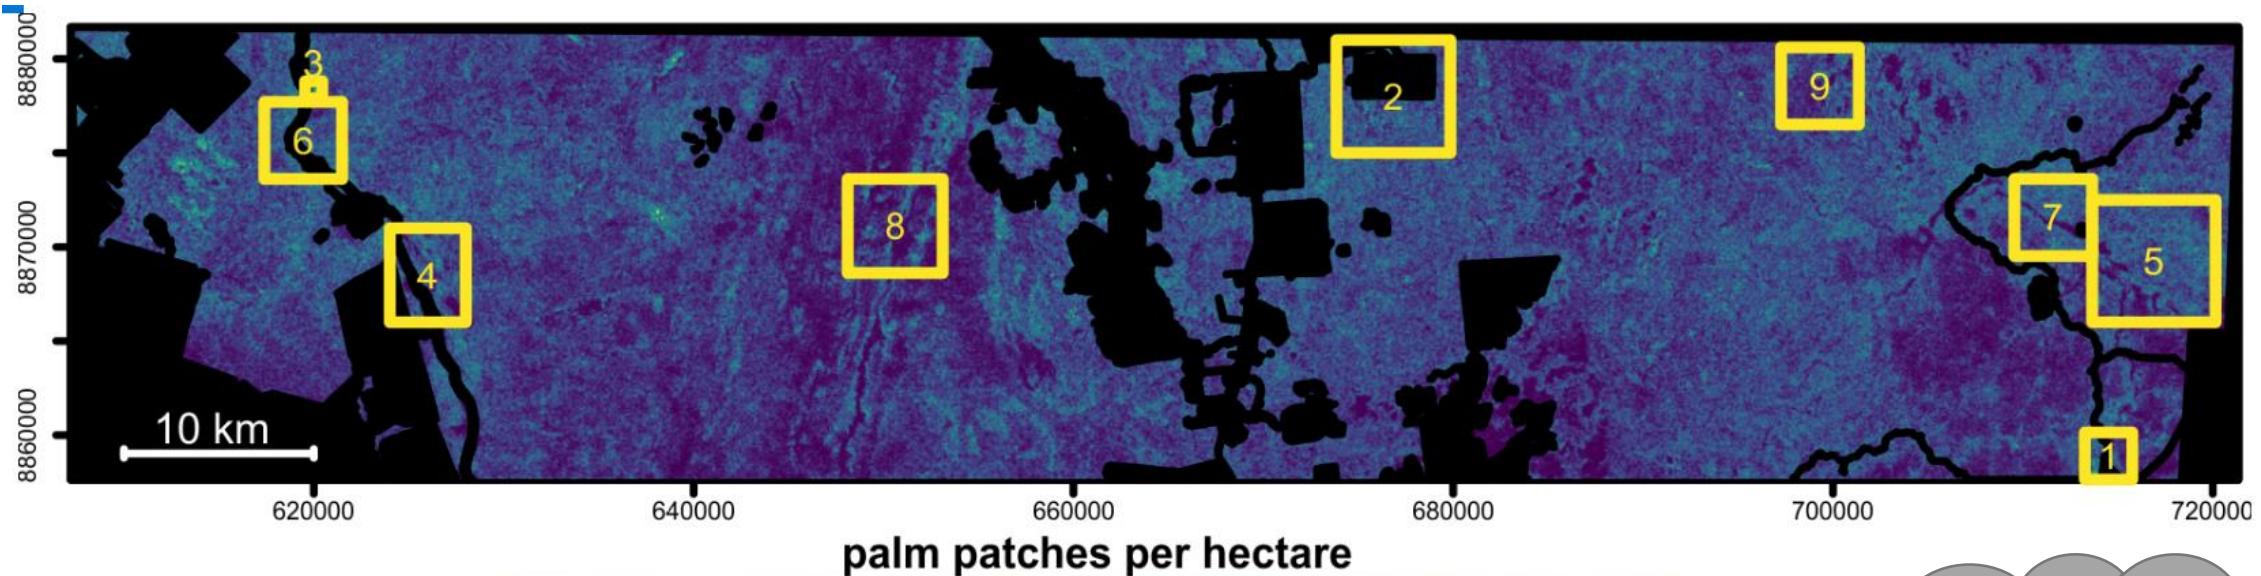


Semantic Segmentation

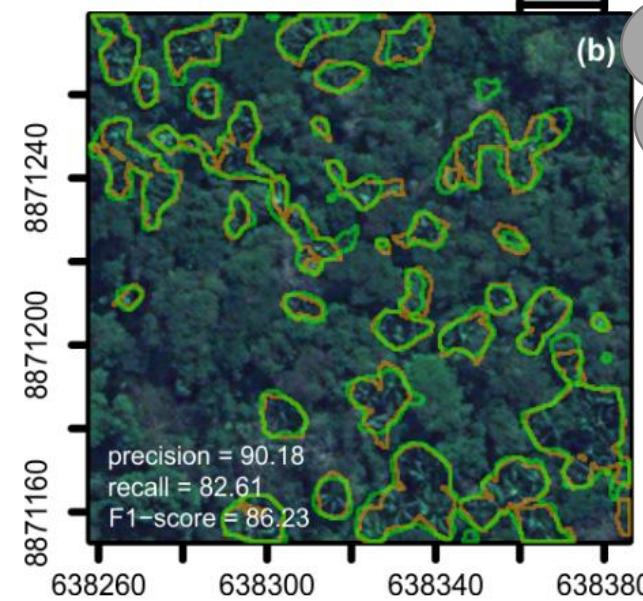
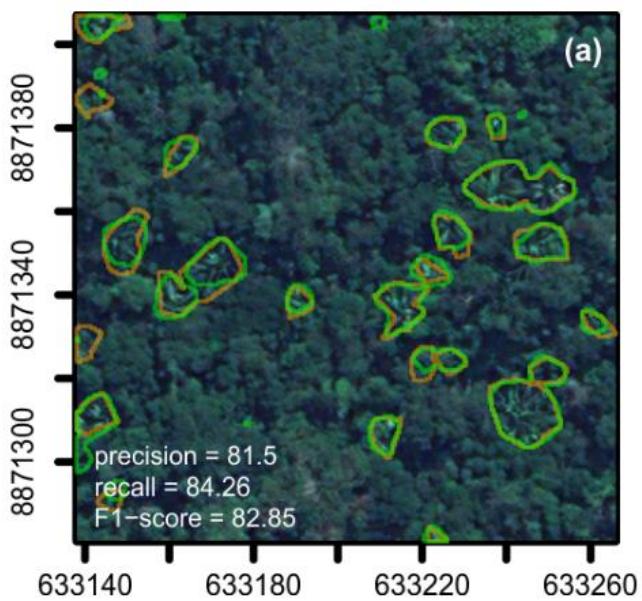


Instance Segmentation





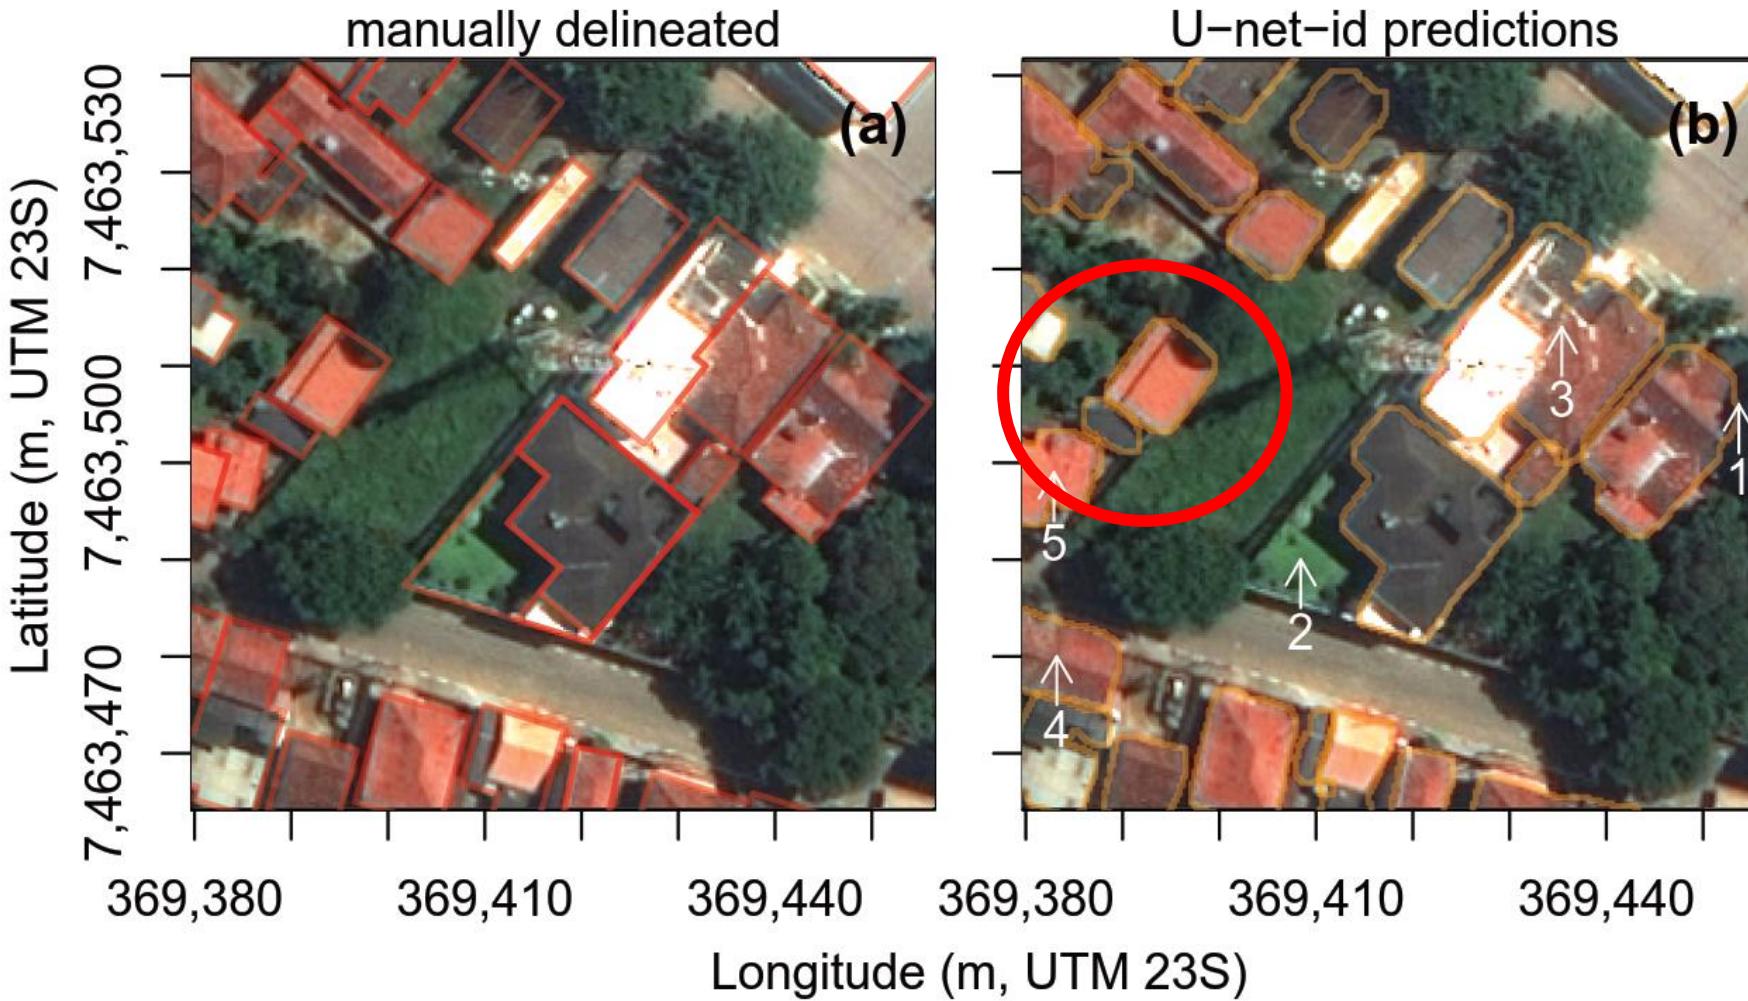
palm patches per hectare



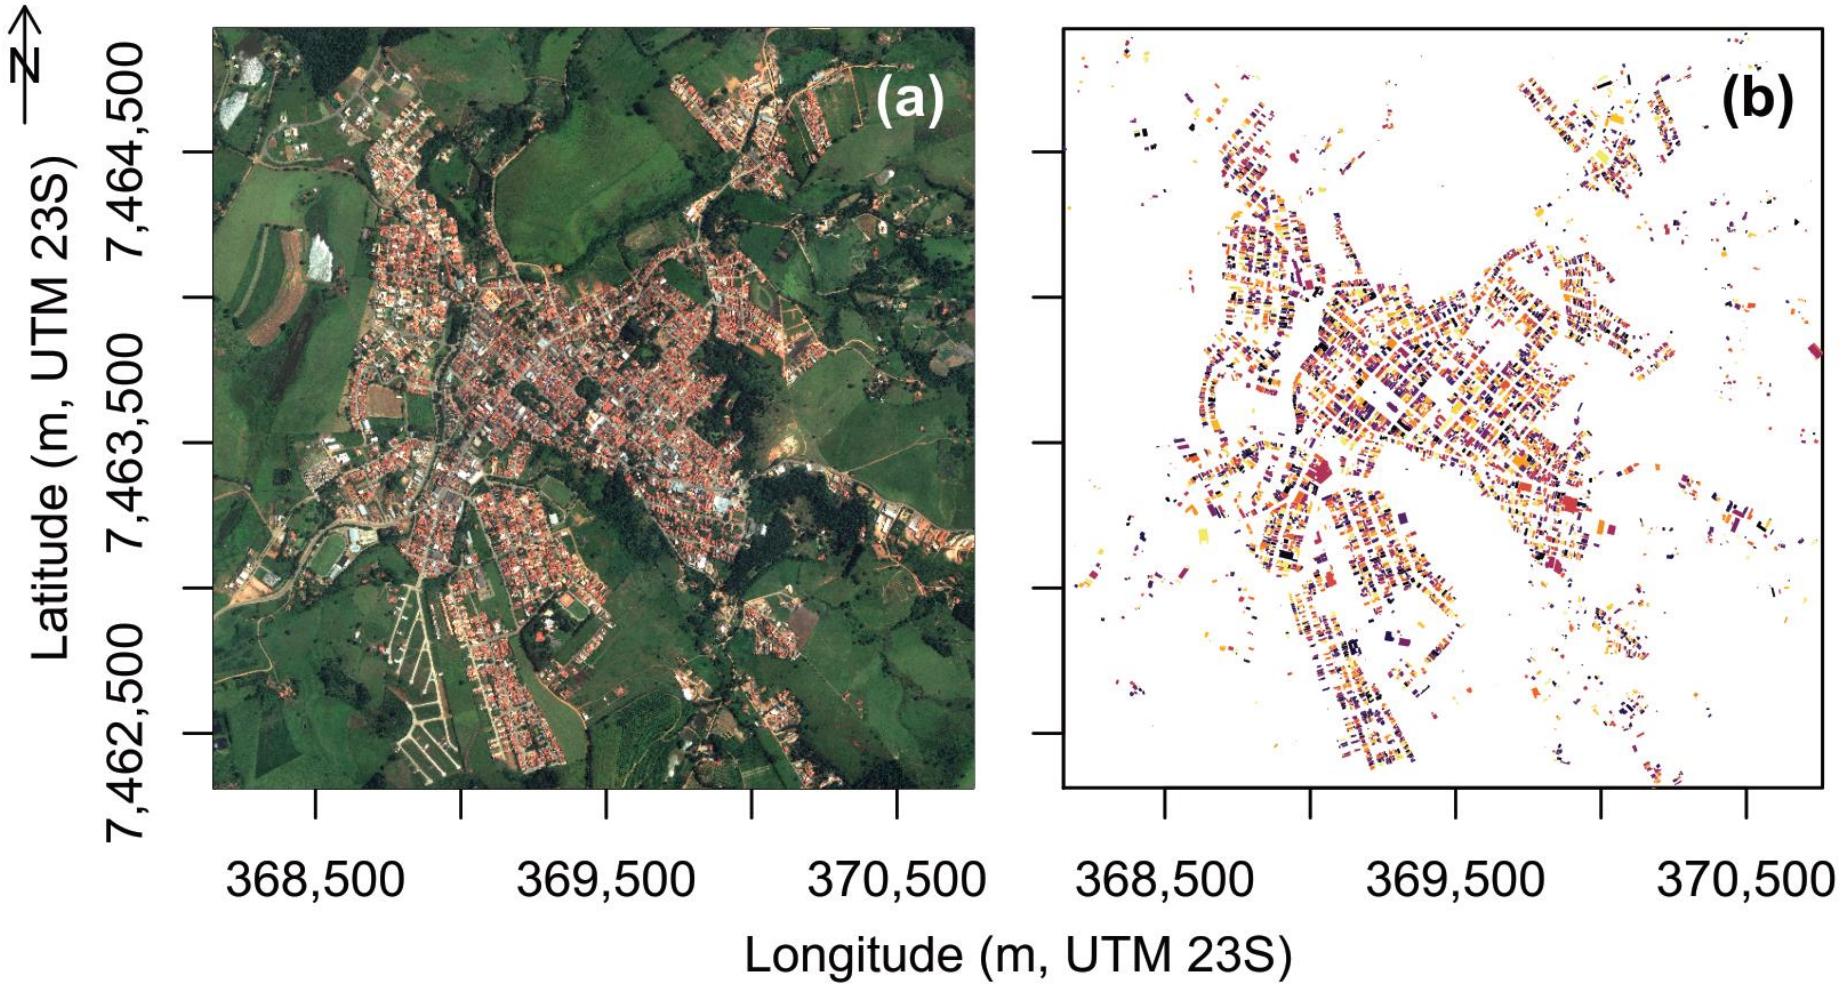
*False Positives are  
very often  
True Positives...*

- manual delineation
- automatic delineation

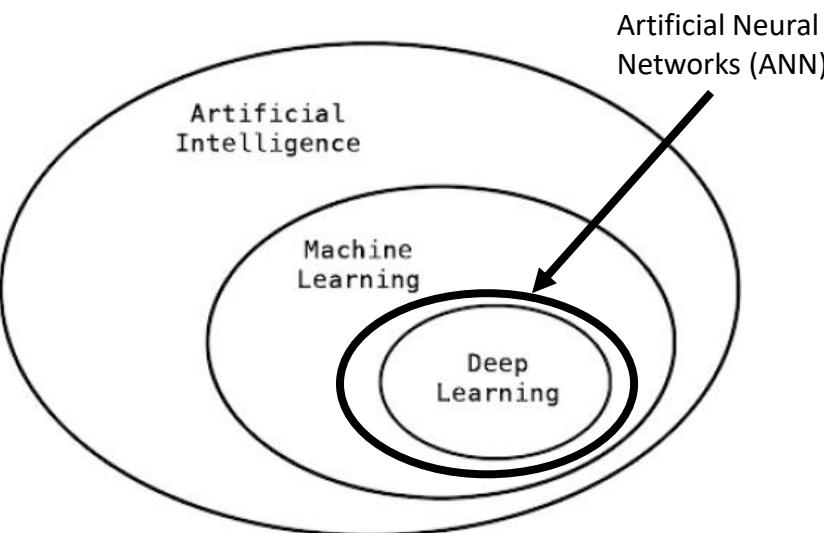
# Building segmentation (Wagner et al. 2020)



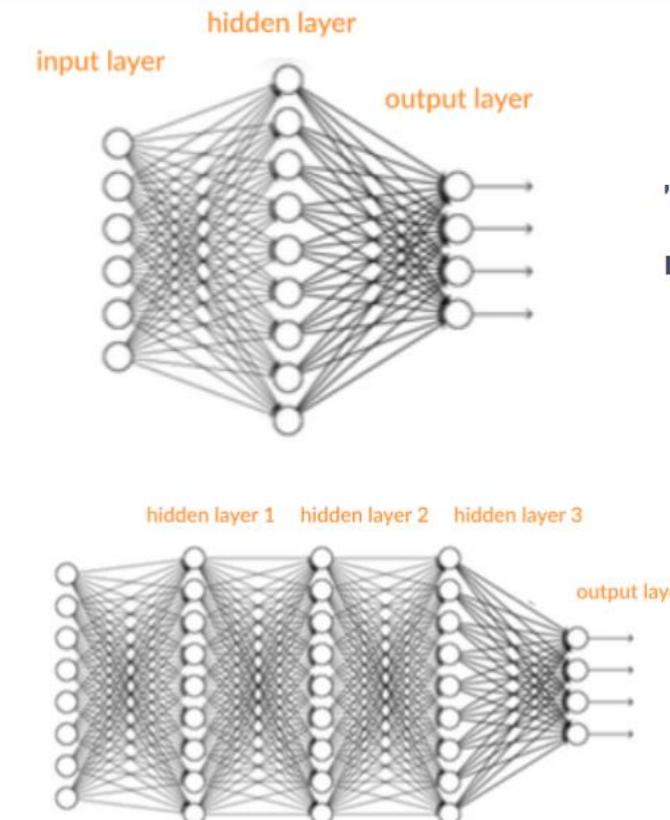
# Building segmentation (Wagner et al. 2020)



AI -> ML -> ANN -> DL



(Adapted from Chollet & Allaire 2017)

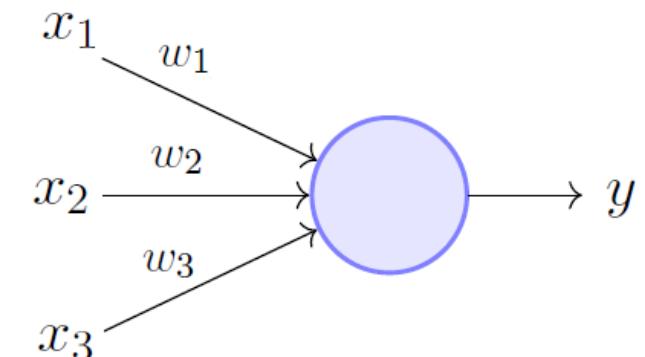


"Non-deep" feedforward neural network

Deep neural network

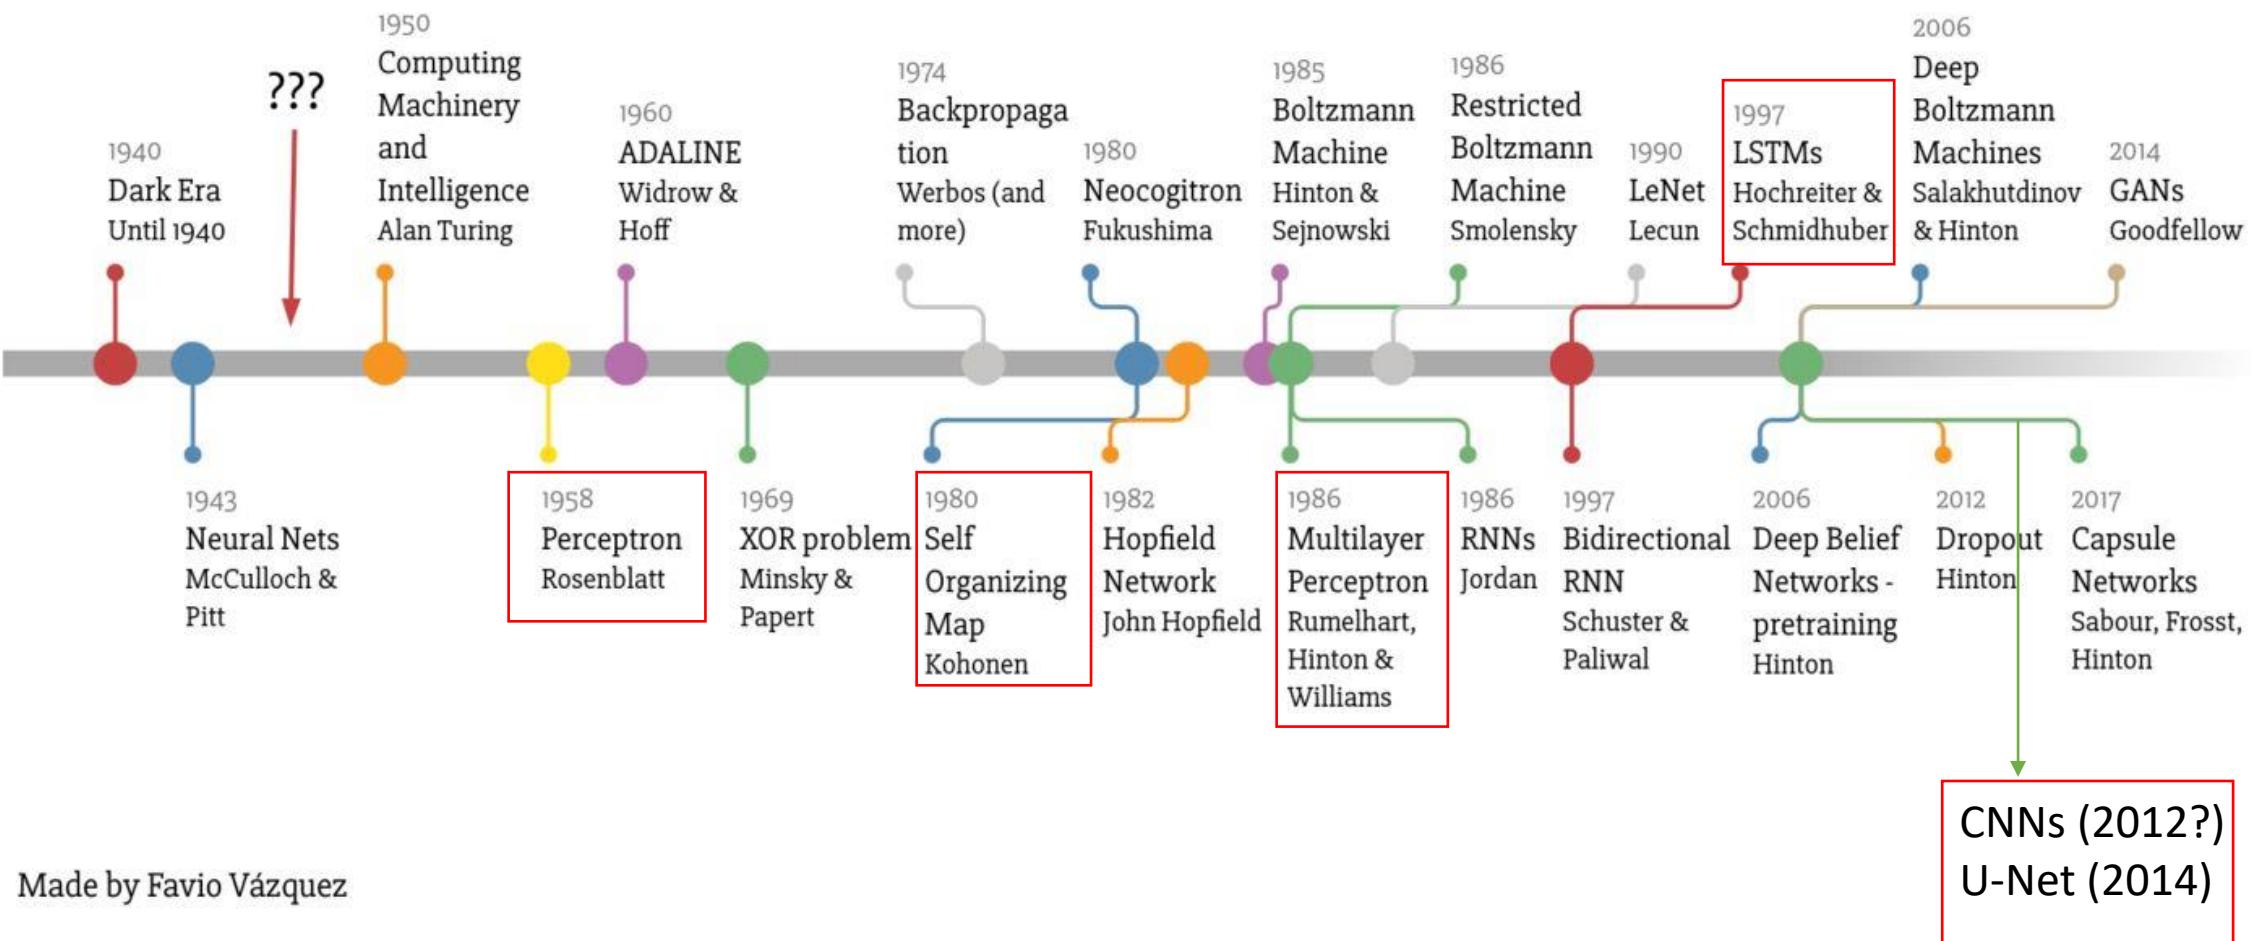
# Neural networks

- Goal: approximate some mathematical function
- Relate inputs ( $x$ ) to outputs ( $y$ ) using weights ( $w$ ) in between neurons
- Can have many layers with neurons
- Applications for regression or classification problems
- Usually supervised but can also be unsupervised



Perceptron Model (Minsky-Papert in 1969)

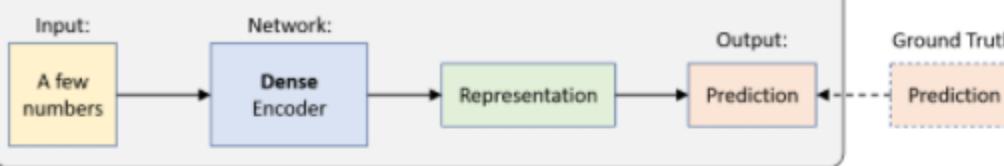
# Deep Learning Timeline



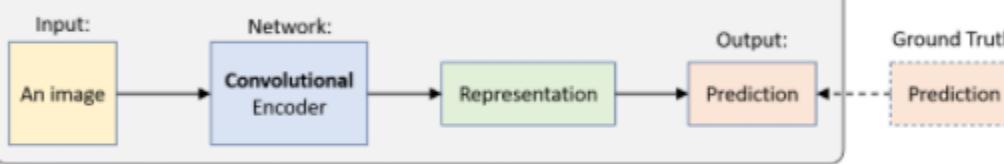
# Main DL model types

## Supervised Learning

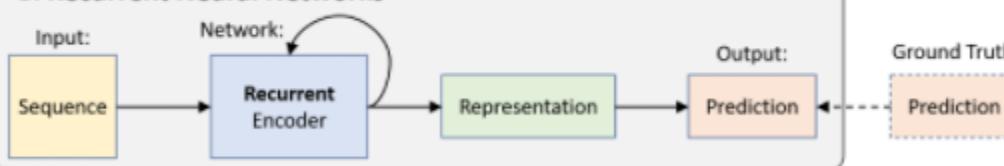
### 1. Feed Forward Neural Networks



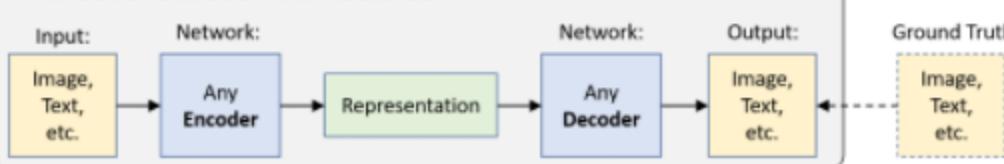
### 2. Convolutional Neural Networks



### 3. Recurrent Neural Networks

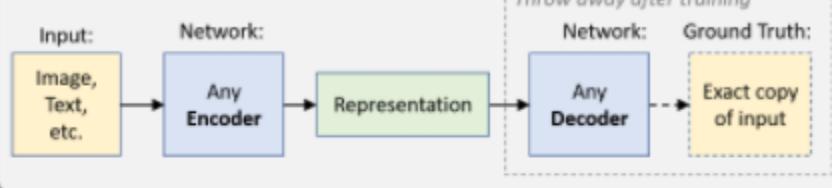


### 4. Encoder-Decoder Architectures

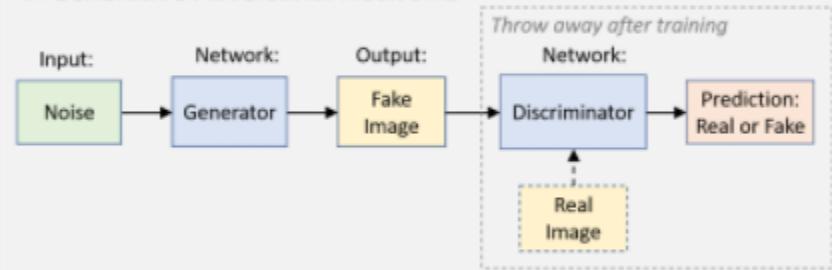


## Unsupervised Learning

### 5. Autoencoder

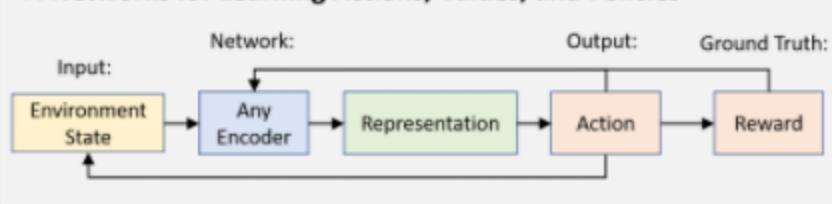


### 6. Generative Adversarial Networks



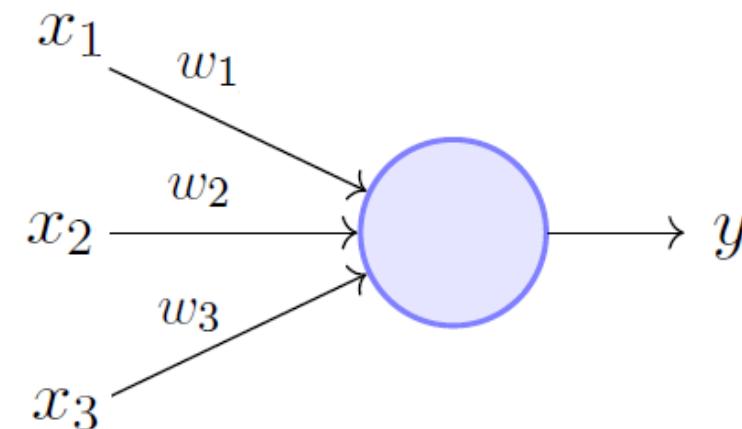
## Reinforcement Learning

### 7. Networks for Learning Actions, Values, and Policies



# Quick intro into ANN - Perceptron 1958

- Basically a linear model
- *feed forward* type
- Weights ( $w$ ) are adjusted and multiplied by inputs ( $x$ ) to obtain the output ( $y$ )

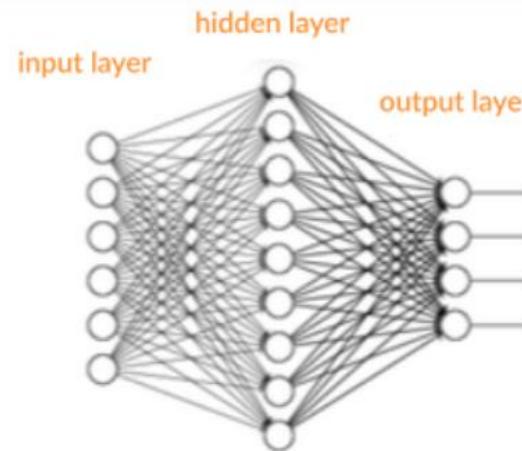


*Simplifying to:*  
 $y = x_1 * w_1 + x_2 * w_2 + x_3 * w_3$

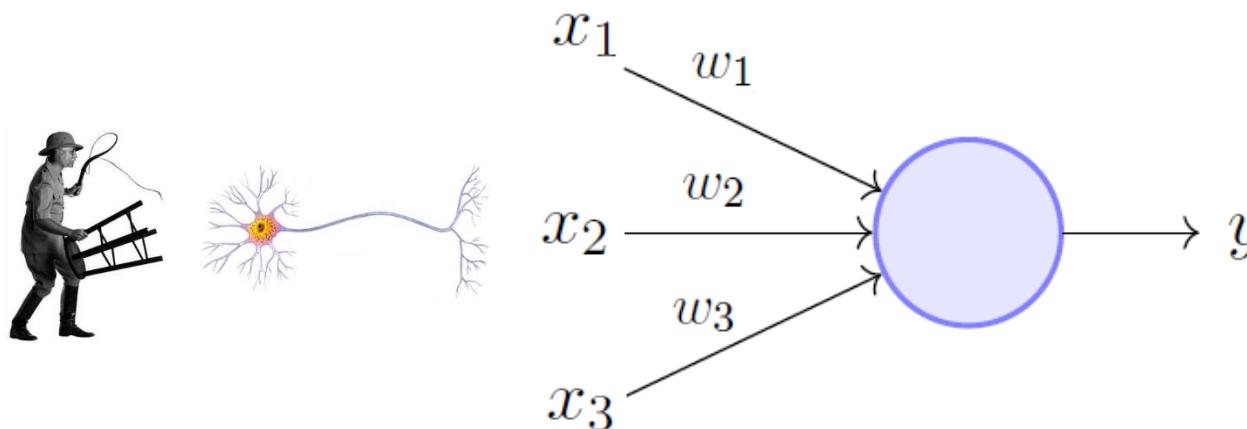
Perceptron Model (Minsky-Papert in 1969)

# Perceptron 1958

- Basically a linear model
- *feed forward* type
- Weights ( $w$ ) are adjusted and multiplied by inputs ( $x$ ) to obtain the output ( $y$ )



Multi-Layer Perceptron (MLP)  
1986

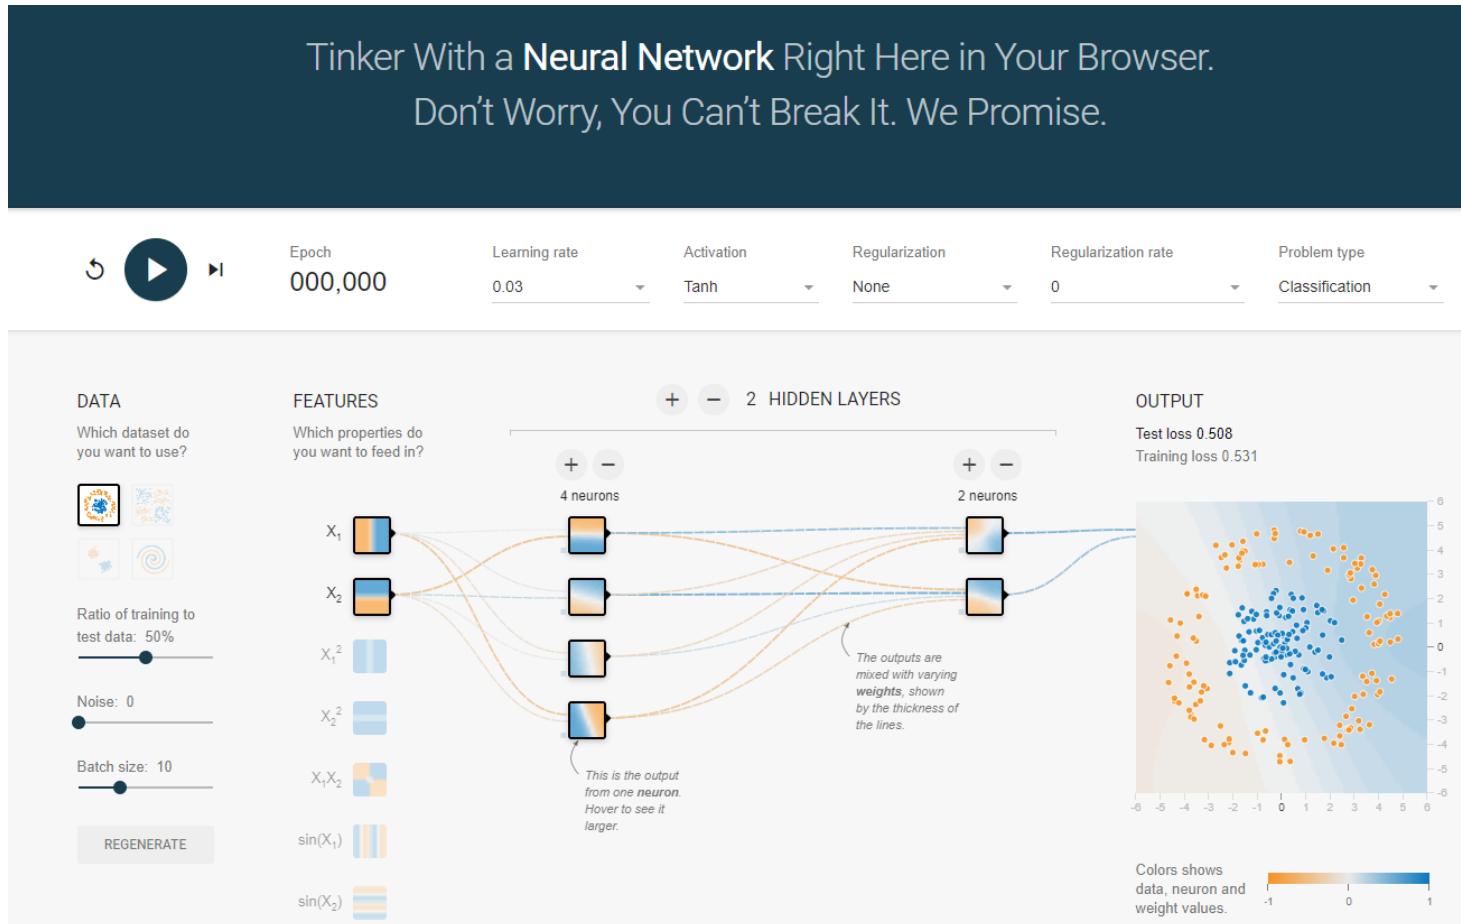


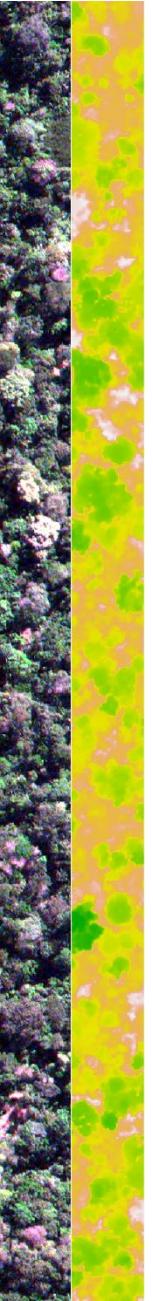
*Simplifying to:*  
 $y = x_1 * w_1 + x_2 * w_2 + x_3 * w_3$

Perceptron Model (Minsky-Papert in 1969)

# Playing with a MLP

- <https://playground.tensorflow.org/>





# Why deep learning for remote sensing?

- Classification and segmentation, but also regression

>	
<b>Method</b>	
<b>Number of layers</b>	
<b>Accuracy</b>	
<b>Domain</b>	

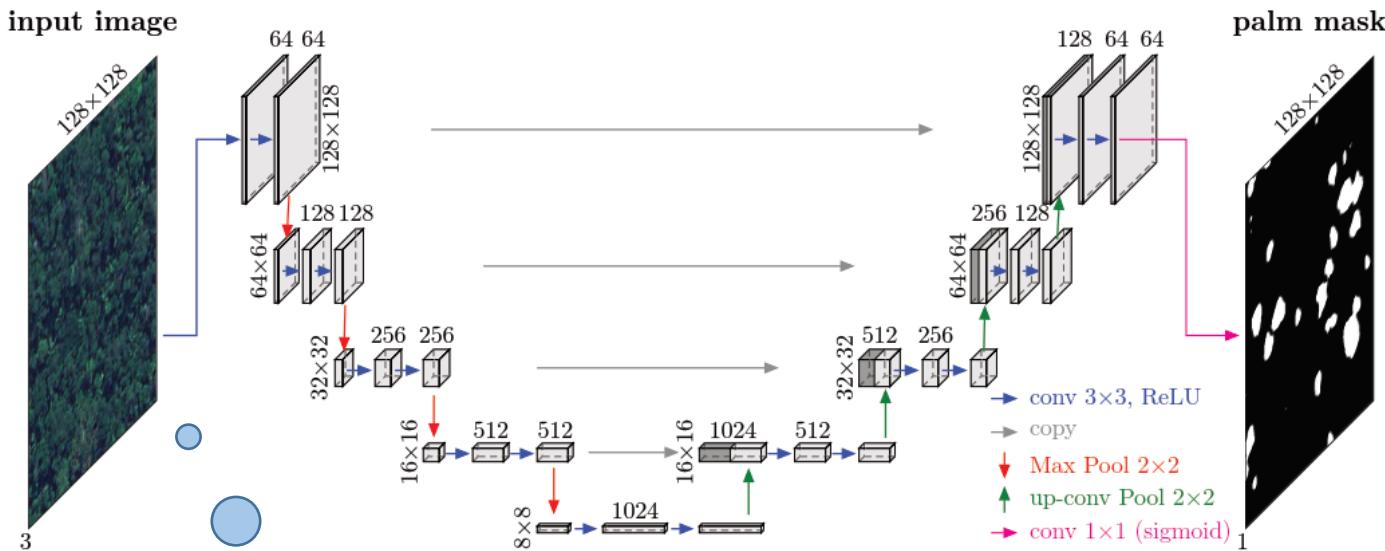
<b>Past</b>
• Multi-layer Perceptron (MLP)
• A few, e.g. usually < 5
• Similar accuracy to traditional machine learning algorithms
• Pixel-level (or region)

# Why deep learning for remote sensing?

- Classification and segmentation, but also regression

>	Past	Present
<b>Method</b>	<ul style="list-style-type: none"><li>• Multi-layer Perceptron (MLP)</li></ul>	<ul style="list-style-type: none"><li>• Convolutional Neural Networks (CNN)</li></ul>
<b>Number of layers</b>	<ul style="list-style-type: none"><li>• A few, e.g. usually &lt; 5</li></ul>	<ul style="list-style-type: none"><li>• Many and growing, e.g. &gt;20</li></ul>
<b>Accuracy</b>	<ul style="list-style-type: none"><li>• Similar accuracy to traditional machine learning algorithms</li></ul>	<ul style="list-style-type: none"><li>• <u>Above traditional machine learning methods and close to human level</u></li></ul>
<b>Domain</b>	<ul style="list-style-type: none"><li>• Pixel-level (or region)</li></ul>	<ul style="list-style-type: none"><li>• Patch-level</li></ul>

# Convolutional neural networks: state-of-the-art methods for semantic segmentation



**Figure 2.** U-net architecture used for canopy palm segmentation, adapted from [19]. The number of channels and the row  $\times$  column size in pixels are indicated for each cuboid.

The inputs are  
image patches  
instead of  
pixels



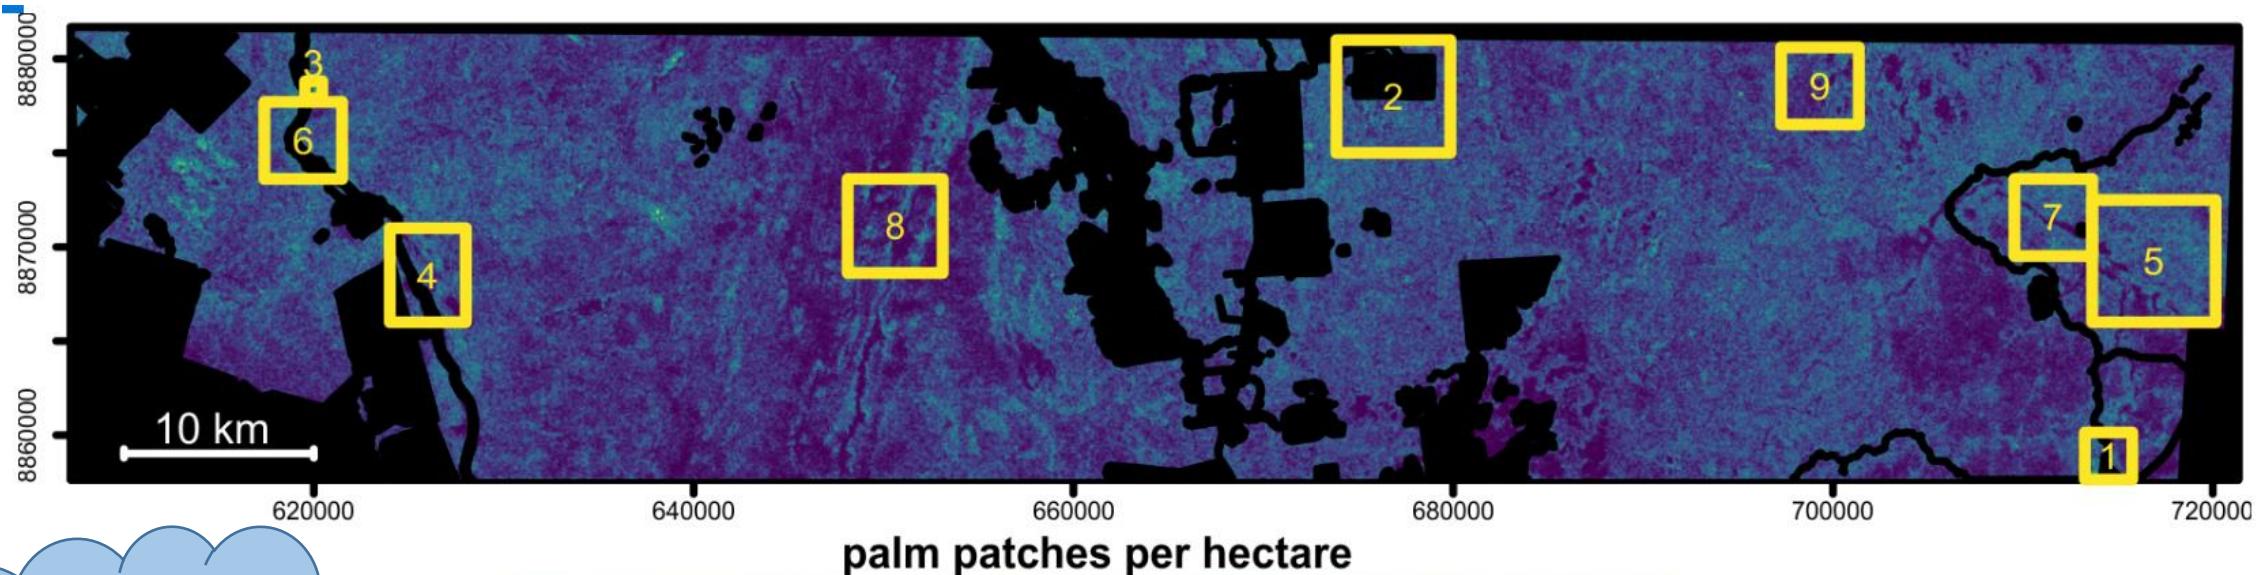
remote sensing

Article

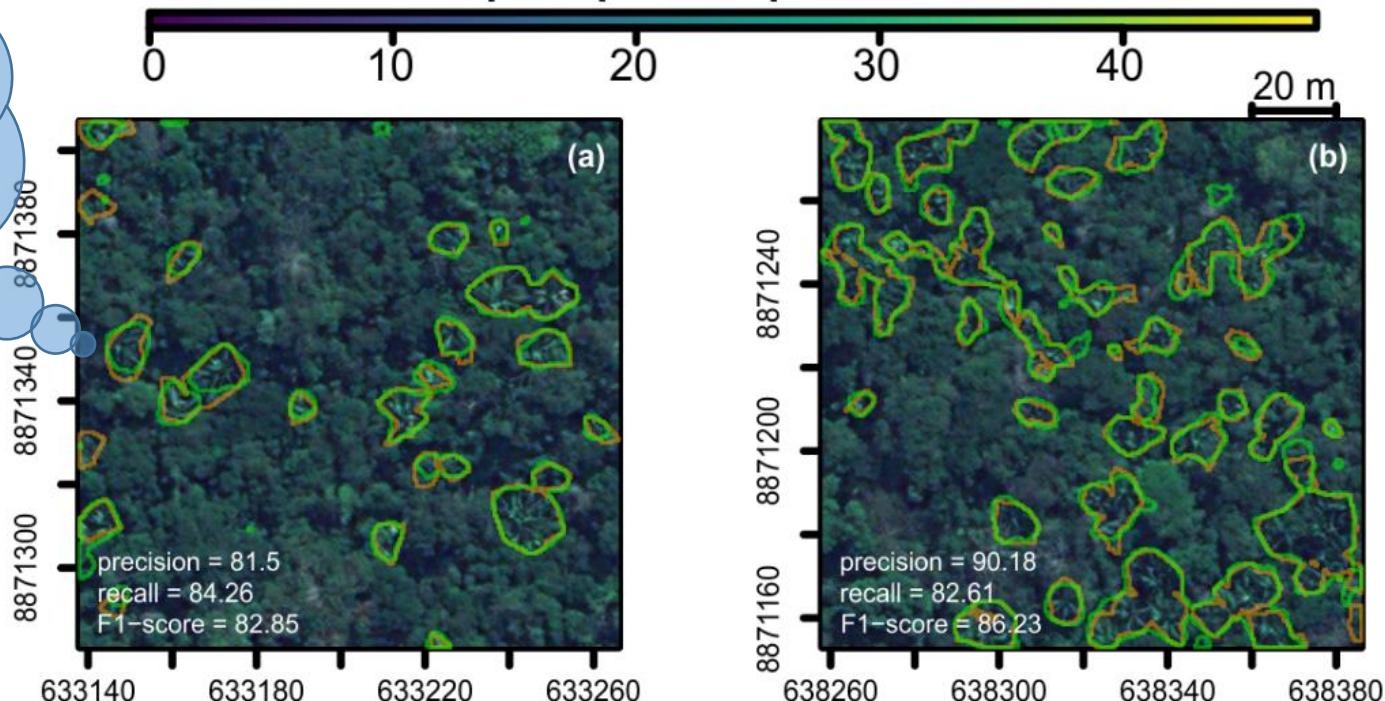
Regional Mapping and Spatial Distribution Analysis of Canopy Palms in an Amazon Forest Using Deep Learning and VHR Images



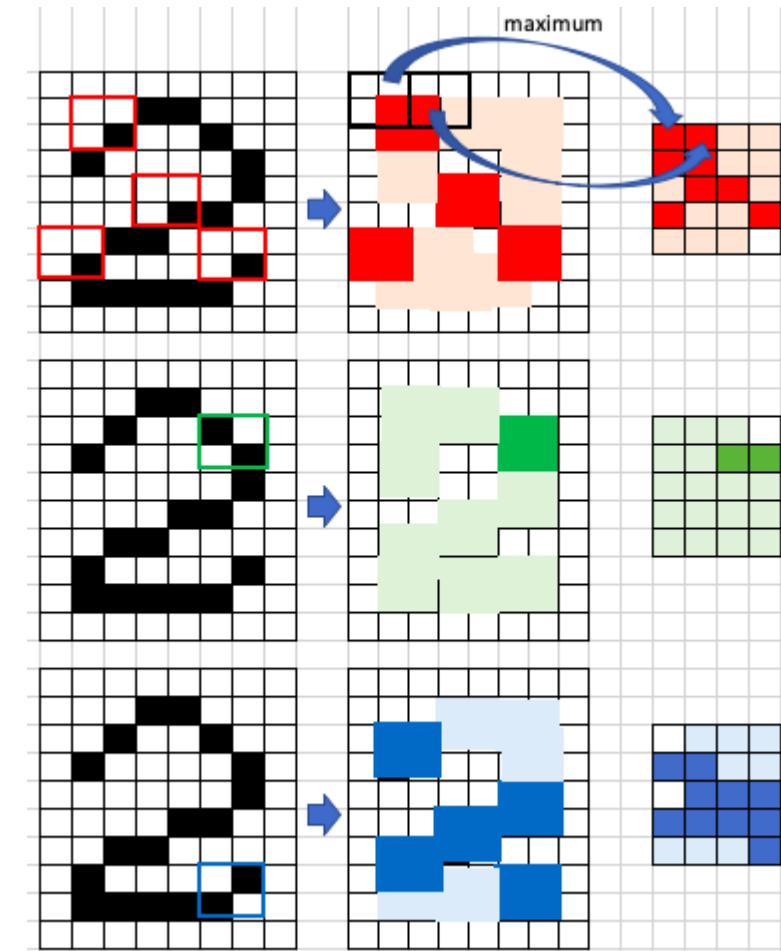
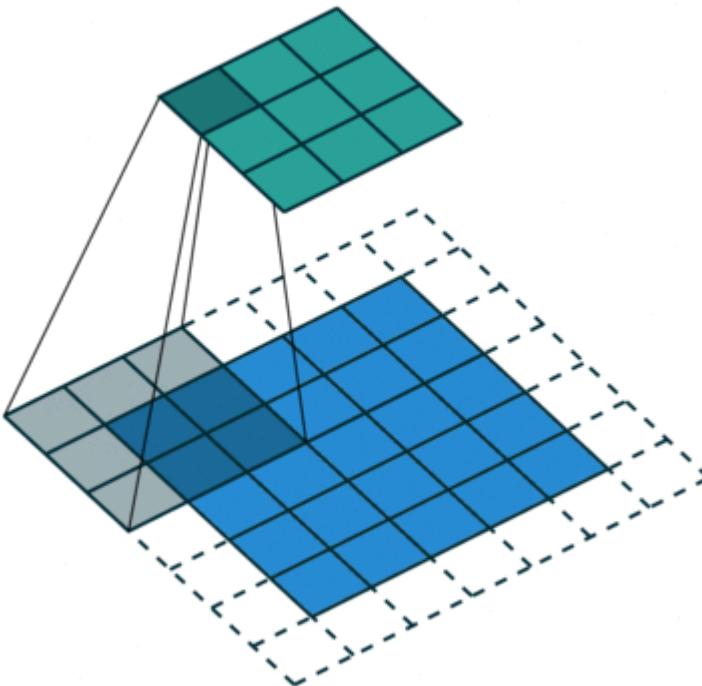
Fabien H. Wagner <sup>1,\*</sup> , Ricardo Dalagnol <sup>2</sup> , Ximena Tagle Casapia <sup>3,4</sup> , Annia S. Streher <sup>2</sup>, Oliver L. Phillips <sup>5</sup> , Emanuel Gloor <sup>5</sup> and Luiz E. O. C. Aragão <sup>2,6</sup>



If you can see  
the target in  
the image,  
deep learning  
can map it



# Convolution



# Convolution

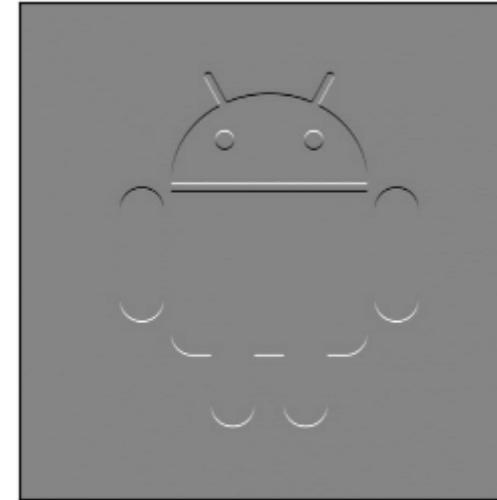


-1	-2	-1
0	0	0
1	2	1

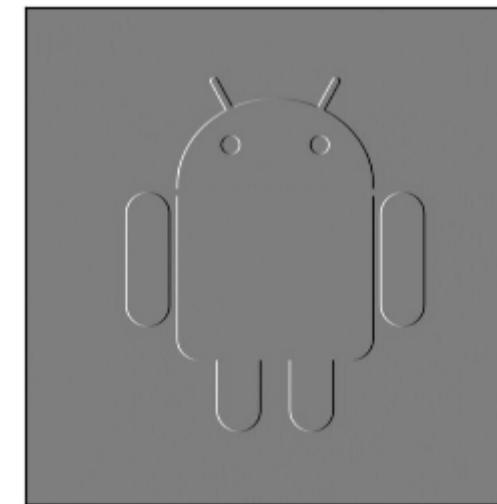
Finds horizontals

-1	0	1
-2	0	2
-1	0	1

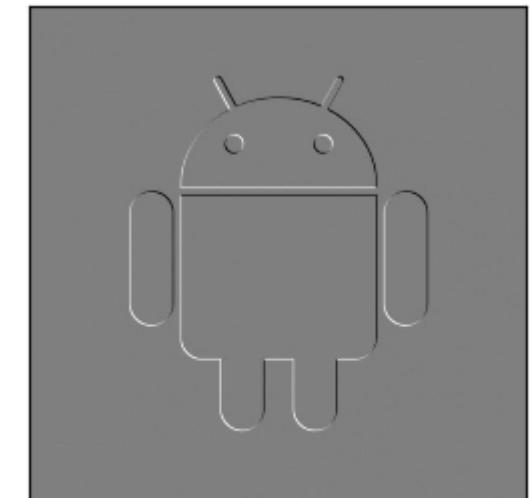
Finds verticals



Horizontal Sobel

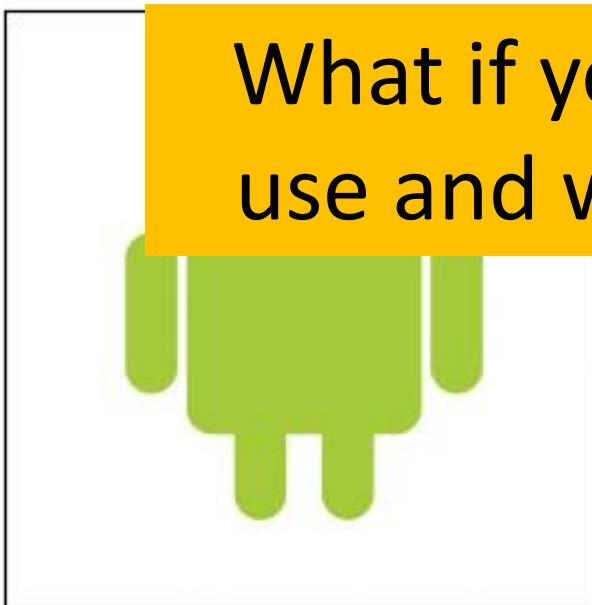


Vertical Sobel



Combined Sobel

# Convolution

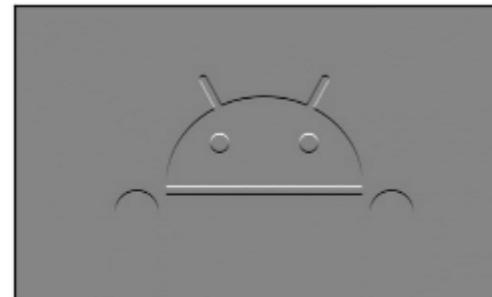


-1	-2	-1
----	----	----

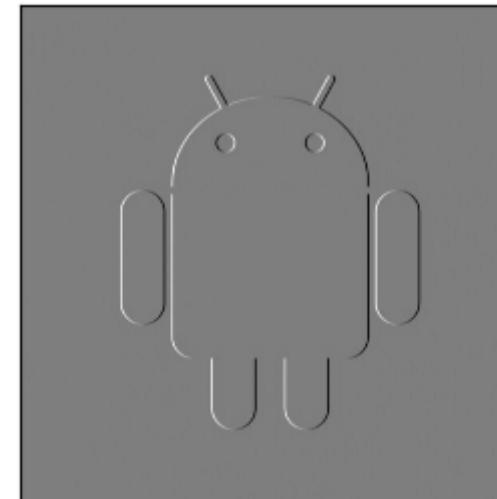
What if you don't know which *filters* to use and which *features* are important?

-1	0	1
-2	0	2
-1	0	1

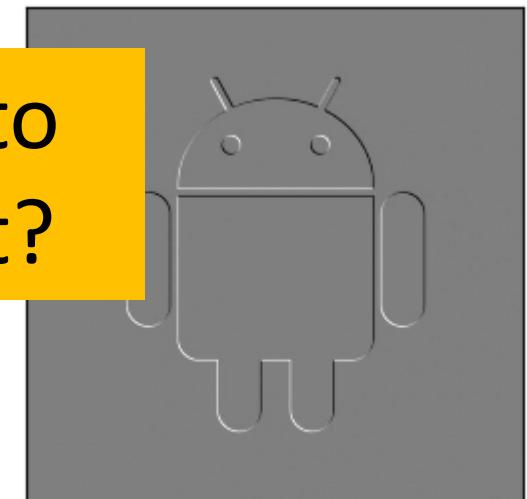
Finds verticals



Horizontal Sobel

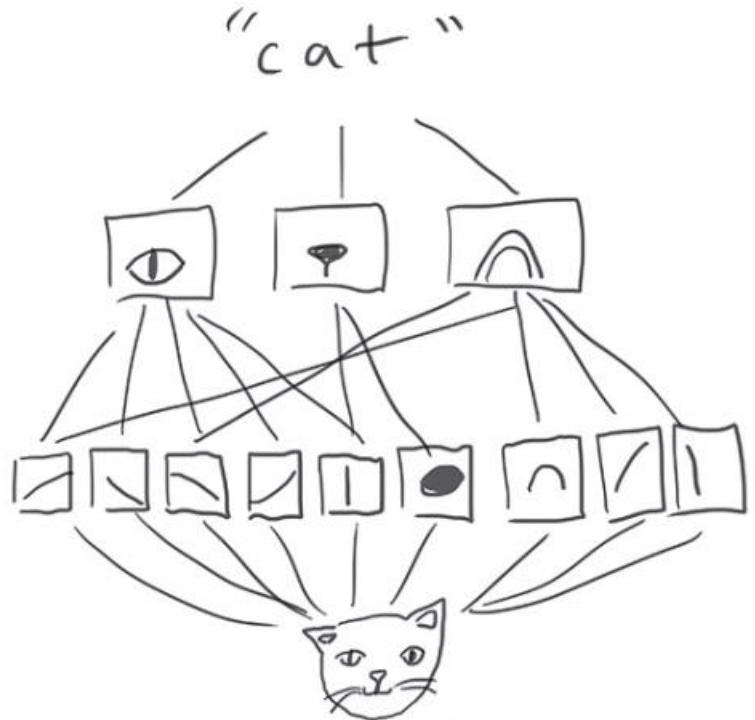
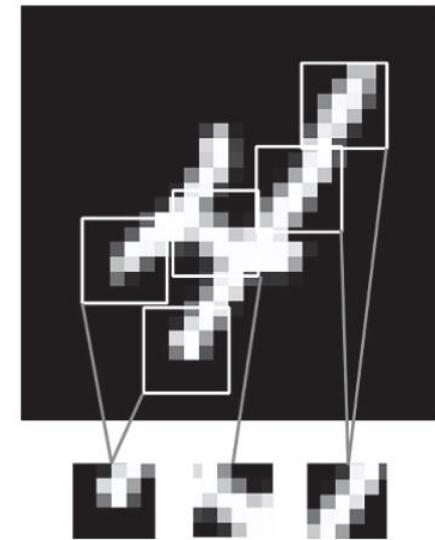
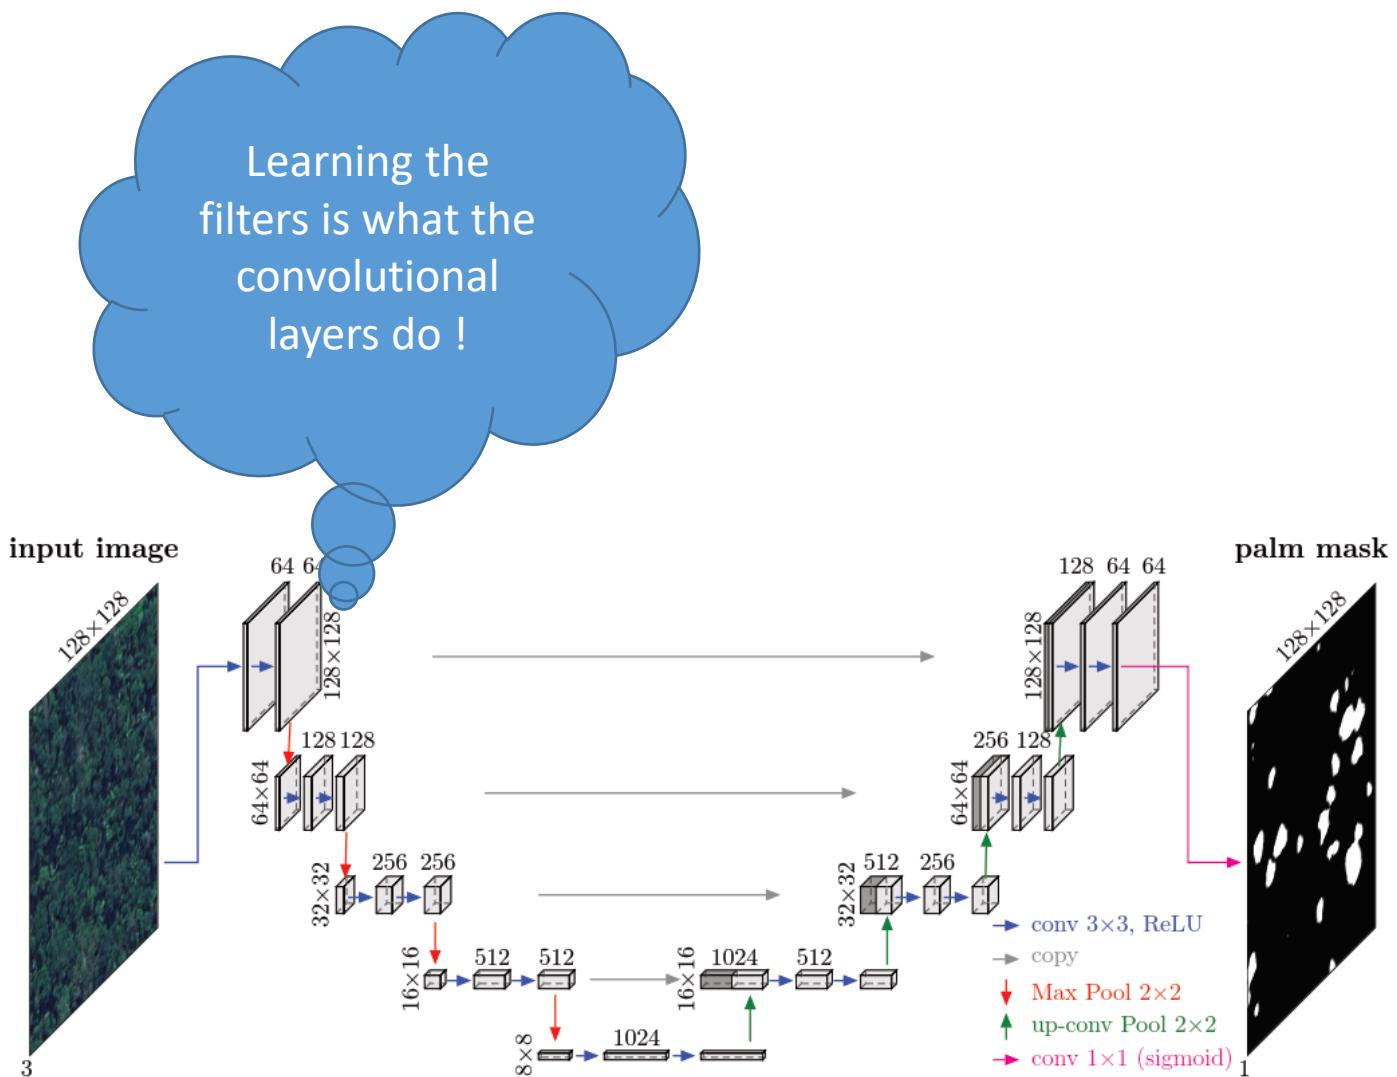


Vertical Sobel



Combined Sobel

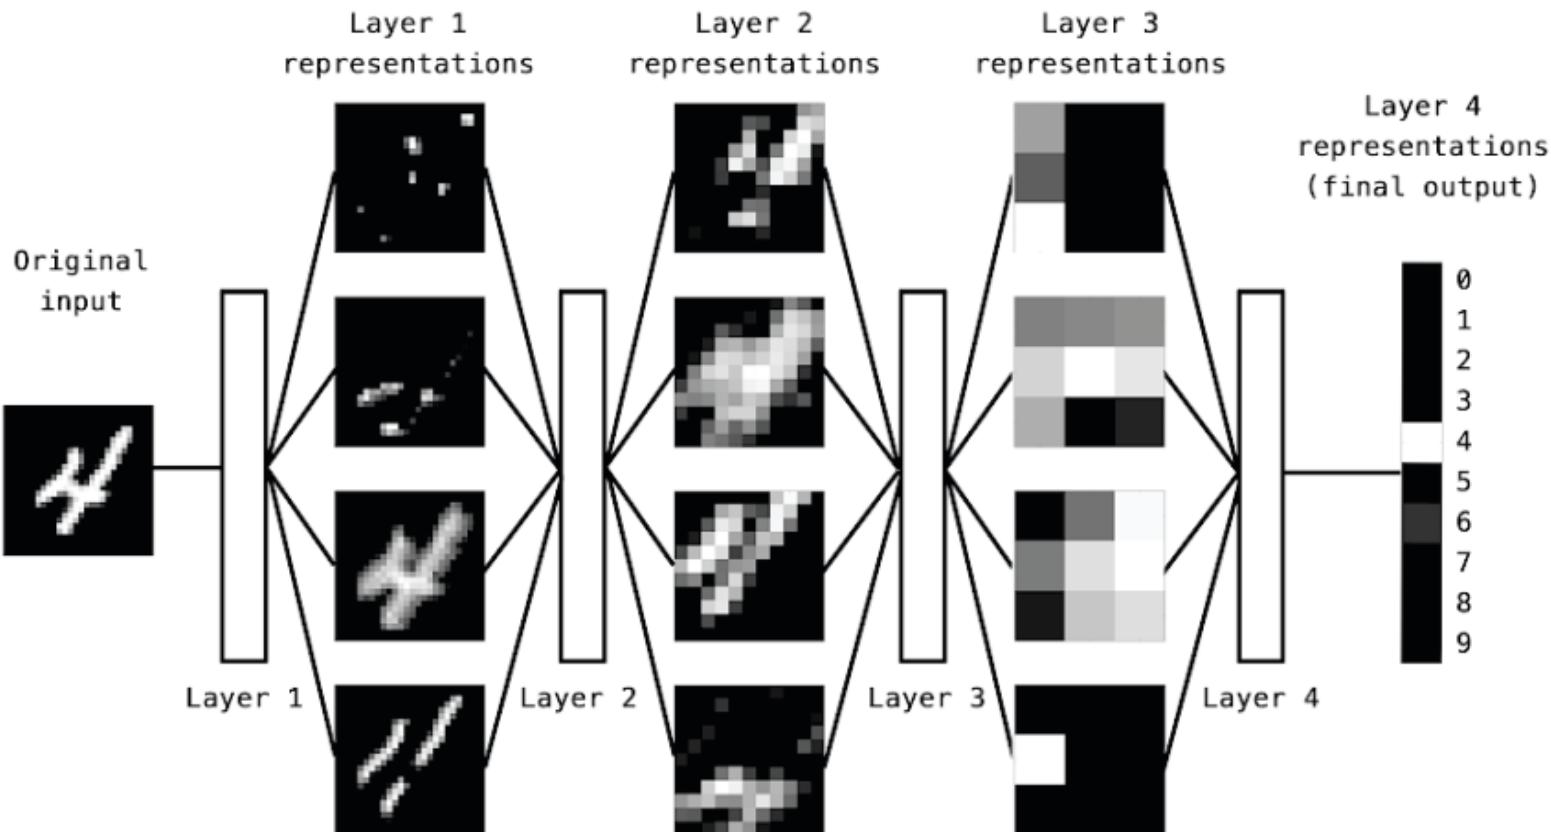
# CNN



**Figure 2.** U-net architecture used for canopy palm segmentation, adapted from [19]. The number of channels and the row  $\times$  column size in pixels are indicated for each cuboid.

(Chollet & Allaire 2017)

# Example with *mnist* dataset: colab python code



# Attention for performance metrics

<https://towardsdatascience.com/metrics-to-evaluate-your-semantic-segmentation-model-6bcb99639aa2>

- Balanced-classification problems
  - Accuracy, Kappa, etc.
- Class-imbalance
  - Precision, Recall, and **F1-score**
- Multi-class segmentation
  - Average precision between classes
- Single-class segmentation
  - **Intersection over Union (IoU)**

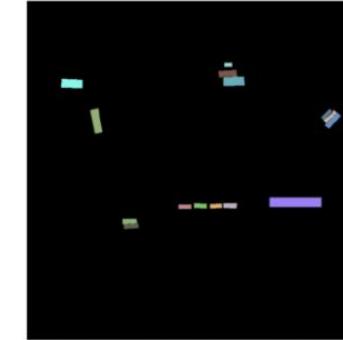
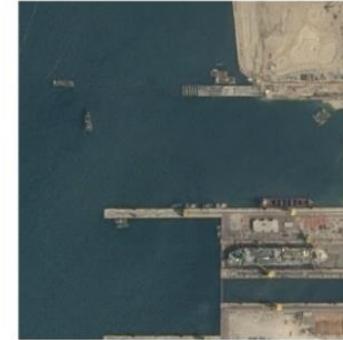
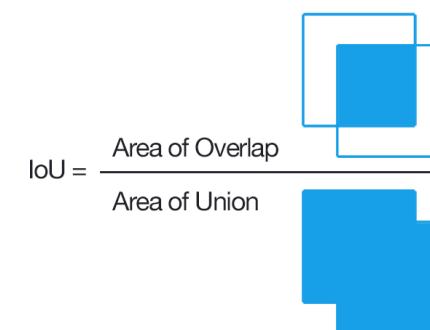
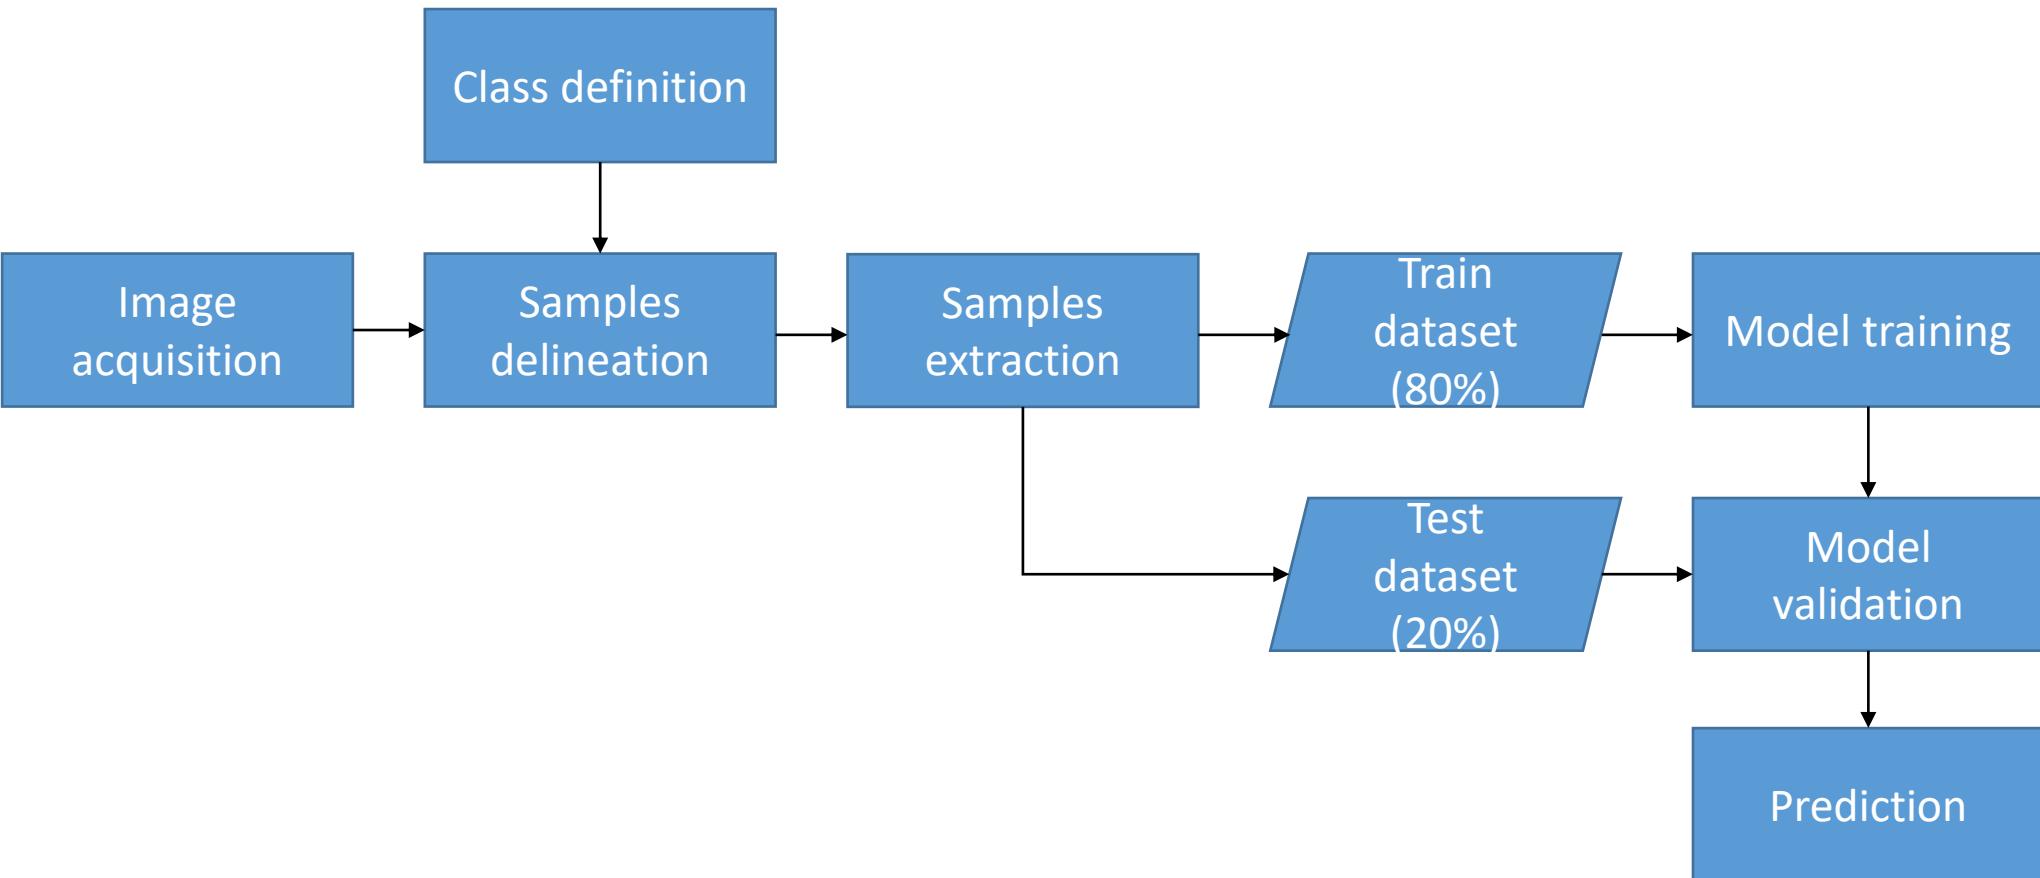


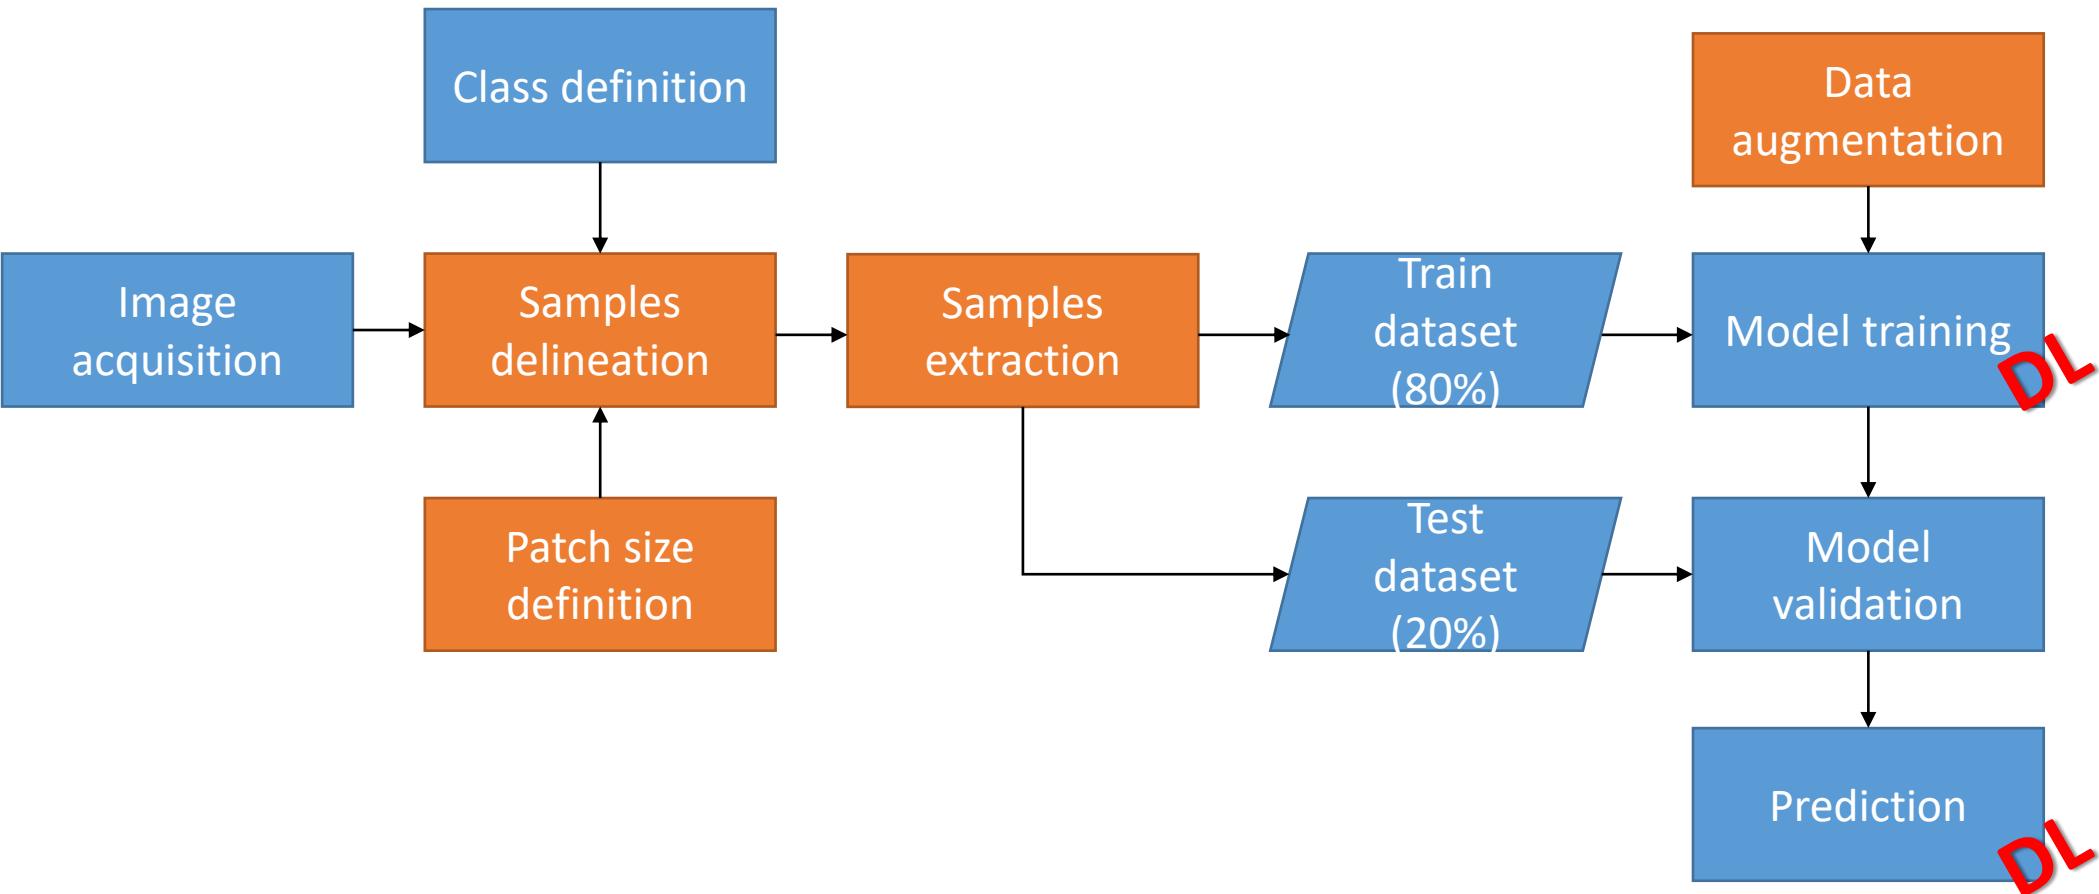
Image from [Vlad Shmyhlo](#) in article: Image Segmentation: Kaggle experience (Part 1 of 2) in TDS



# Main steps for RS image segmentation

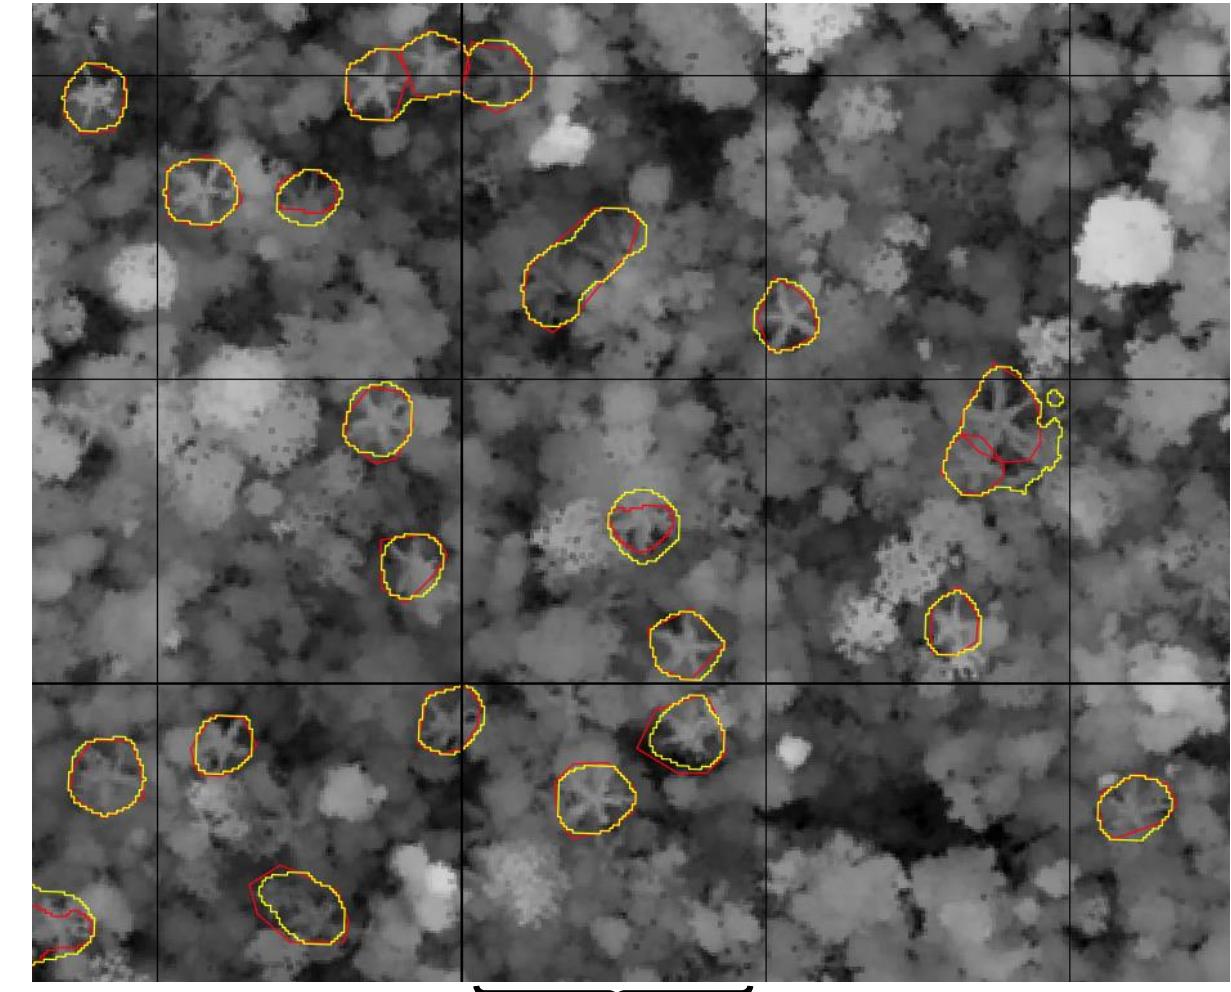


# Main steps for RS image segmentation with DL



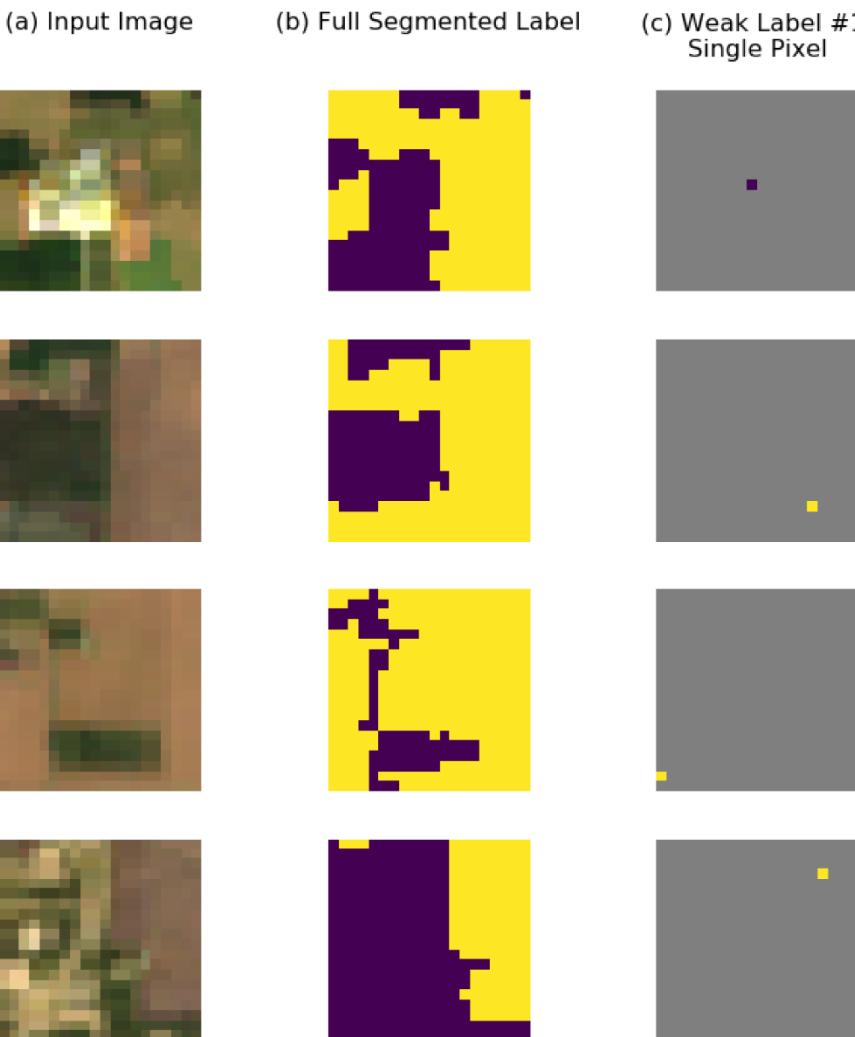
# DL vs ML differences: Patch size definition

- Patch of size  $m \times m$  pixels
  - $m = \text{multiple of } 32$
  - E.g.  $128 \times 128$  pixels
- No optimal definition
- The patch should bring contextual information of the target



64 meters  
(128x128 pixels of 0.5m)

# DL vs ML differences: Samples delineation



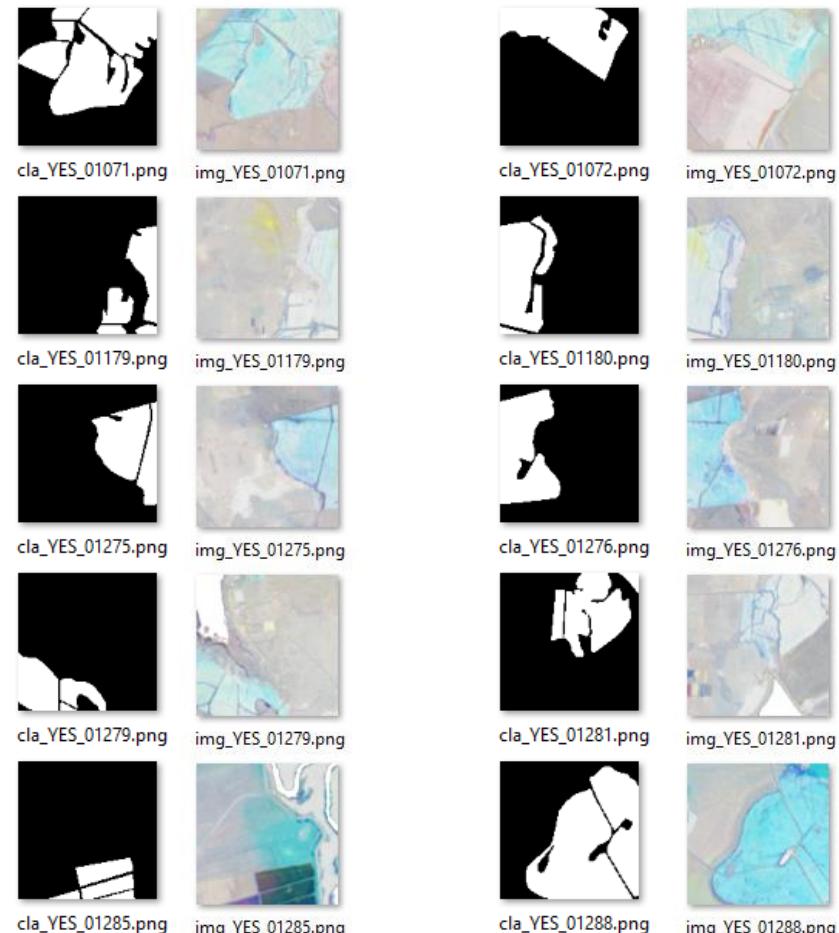
- Samples are usually ‘dense’
- Delineated in GIS or other labelling softwares
- Recent research shows ways to use single pixels called ‘weak supervised learning’ (Wang et al., 2020)

# DL vs ML differences: Samples extraction

## Traditional ML: pixels to table

	X1	X2	X3	X4	class
2129	0.18409561	0.35243168	0.8963459	0.9284320	Rice
2130	0.18615542	0.48938116	0.9341421	0.9476923	Rice
2131	0.24692002	0.25649321	0.5984405	0.8030888	Rice
2132	0.25300574	0.79422522	0.9033156	0.8696288	Rice
2133	0.30483481	0.69084626	0.9067416	0.9049394	Rice
2134	0.25143707	0.25066489	0.8754706	0.9301759	Rice
2135	0.31929952	0.73389256	0.9191375	0.9151656	Rice
2136	0.44001359	0.47436264	0.4754879	0.4471128	Other
2137	0.62812239	0.60058695	0.4603175	0.4119138	Other
2138	0.66906476	0.59619188	0.5153217	0.4031008	Other
2139	0.51898736	0.56076455	0.2089286	0.2098260	Other
2140	0.47804552	0.37455633	0.3560375	0.3706468	Other
2141	0.57003969	0.56485832	0.5612745	0.6241940	Other

## Deep Learning: patches to files



When predicting, we crop, apply model, then mosaic.

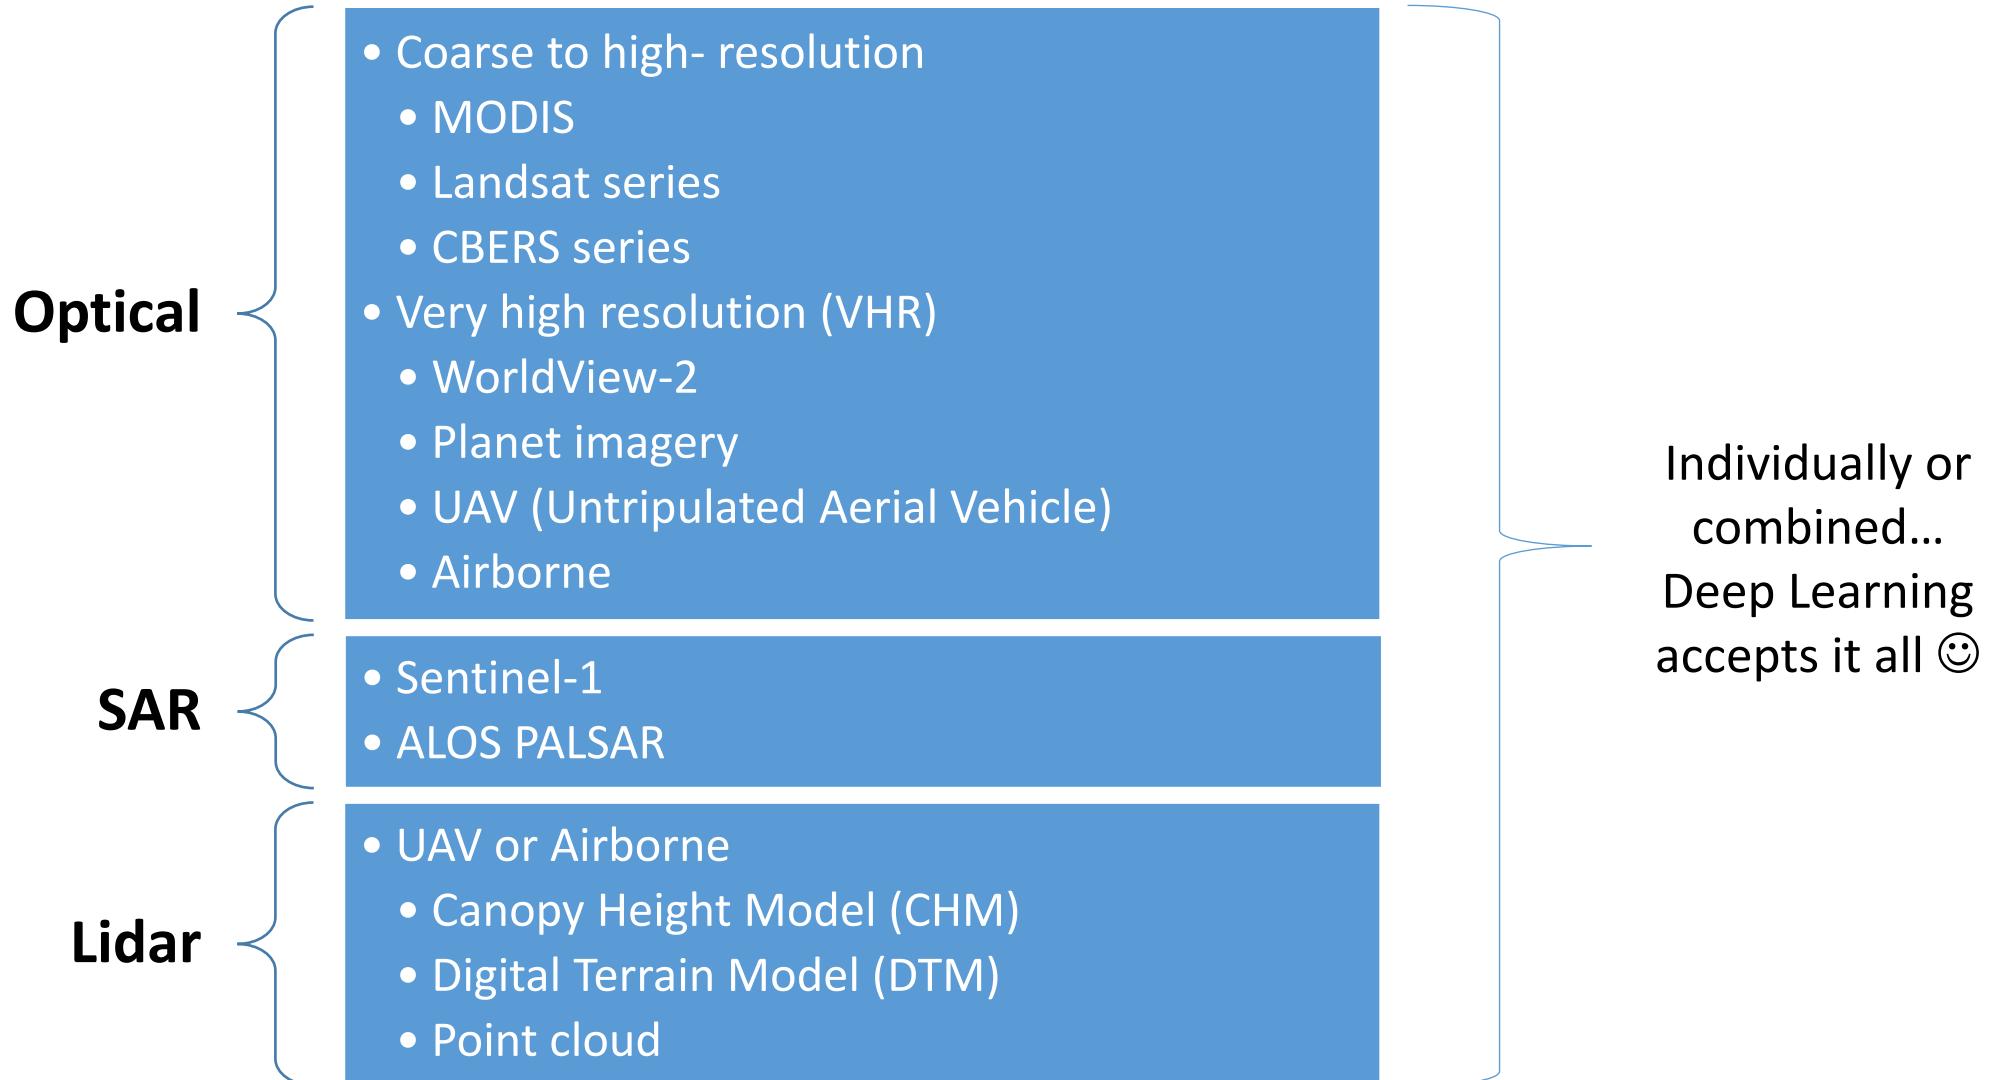
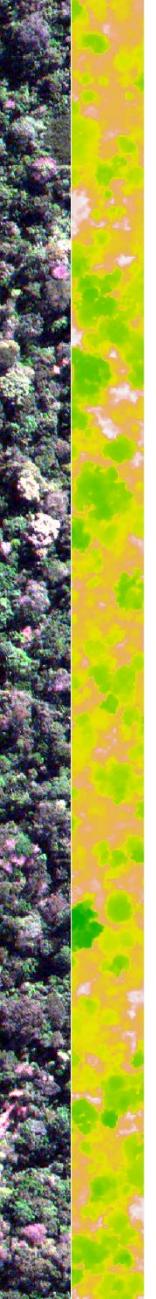
# DL vs ML differences: Data augmentation

- Randomized processes that introduces variability in the dataset to teach the model that patterns are not limited to certain angles, values, scales, etc.
  - 0-360° rotations
  - Horizontal and vertical flips
  - Signal change in brightness and saturation (e.g. 90 to 110%)
  - Zoom



(Chollet & Allaire 2017)

# What types of RS data can I use? All types!



# CPU vs GPU

<https://colab.research.google.com/notebooks/gpu.ipynb#scrollTo=Y04m-jvKRDsJ>

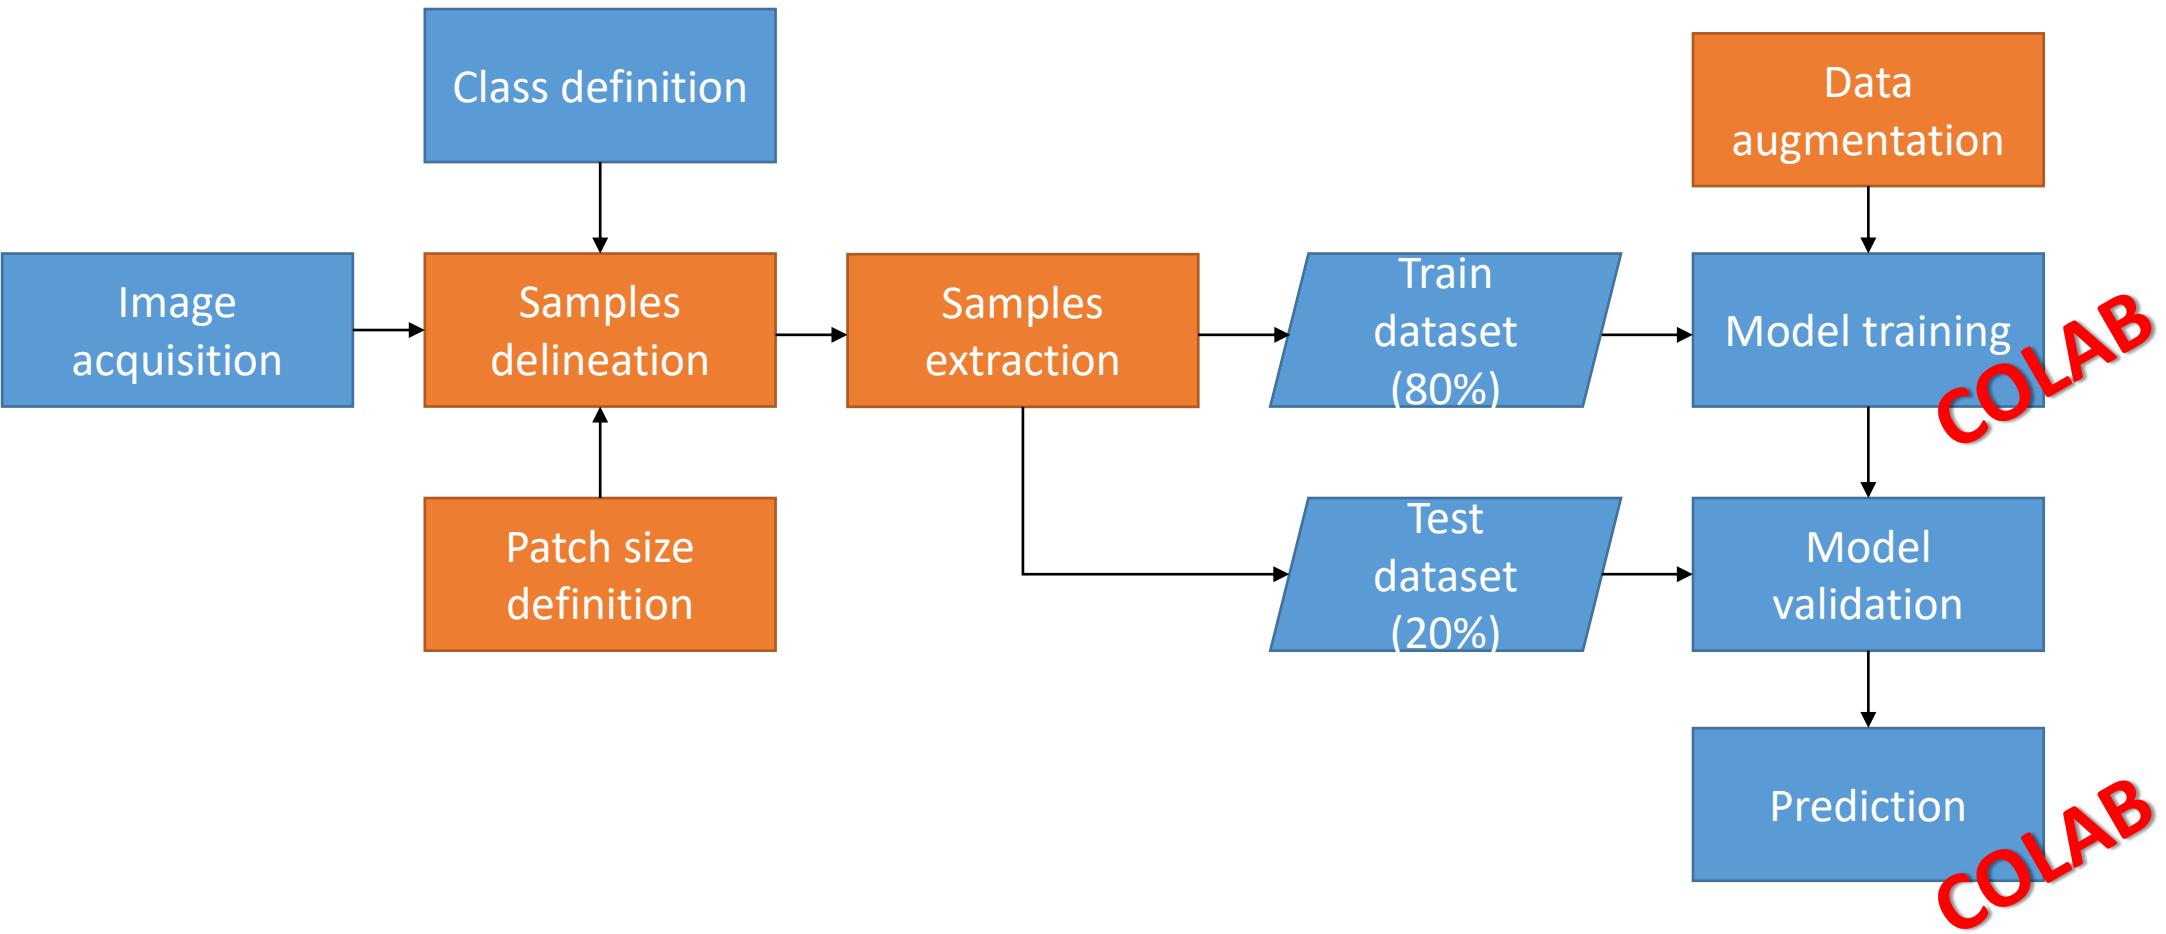
**GPU was 85x faster than CPU at convolution applications**



# Which language and platform to run DL?

- Python: tensorflow, keras, pytorch
- R: with reticulate / keras <- tensorflow
- Your own computer with GPU
- Google Colab: free, but limited – ideal for beginners
- Paid-virtual machines with GPU (from Amazon, Google, etc.)

# Main steps for RS image segmentation with DL





zenodo

July 3, 2020

Lesson  Open Access

Edit

New version

1,374

views

436

downloads

[See more details...](#)

Indexed in

OpenAIRE

Download the PDF file and follow the instructions. This is for NVIDIA GPU and Windows OS.

**Contacts :** Ricardo Dalagnol (ricds@hotmail.com)

Fabien Wagner (wagner.h.fabien@gmail.com)

v2.0 Updated for Python 3.9 and Tensorflow 2.5.

Send an email for collaboration or if you have suggestions for the improvement of this manual.

Preview

Página: 1 de 4

Zoom automático

Version 2.0 – 13 July 2021

Tensorflow and Keras installation steps for Deep Learning applications in Rstudio

Ricardo Dalagnol<sup>1</sup> and Fabien Wagner

1. Install OSGEO (<https://trac.osgeo.org/osgeo4w>) to have GDAL utilities and QGIS. Do not modify the installation path. This is mainly for the remote sensing community, to work with Geotiff.
2. Install ImageMagick (<https://imagemagick.org/index.php>) for image visualization.
3. Install the latest Miniconda (<https://docs.conda.io/en/latest/miniconda.html>)
4. Install Visual Studio
  - a. Install the 2019 or 2017 free version. The free version is called ‘Community’.

Publication date:

July 3, 2020

DOI:

Keyword(s):

License (for files):

Creative Commons Attribution 4.0 International

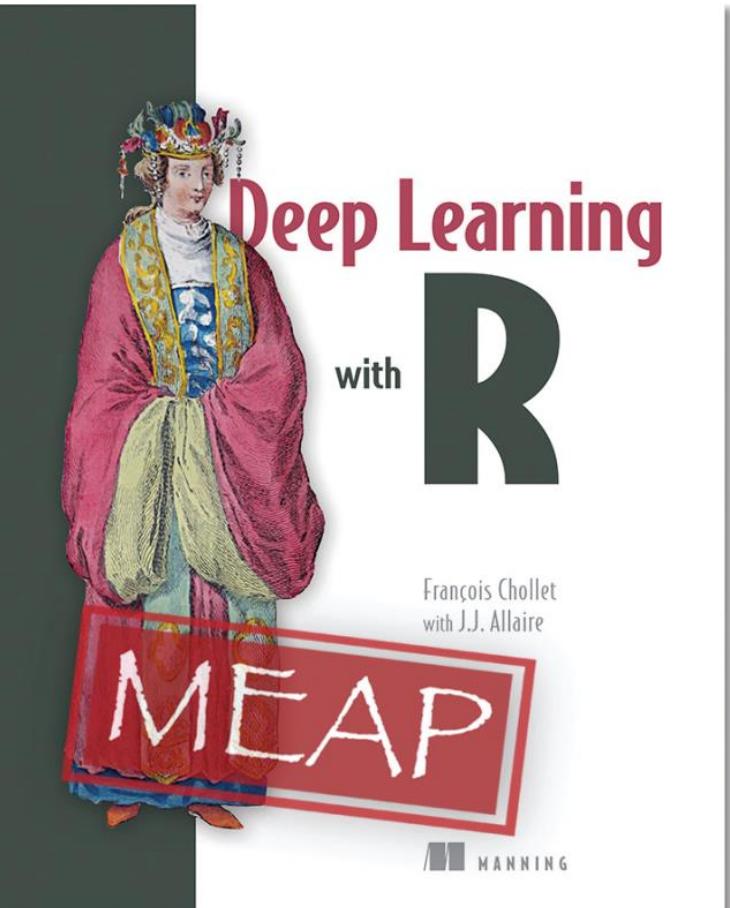
Versions

# End-to-end Deep Learning in the cloud

- Google Colab: limited time, may have installation issues
- Paid-virtual machines with GPU (from Amazon, Google, etc.): should work better

# References

**“DL Bible”**



(Chollet & Allaire 2017)

Overview

Keras

Getting started

Overview

Keras basics

Sequential API

Functional API

Saving and serializing models

Frequently Asked Questions

Advanced

## Getting Started with Keras

### Overview

Keras is a high-level neural networks API developed with a focus on enabling fast experimentation. *Being able to go from idea to result with the least possible delay is key to doing good research.* Keras has the following key features:

- Allows the same code to run on CPU or on GPU, seamlessly.
- User-friendly API which makes it easy to quickly prototype deep learning models.
- Built-in support for convolutional networks (for computer vision), recurrent networks (for sequence processing), and any combination of both.
- Supports arbitrary network architectures: multi-input or multi-output models, layer sharing, model sharing, etc. This means that Keras is appropriate for building essentially any deep learning model, from a memory network to a neural Turing machine.

This website provides documentation for the R interface to Keras. See the main Keras website at <https://keras.io> for additional information on the project.

## Object Detection and Image Segmentation with Deep Learning on Earth Observation Data: A Review-Part I: Evolution and Recent Trends

by  Thorsten Hoeser <sup>1,\*</sup>  and  Claudia Kuenzer <sup>1,2</sup> 

<sup>1</sup> German Remote Sensing Data Center (DFD), German Aerospace Center (DLR), Münchner Straße 20, D-82234 Wessling, Germany

<sup>2</sup> Department of Remote Sensing, Institute of Geography and Geology, University Würzburg, Am Huband, D-97074 Wuerzburg, Germany

\* Author to whom correspondence should be addressed.

*Remote Sens.* **2020**, *12*(10), 1667; <https://doi.org/10.3390/rs12101667>

Received: 27 April 2020 / Revised: 19 May 2020 / Accepted: 19 May 2020 / Published: 22 May 2020

Hoeser & Kuenzer (2020) *Remote sensing*

<https://www.mdpi.com/2072-4292/12/10/1667/htm>

# Summary

- Deep learning are state-of-the-art models for image segmentation. You should start to do DL now 😊
- Hope you learned something new today, thanks !

Contact



ricds@hotmail.com



@RicardoDalagnol